

Problem Statement

- It is a critical requirement for business to understand the value derived from a customer. RFM is a method used for analyzing customer value.
- Customer segmentation is the practice of segregating the customer base into groups of individuals based on some common characteristics such as age, gender, interests, and spending habits
- Perform customer segmentation using RFM analysis. The resulting segments can be ordered from most valuable (highest recency, frequency, and value) to least valuable (lowest recency, frequency, and value).

Dataset Description

This is a transnational data set which contains all the transactions that occurred between 01/12/2010 and 09/12/2011 for a UK-based and registered non-store online retail. The company mainly sells unique and all-occasion gifts.

Variables	Description
InvoiceNo	Invoice number. Nominal, a six digit integral number uniquely assigned to each transaction. If this code starts with letter 'c', it indicates a cancellation
StockCode	Product (item) code. Nominal, a five digit integral number uniquely assigned to each distinct product
Description	Product (item) name. Nominal
Quantity	The quantities of each product (item) per transaction. Numeric
InvoiceDate	Invoice Date and time. Numeric, the day and time when each transaction was generated
UnitPrice	Unit price. Numeric, product price per unit in sterling
CustomerID	Customer number. Nominal, a six digit integral number uniquely assigned to each customer
Country	Country name. Nominal, the name of the country where each customer resides

Write Up

Week 1

Data Cleaning:

1. Perform a preliminary data inspection and data cleaning.
 - a. Check for missing data and formulate an apt strategy to treat them.
 - b. Remove duplicate data records.
 - c. Perform descriptive analytics on the given data.

Data Transformation:

2. Perform cohort analysis (a cohort is a group of subjects that share a defining characteristic). Observe how a cohort behaves across time and compare it to other cohorts.
 - a. Create month cohorts and analyze active customers for each cohort.
 - b. Analyze the retention rate of customers.

Data Cleaning:

1. Perform a preliminary data inspection and data cleaning.
 - a. Check for missing data and formulate an apt strategy to treat them.
 - b. Remove duplicate data records.
 - c. Perform descriptive analytics on the given data.

541909 rows x 8 columns

8 columns are available.

The item related are - Stock Code (Quantifiable), & Description

Sale Related are - Invoice number and Invoice Date & Quantity & Unit Price

Customer Related are - Customer Id & Country

The main or basic inferred data are - Spending pattern, Spending categories, Customer Spending Behavior

- a. Check for missing data and formulate an apt strategy to treat them.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	InvoiceNo	541909	non-null object
1	StockCode	541909	non-null object
2	Description	540455	non-null object
3	Quantity	541909	non-null int64
4	InvoiceDate	541909	non-null datetime64[ns]
5	UnitPrice	541909	non-null float64
6	CustomerID	406829	non-null float64
7	Country	541909	non-null object

dtypes: datetime64[ns] (1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
Column - Description, & CustomerID, have some Null Values in them

Description Column is No Value add so there is no problem in dropping it. Since Customer Id is the main identifying element, Unique Identifying Entity, its absence would be difficult to fill through Unlike cost or Sale Unit, where we could use average, we cannot use any other means to treat this except deletion.

Identifying & Removing Null Values. Starting from the column that has max null values. If it clears the null values in other columns, we would not need to repeat the activity with other columns

InvoiceNo	0
StockCode	0
Description	0
Quantity	0
InvoiceDate	0
UnitPrice	0
CustomerID	0
Country	0

dtype: int64

b. Identify & Remove duplicate data records.

Get names of indexes for which Column Unit Price which has value negative value. Assuming that the shop keeper does not pay customer to purchase. There is one such instance, removing this data set as incorrect

c. Perform descriptive analytics on the given data.

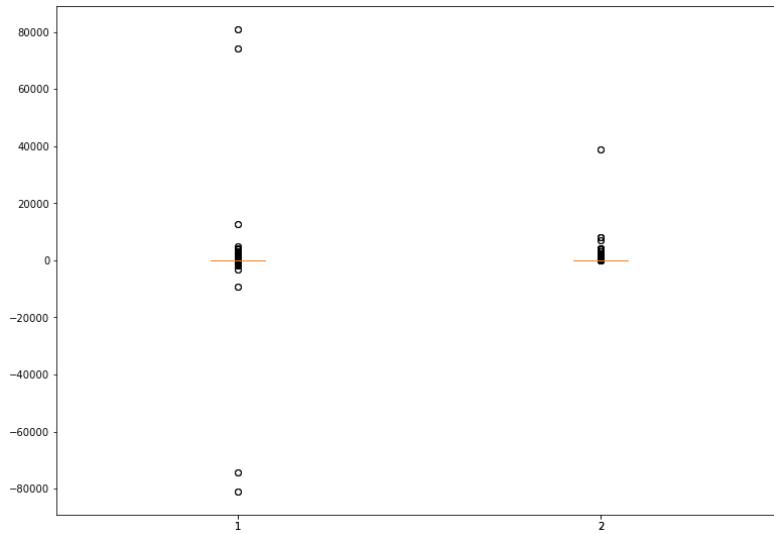
However, not doing the same with Quantity. The assumption here is that these might be billed in previous cycle which is not included in this database and were returned to seller in this cycle. Approx ~2% was returned

	Quantity	UnitPrice	CustomerID
count	401604.000000	401604.000000	401604.000000
mean	12.183273	3.474064	15281.160818
std	250.283037	69.764035	1714.006089
min	-80995.000000	0.000000	12346.000000
25%	2.000000	1.250000	13939.000000
50%	5.000000	1.950000	15145.000000
75%	12.000000	3.750000	16784.000000
max	80995.000000	38970.000000	18287.000000

#Counting % of Returns

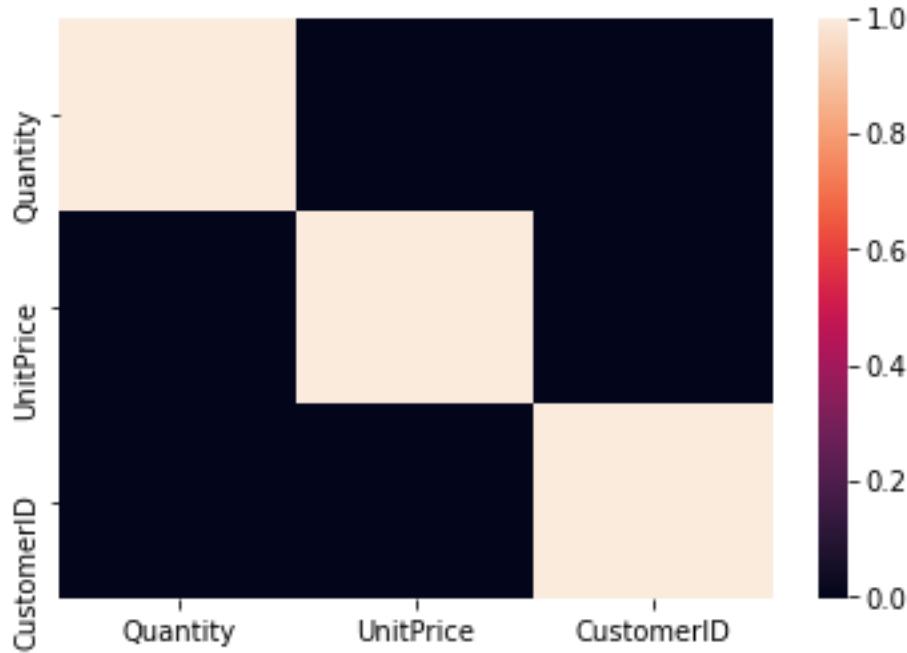
2.209141343213713

Quantity	UnitPrice	CustomerID	
count	401604.000000	401604.000000	401604.000000
mean	12.183273	3.474064	15281.160818
std	250.283037	69.764035	1714.006089
min	-80995.000000	0.000000	12346.000000
25%	2.000000	1.250000	13939.000000
50%	5.000000	1.950000	15145.000000
75%	12.000000	3.750000	16784.000000
max	80995.000000	38970.000000	18287.000000



AS we see from Standard Deviation and also the boxplot that there are lot of outliers in Quantity ordered and in Unit Price Now for Descriptive Analysis.

	Quantity	UnitPrice	CustomerID
Quantity	1.000000	-0.001243	-0.003457
UnitPrice	-0.001243	1.000000	-0.004524
CustomerID	-0.003457	-0.004524	1.000000



Buyers are from 36 different countries

#Unique Customers are 4372 out of which #Only 79 buyer purchased once

#Percentage of single purchaser is, 1.80

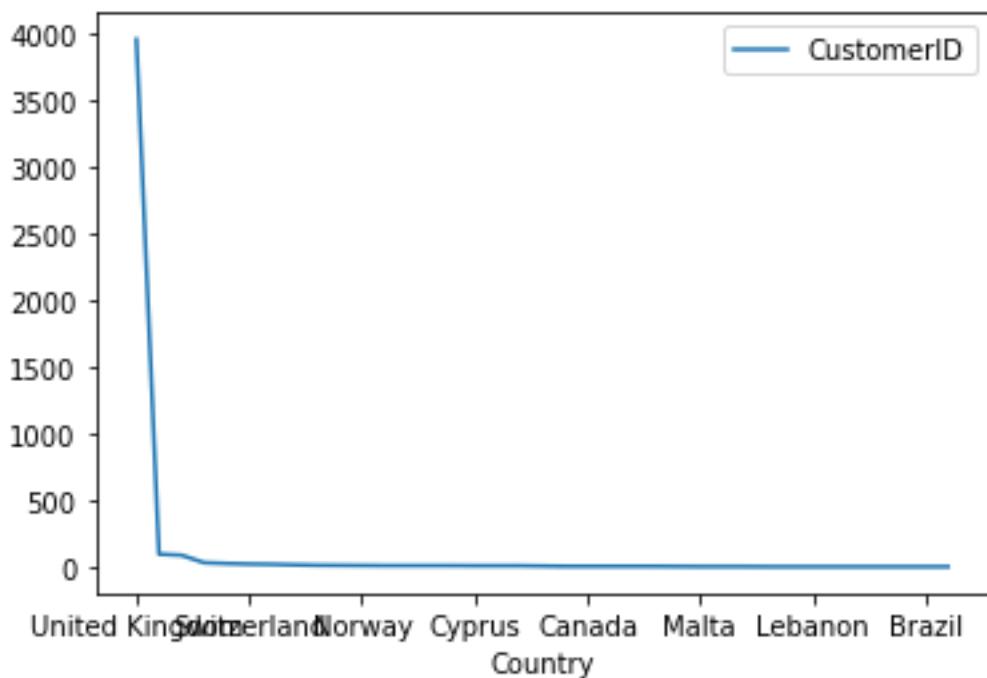
#Repeat Customers are 4293

Freq	CustomerID
17841.0	7812 17841.0
14911.0	5898 14911.0
14096.0	5128 14096.0
12748.0	4459 12748.0
14606.0	2759 14606.0
...	...
15423.0	2 15423.0
14642.0	2 14642.0
13130.0	2 13130.0
13298.0	2 13298.0
14821.0	2 14821.0

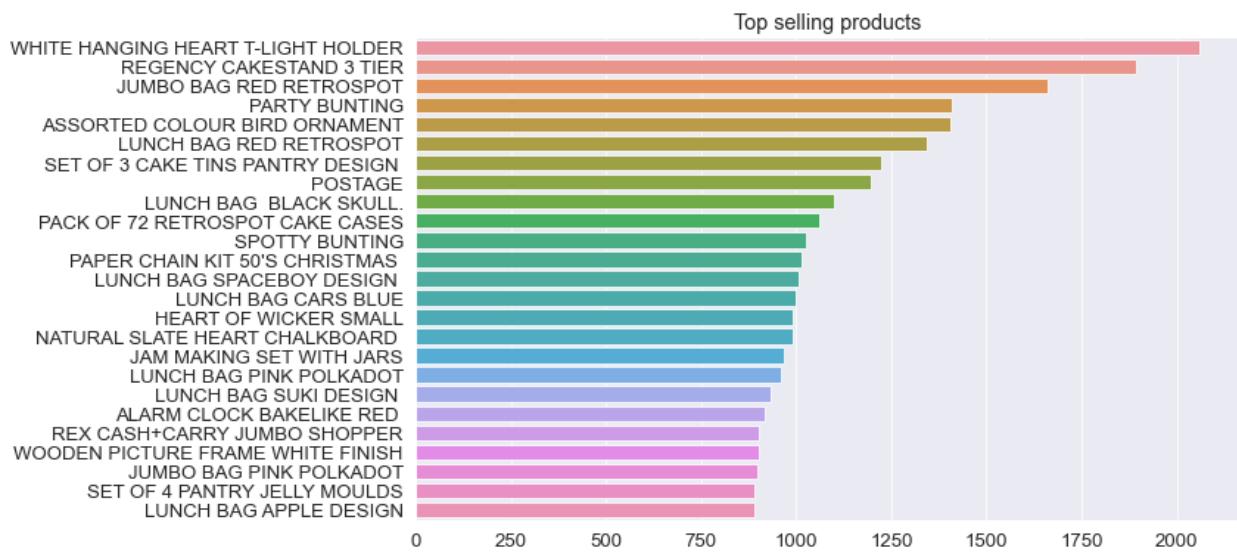
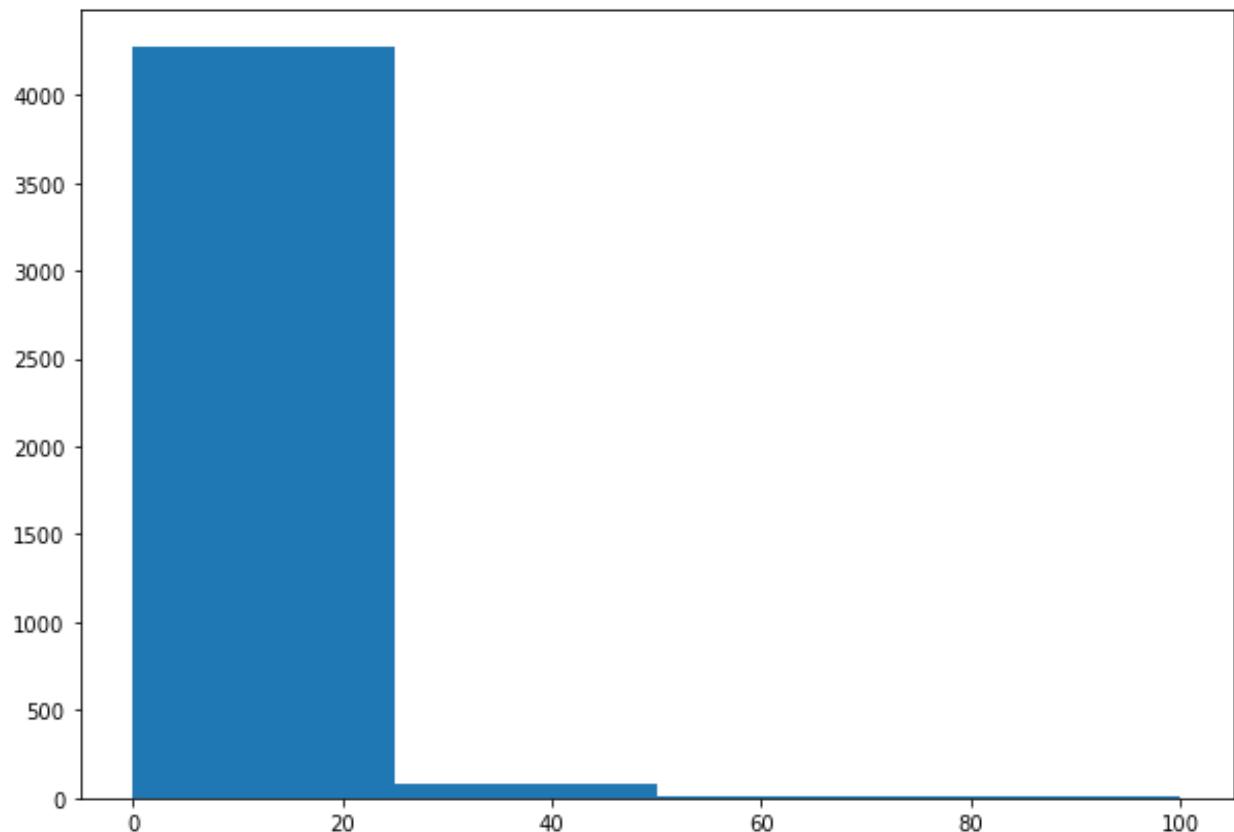
4293 rows x 2 columns

Change in Corel seen slightly when we plot only the customer who purchased once

	Quantity	UnitPrice	CustomerID
Quantity	1.000000	-0.048550	-0.201286
UnitPrice	-0.048550	1.000000	0.082662
CustomerID	-0.201286	0.082662	1.000000



69.97% of customers ordered more than one item.

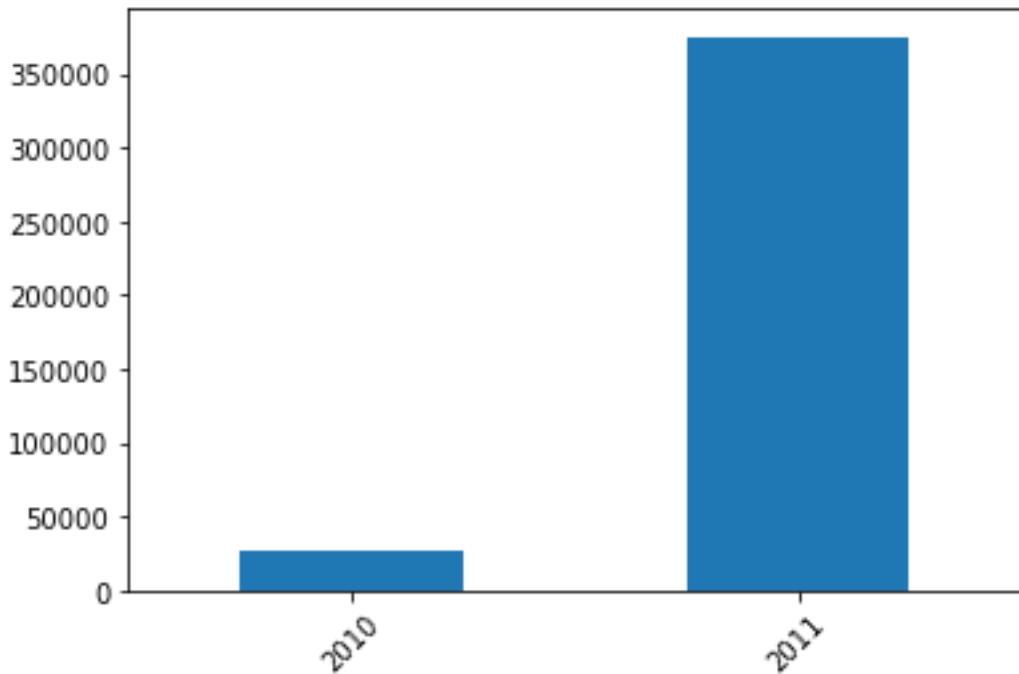


Check the oldest and latest date in the dataset.

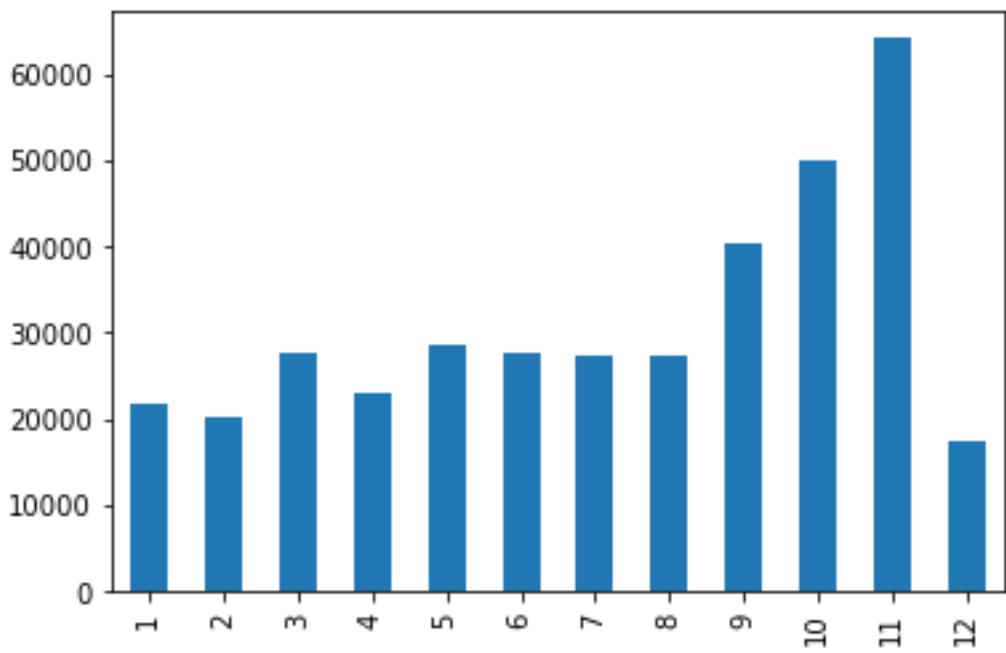
Oldest date is - 2010-12-01 08:26:00

Latest date is - 2011-12-09 12:50:00

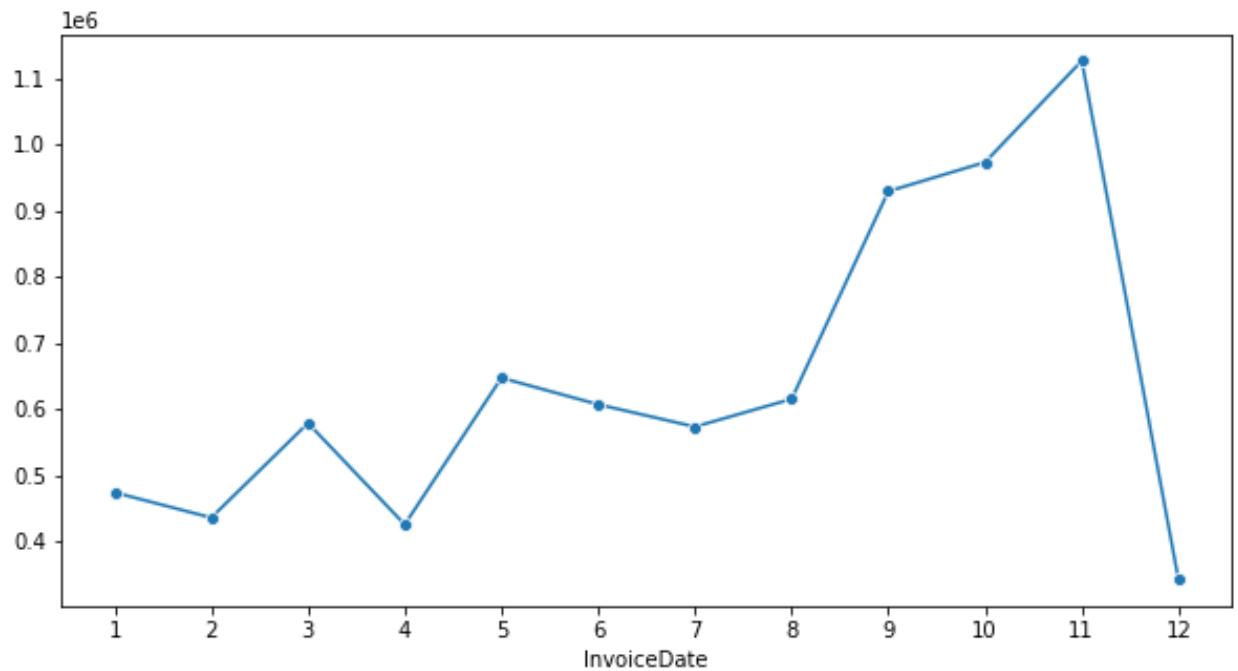
Count of transactions in different years



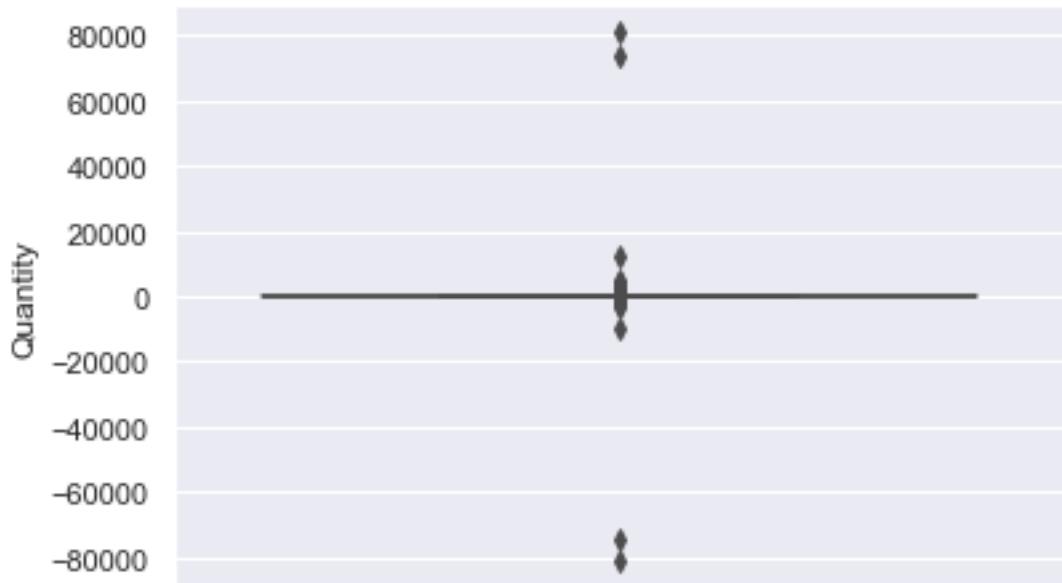
Most of the records belong to 2011. Now doing monthly break up. And we see that max transaction is in Nov & Oct. Could be Black Friday or Halloween. Dec is lesser does not indicating advance buying, because in this data, sales till 09th is considered.

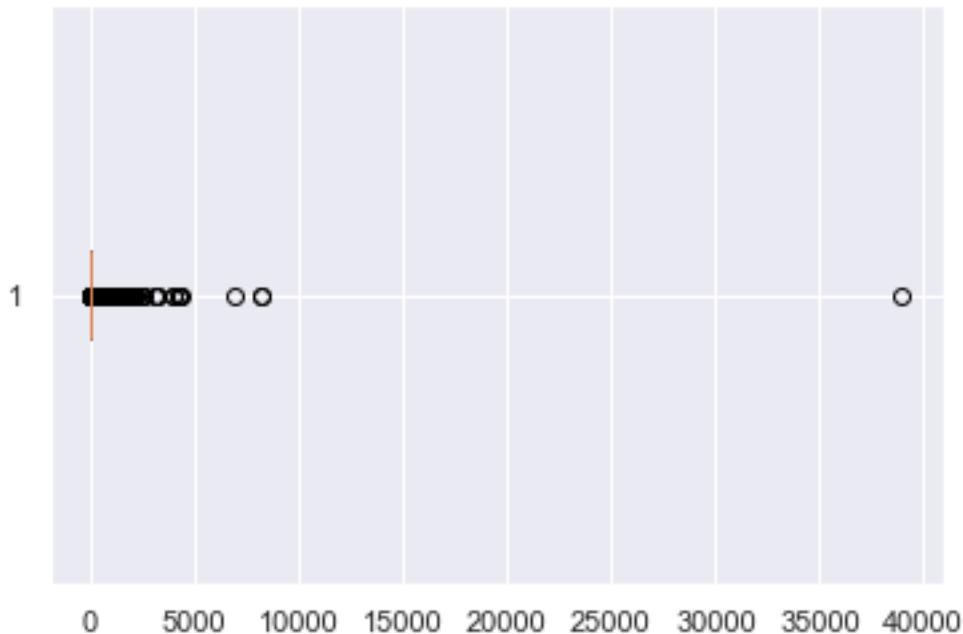


#Monthly Sales & monthly_gross is analysis. It follows the same pattern



Boxplot to for Quantity distribution is checked for outliers. Outliers seen in #Unit price distribution





Data Transformation:

2. Perform cohort analysis (a cohort is a group of subjects that share a defining characteristic). Observe how a cohort behaves across time and compare it to other cohorts.

- a. Create month cohorts and analyze active customers for each cohort.
- b. Analyze the retention rate of customers.

Perform cohort analysis (a cohort is a group of subjects that share a defining characteristic). Observe how a cohort behaves across time and compare it to other cohorts. a. Create month cohorts and analyze active customers for each cohort. b. Analyze the retention rate of customers.

Time cohorts Time cohorts are customers who signed up for a product or service during a particular time frame. Analysing these cohorts shows the customers' behaviour depending on the time they started using the company's products or services. The time may be monthly or quarterly, even daily.

#Assigning Cohort to each group

#monthly cohorts based on the month each customer has made their first transaction.

Group by CustomerID and select the InvoiceMonth value

Calculate time offset in months

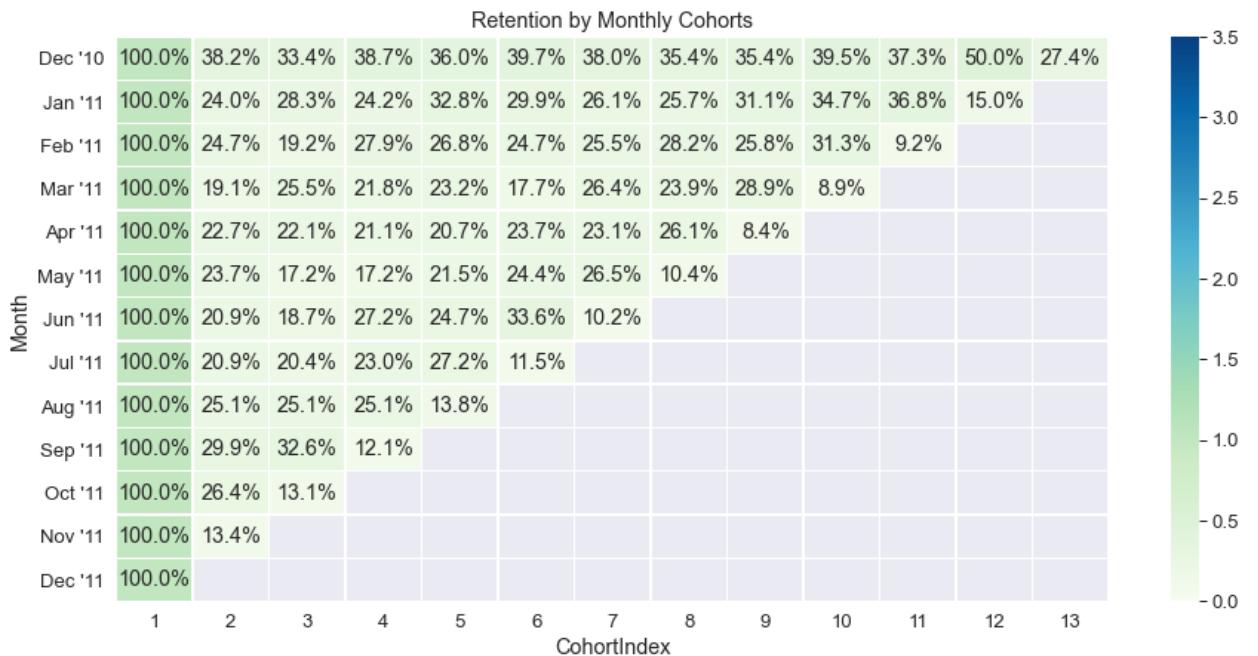
Calculating time offset for each transaction allows you to report the metrics for each cohort in a comparable fashion.

First, we will create some variables that capture the integer value of years and months for Invoice and Cohort Date

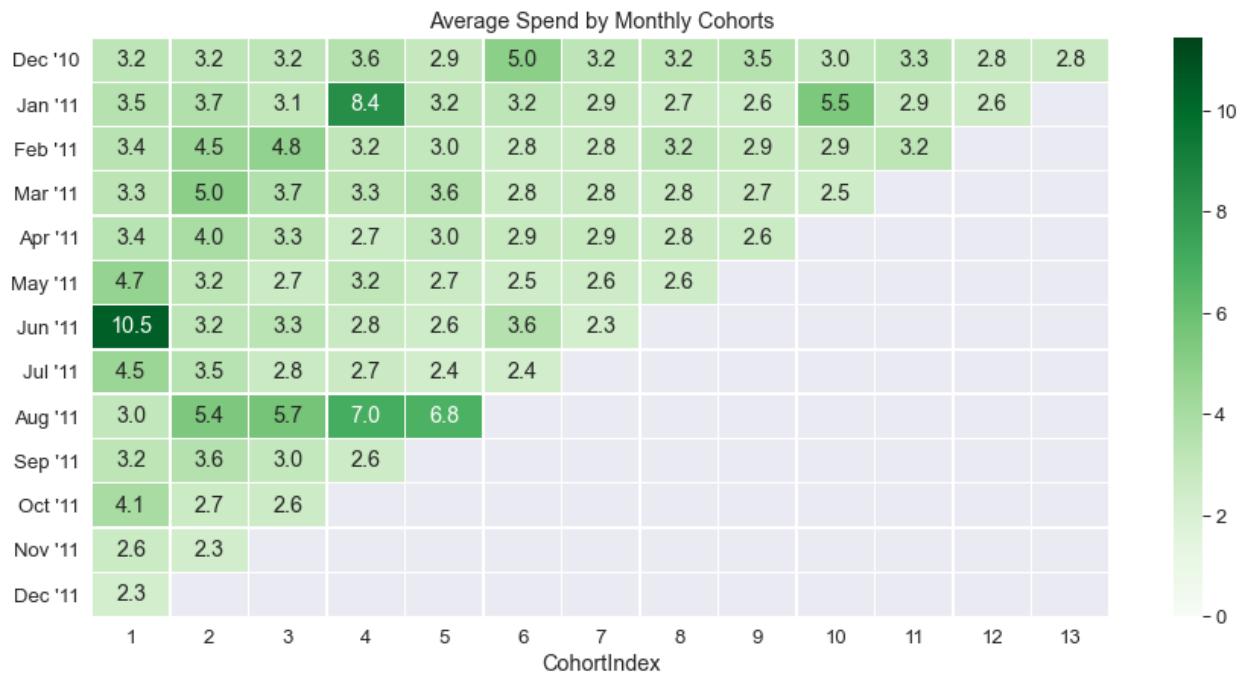
Min & max in each cohort is

```
(CohortIndex
1      1.000000
2      0.133956
3      0.130682
4      0.120805
5      0.137725
6      0.115183
7      0.102128
8      0.103943
9      0.083612
10     0.088636
11     0.092105
12     0.149644
13     0.274262
dtype: float64,
CohortIndex
1      1.000000
2      0.381857
3      0.334388
4      0.387131
5      0.359705
6      0.396624
7      0.379747
8      0.354430
9      0.354430
10     0.394515
11     0.373418
12     0.500000
13     0.274262
```

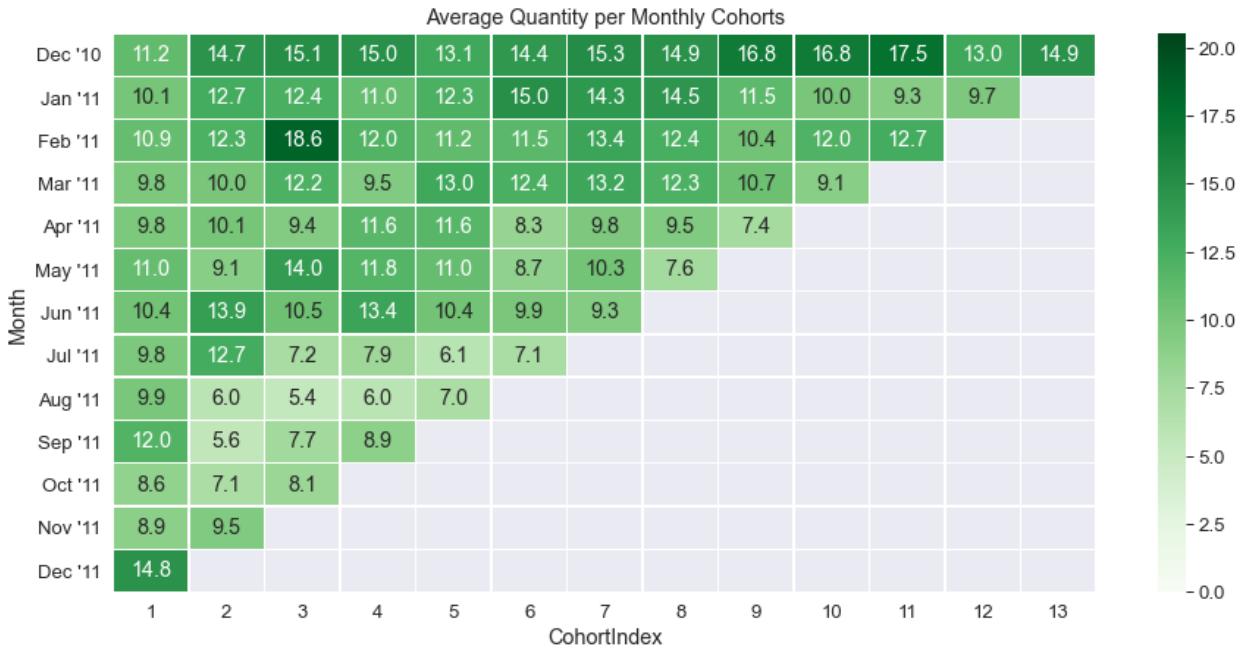
#retention rate or Active Customers



#average price per cohort



Avg Quantity by purchased per month



Week 2

Data Modeling :

1. Build a RFM (Recency Frequency Monetary) model. Recency means the number of days since a customer made the last purchase. Frequency is the number of purchase in a given period. It could be 3 months, 6 months or 1 year. Monetary is the total amount of money a customer spent in that given period. Therefore, big spenders will be differentiated among other customers such as MVP (Minimum Viable Product) or VIP.

2. Calculate RFM metrics.

3. Build RFM Segments. Give recency, frequency, and monetary scores individually by dividing them into quartiles.

b1. Combine three ratings to get a RFM segment (as strings).

b2. Get the RFM score by adding up the three ratings.

b3. Analyze the RFM segments by summarizing them and comment on the findings.

Note: Rate "recency" for customer who has been active more recently higher than the less recent customer, because each company wants its customers to be recent.

Note: Rate "frequency" and "monetary" higher, because the company wants the customer to visit more often and spend more money

Building RFM

1. Build a RFM (Recency Frequency Monetary) model. Recency means the number of days since a customer made the last purchase. Frequency is the number of purchase in a given period. It could be 3 months, 6 months or 1 year. Monetary is the total amount of money a customer spent in that given

period. Therefore, big spenders will be differentiated among other customers such as MVP (Minimum Viable Product) or VIP.

RFM Analysis

RFM analysis is a customer segmentation technique that uses past purchase behavior to divide customers into groups. RFM helps divide customers into various categories or clusters to identify customers who are more likely to respond to promotions and also for future personalization services.

2. Calculate RFM metrics.

Recency (R): Time since last purchase

Frequency (F): Total number of purchases

Monetary (M): Total purchase value

For RFM need to divide customers into four equal groups according to the distribution of values for recency, frequency, and monetary value. Four equal groups across three variables create 64 ($4 \times 4 \times 4$) different customer segments.

For example:

Customer with most recent purchase (R=4),

Customer with most quantity (F=4),

Customer who spent the most (M=4)

This customer belongs to RFM segment 4-4-4 (Best Customers), (R=4, F=4, M=4)

3. Build RFM Segments. Give recency, frequency, and monetary scores individually by dividing them into quartiles.

b1. Combine three ratings to get a RFM segment (as strings).

b2. Get the RFM score by adding up the three ratings.

b3. Analyze the RFM segments by summarizing them and comment on the findings.

#For recency, need to get the date difference since the last purchase.

#For this using the last purchase date on the database as today's date

#last date available in our dataset

Drop Current to date that we took for calculation. That is no more required

#Now finding out Frequency - how often or how many a customer used the product of a company.

Doing the same on the Spending or Monetary

Monetary is the total amount of money a customer spent in that given period. Therefore, big spenders will be differentiated with other customers such as MVP or VIP.

#Also creating a separate table for Customers for RFM. First merging frequency with regency and then that table with monetary

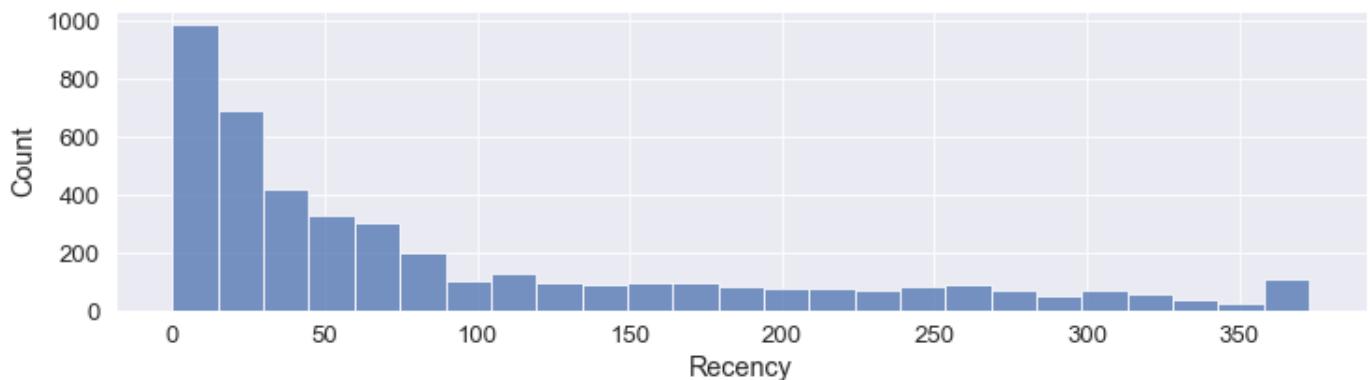
#RFM Table integrity Check is done - Passed

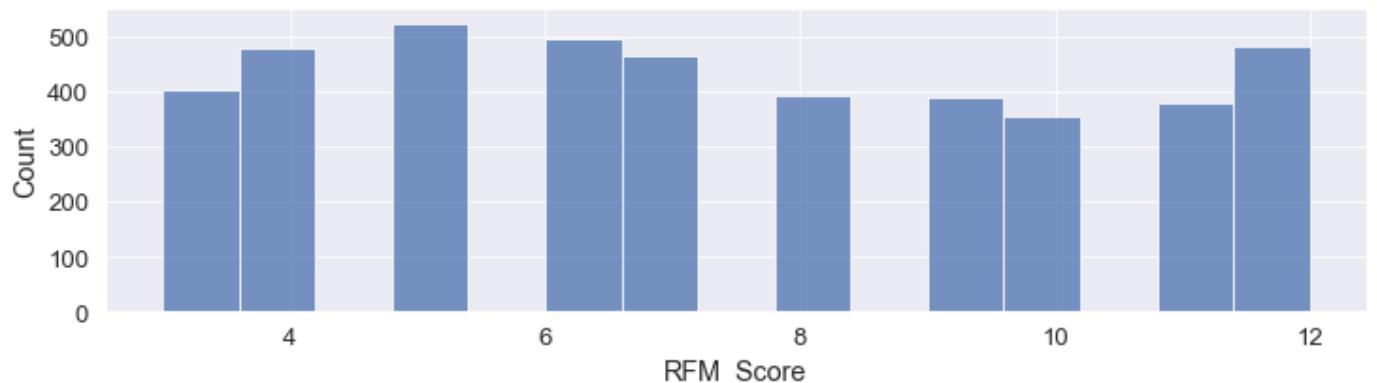
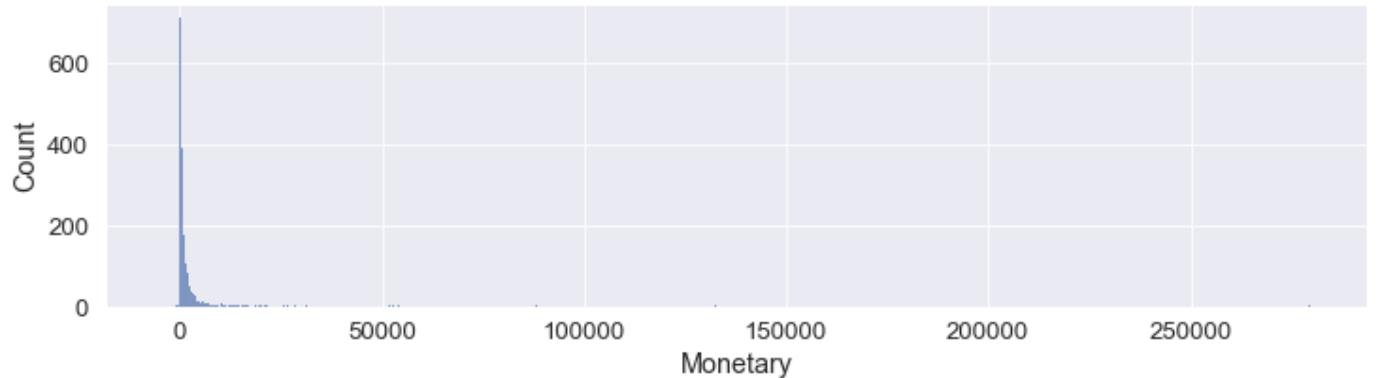
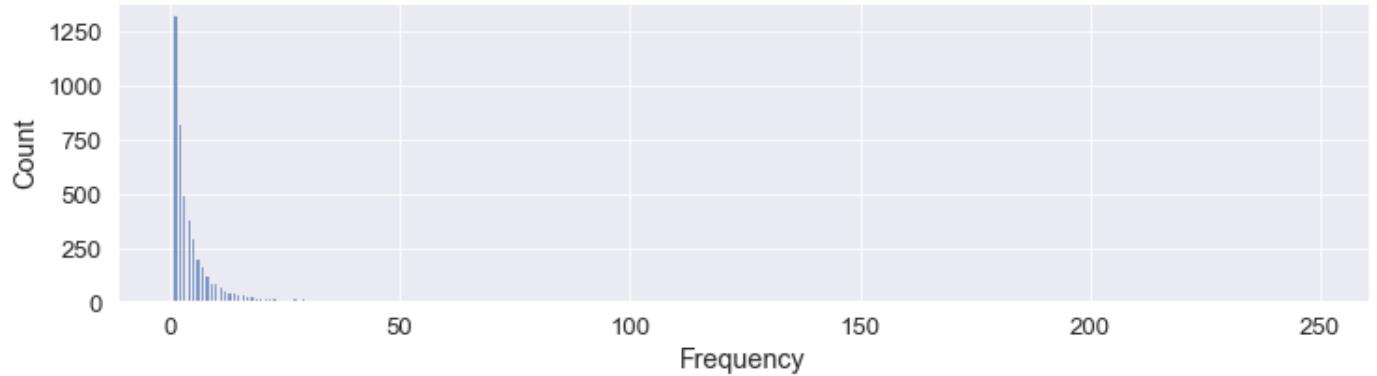
Then Assigning Score to the RFM before categorization. Will help later in plotting & Creating the Segments

```
'Best Customers':'444',    # Highest frequency as well as monetary value with least recency  
'Loyal Customers':'344',   # High frequency as well as monetary value with good recency  
'Big Spenders':'334',     # High monetary value but good recency and frequency values  
'Almost Lost':'244',      # Customer's shopping less often now who used to shop a lot  
'Lost Customers':'144',    # Customer's shopped long ago who used to shop a lot.  
'Recent Customers':'443',  # Customer's who recently started shopping a lot but with less monetary value  
'Lost Cheap Customers':'122', # Customer's shopped long ago but with less frequency and monetary value  
'No Harm to Lose Cheap Customers':'211' # Customer's shopped sometime back ago but with less frequency and monetary value
```

Swap the key and value of dictionary. So that Lookup is from value to Customer type and not vice-versa

Segment	#	Reference
Almost Lost	90	
Best Customers	482	These need Promotion materials and other engagement
Big Spenders	55	These need exclusive product - high end newsletters
Lost Cheap Customers	151	They may come back but focus spending is not recommended
Lost Customers	13	No effort to win them back
Loyal Customers	225	They need focused product list and AI built-in recommendation
No Harm to Lose Cheap Customers	177	Won't spend of this category
Recent Customers	99	Need to keep their interest alive - promote newer things on their spending using product recommendation





	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFM_Score
count	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000
mean	15299.677722	91.581199	5.075480	1893.531433	2.510979	2.349039	2.500000	7.360018
std	1722.390705	100.772139	9.338754	8218.696204	1.117084	1.151264	1.118162	2.872703
min	12346.000000	0.000000	1.000000	-4287.630000	1.000000	1.000000	1.000000	3.000000

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFM_Score
25%	13812.750000	16.000000	1.000000	291.795000	2.000000	1.000000	1.750000	5.000000
50%	15300.500000	50.000000	3.000000	644.070000	3.000000	2.000000	2.500000	7.000000
75%	16778.250000	143.000000	5.000000	1608.335000	4.000000	3.000000	3.250000	10.000000
max	18287.000000	373.000000	248.000000	279489.020000	4.000000	4.000000	4.000000	12.000000

Week 3

Data Modeling :

1. Create clusters using k-means clustering algorithm.

a. Prepare the data for the algorithm. If the data is asymmetrically distributed, manage the skewness with appropriate transformation. Standardize the data.

K-Means Clustering

From the above plots and rfm_table, we see that data is highly skewed.

It needs to be transformed and scale the data first because K-Means assumes that the variables should have a symmetric distributions(not skewed) and they should have same average values as well as same variance. Also, noticed, -ve value in monetary.

minimum range of value starts from 1 otherwise log transformation may lead to errors in graph plotting as well as K-Means clustering. After that we will utilize log transformation and scaling to make data available for K-Means clustering.

The k-means algorithm is an unsupervised clustering algorithm. It takes a bunch of unlabeled points and tries to group them into "k" number of clusters.

It is unsupervised because the points have no external classification.

Step 0: Preparing the data; scaling and removal of -ve values

Step 1: Determine K value by Elbow method and specify the number of clusters K

Step 2: Randomly assign each data point to a cluster

Step 3: Determine the cluster centroid coordinates

Step 4: Determine the distances of each data point to the centroids and re-assign each point to the closest cluster centroid based upon minimum distance

Step 5: Calculate cluster centroids again

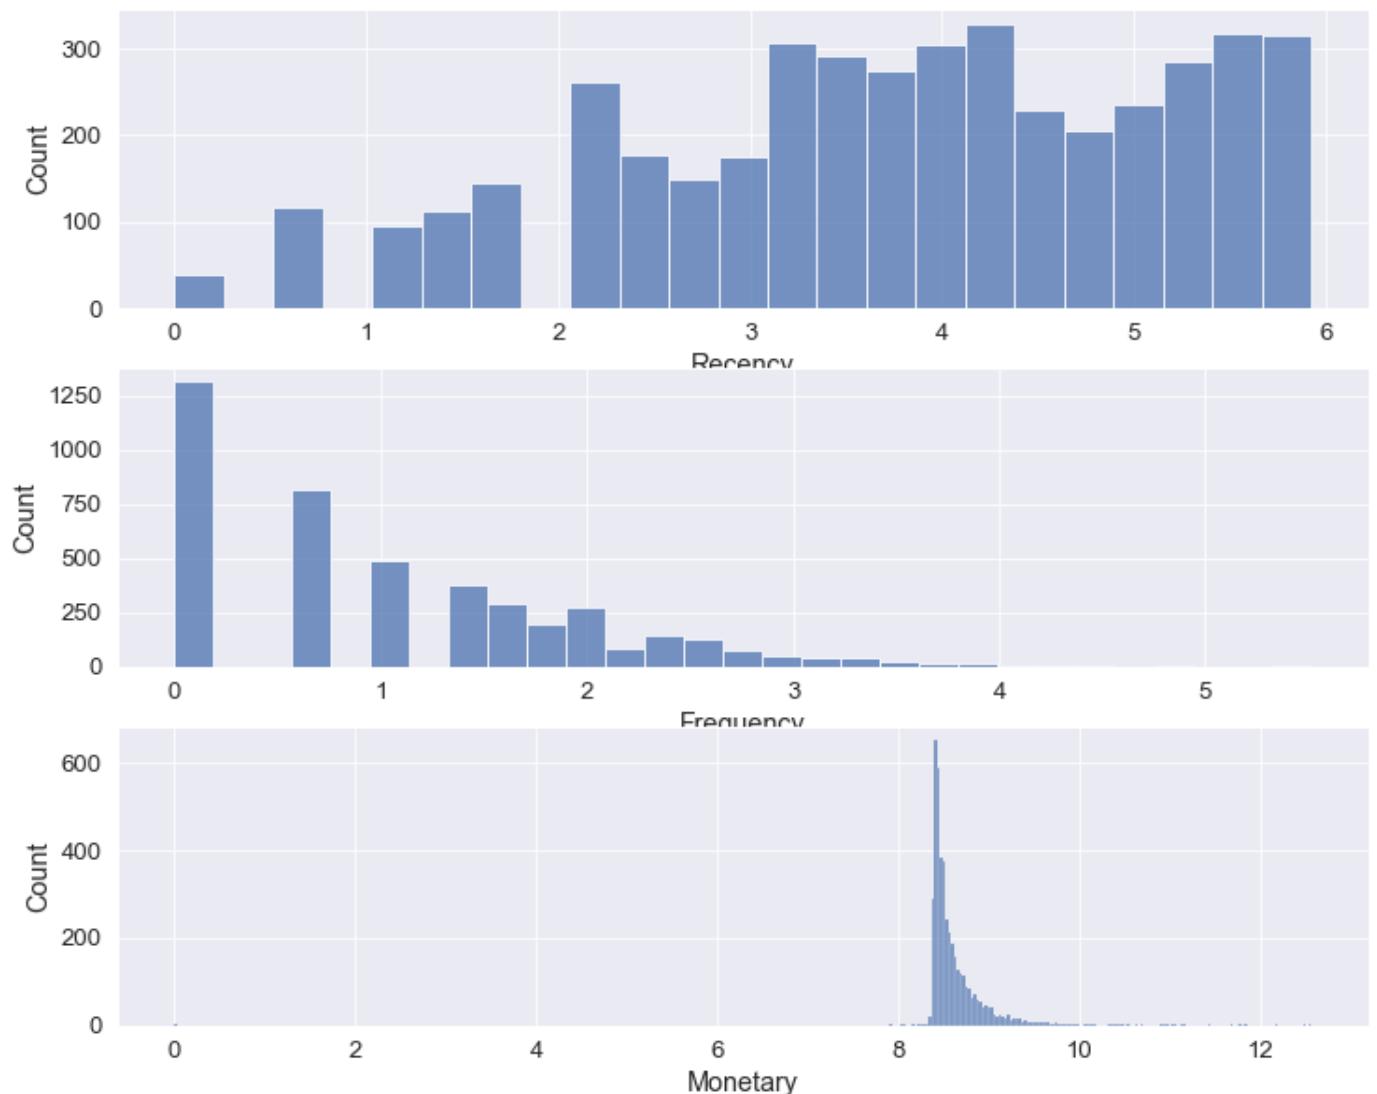
Step 6: Repeat steps 4 and 5 until we reach global optima where no improvements are possible and no switching of data points from one cluster to other.

Create a copy of rfm table for scaled calculation

Since it is unsupervised learning, we do not need to define the Segment & RFM_Score. We need the raw 3 components to find the clusters. Later we would add it in main table to see which cluster the customer belongs to.

Separating the three main inputs for K-Clustering and scale it

Recency	Frequency	Monetary	
0	5.786897	0.693147	8.363723
1	1.098612	1.945910	9.059358
2	4.330733	1.386294	8.713725
3	2.944439	0.000000	8.707182
4	5.739793	0.000000	8.438806



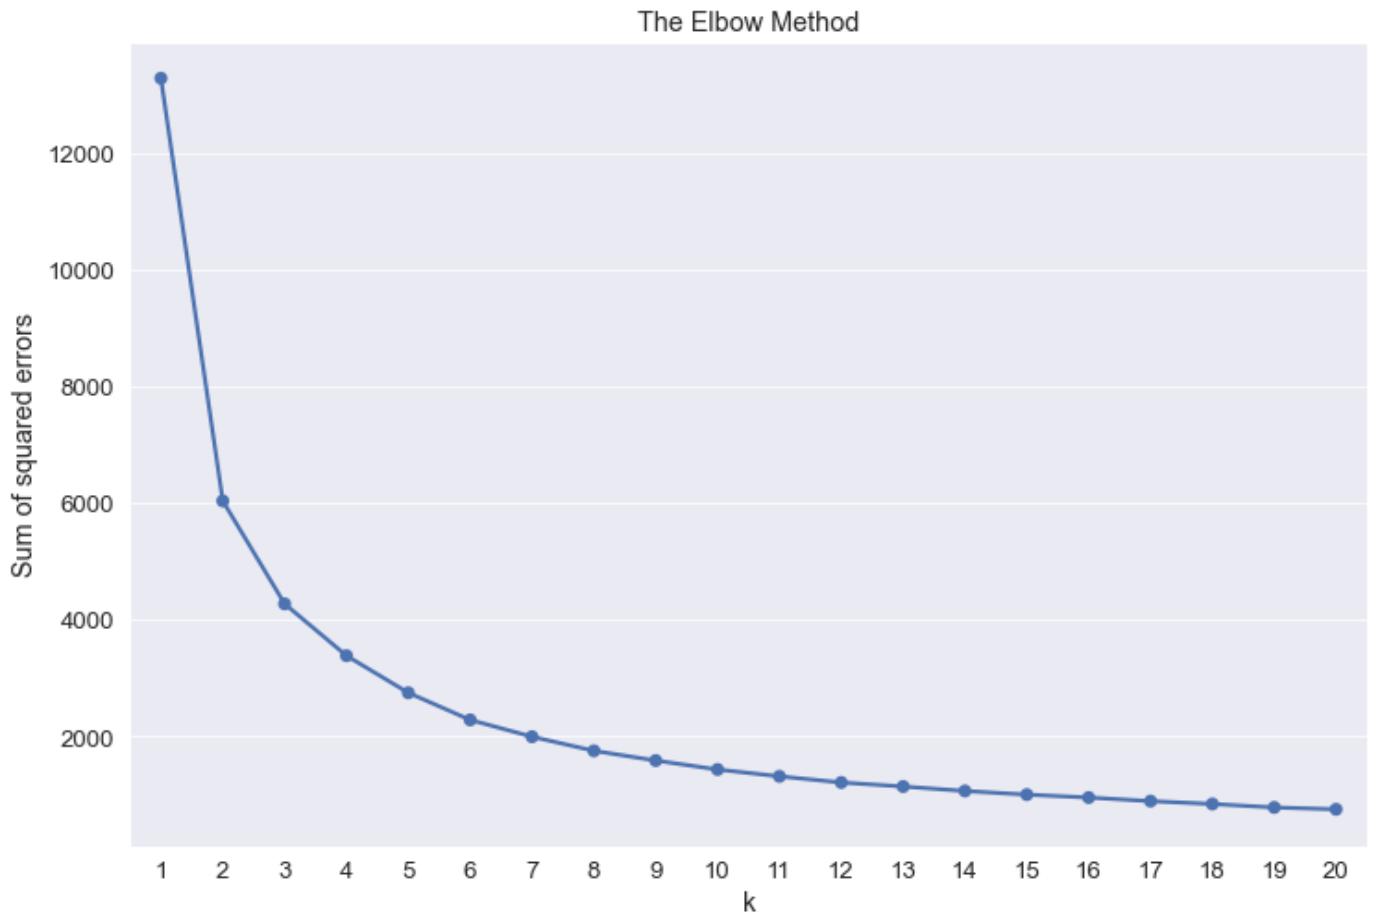
b. Decide the optimum number of clusters to be formed.

Finding out the optimum value of the clusters using elbow method and using the feature in Kmean called `inertia_`

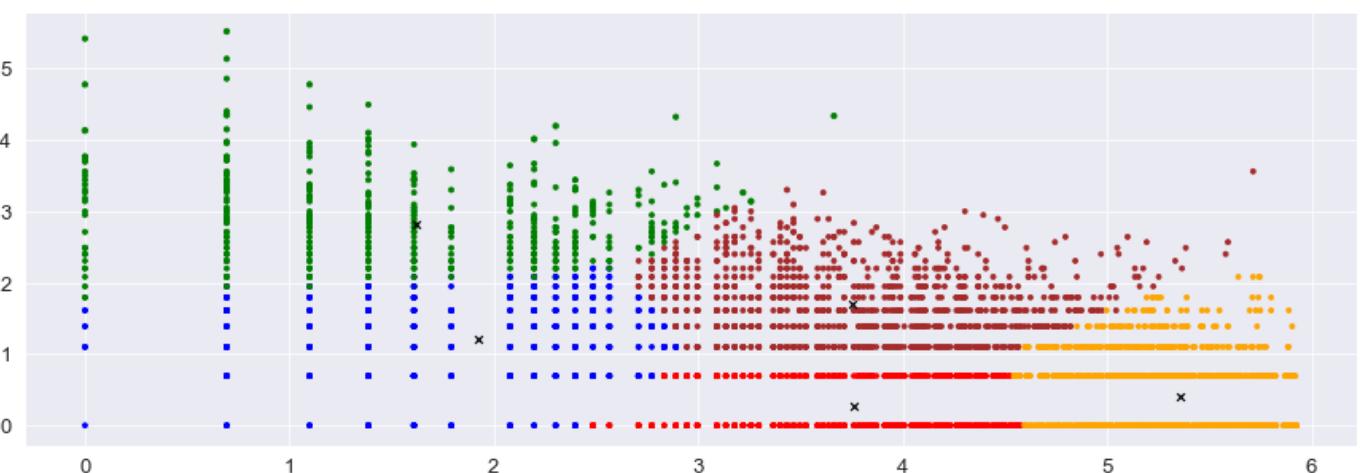
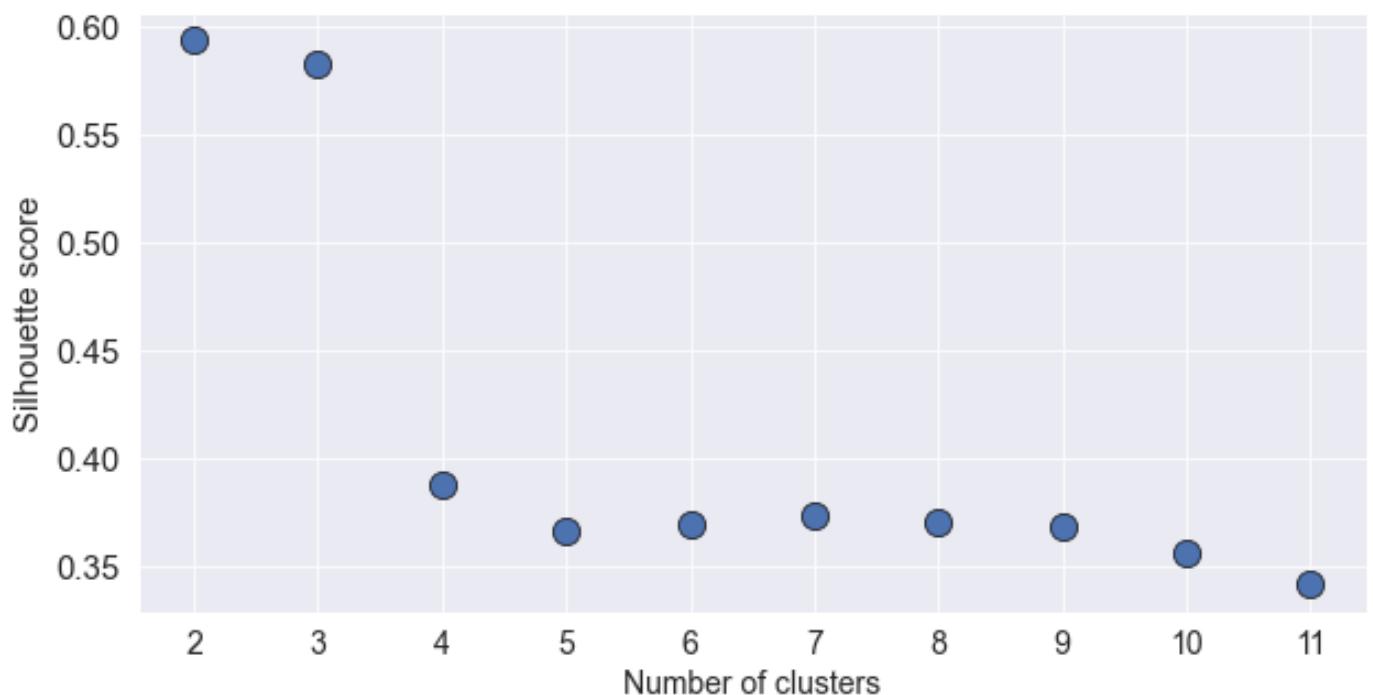
Inertia measures how well a dataset was clustered by K-Means. It is calculated by measuring the distance between each data point and its centroid, squaring this distance, and summing these squares across one cluster.

A good model is one with low inertia AND a low number of clusters (K). However, this is a tradeoff because as K increases, inertia decreases.

In Figure below the slowdown occurs at 5 but sharp cut starts at 3. So, we take 5 or 3 as the number of cluster = $k = 5$



The silhouette coefficient method
for determining number of clusters



CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	Segment	RFMScore	Cluster
0	12346.0	326	2	4288.6	3	1	1	None	4	0

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	Segment	RFM_Score	Cluster
1	12347.0	3	7	8598.63	4	4	4	444	Best Customers	12	2
2	12348.0	76	4	6085.87	2	3	4	234	None	9	1
3	12349.0	19	1	6046.18	3	1	4	314	None	8	3
4	12350.0	311	1	4623.03	1	1	2	112	None	4	0
...
43 67	18280.0	278	1	4469.23	1	1	1	111	None	3	0
43 68	18281.0	181	1	4369.45	1	1	1	111	None	3	0
43 69	18282.0	8	3	4465.23	4	2	1	421	None	7	4
43 70	18283.0	4	16	6334.16	4	4	4	444	Best Customers	12	2
43 71	18287.0	43	3	6125.91	3	2	4	324	None	9	1

c. Analyze these clusters and comment on the results.

Analyze each Cluster

Cluster 0 & 4 does not matter for us. Their RFM Score avg is low, and as we see they do not fall under any specialized marketing plan. They are the ones, who have very low RF&M. Though their number is high. They may be chanced customer who happen to drop in by some add etc.

1 & 3 has few categories that we had defined. Still a lot of effort is not to be directed on this cluster. Normal exposure to brand is good enough.

2 is the category we should be focusing our attention to. Their RFM avg is 12

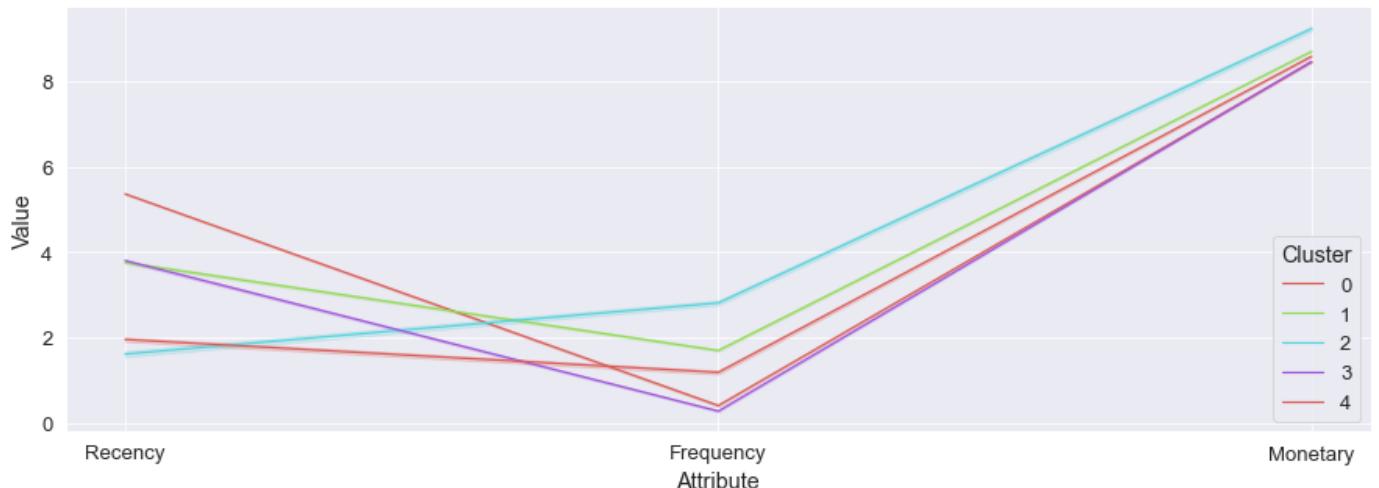
Snake plots

Market research technique to compare different segments

Visual representation of each segment's attributes

Plot each cluster's average normalized values of each attribute

To plot this we should have normalized data distribution and all the attributes in a single column. We will use pandas melt facility to achieve that



Also, Relative importance is found by dividing Cluster avg with Population Avg

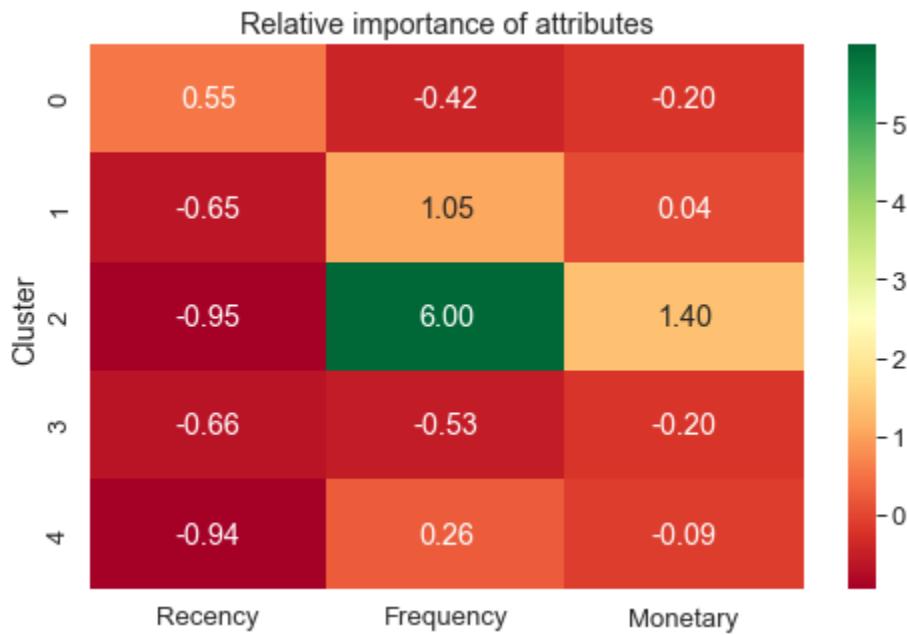
0, 1, 2, & 3 show similar spending.

4 & 3 have very infrequent. 0 shows average frequency of visit. - These are low spending but have a good volume of transaction. 0 cannot be ignored, thought a lot of effort or resources may not be given.

but highest is 2, followed by 1 - they are good spenders with good frequency. We would need to make sure we retain them.

Surprisingly 2 shows low recency. So these are planned buyers, not impulsive ones, while 0 are the impulsive ones.

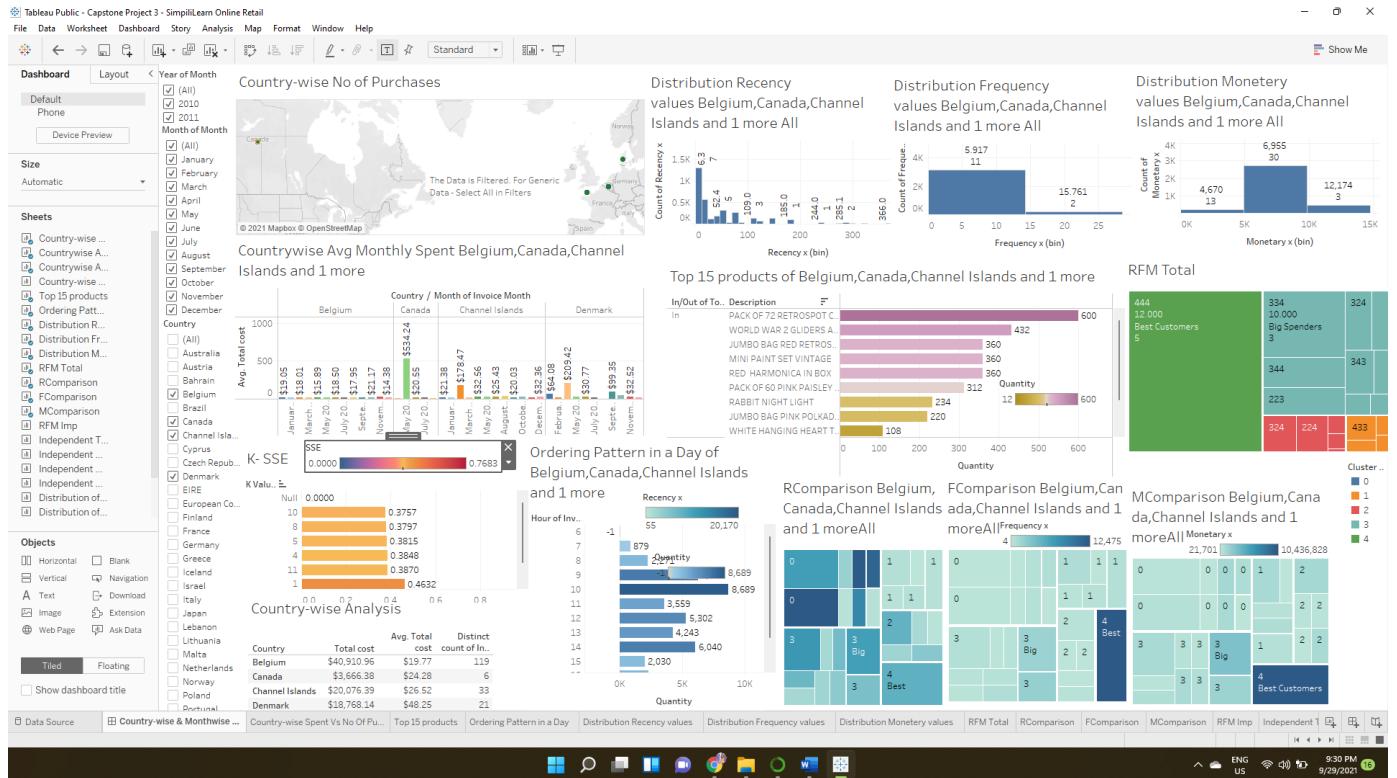
Now we get cluster average and population average to see the relative importance of each cluster
Then plot it in heat map



Week 4

Data Reporting:

1. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:
 - a. Country-wise analysis to demonstrate **average spend**. - Use a bar chart to show the monthly figures
 - b. Bar graph **of top 15 products** which are mostly ordered by the users to show the number of products sold
 - c. Bar graph to show **the count of orders vs. hours throughout the day**
 - d. Plot the **distribution of RFM values using histogram** and frequency charts
 - e. **Plot error (cost) vs. number of clusters selected**
 - f. **Visualize to compare the RFM values** of the clusters using heatmap



Code & Screen Shots



Capstone Project 3 - Retail .ipynb



Capstone Project
Code and Output.pdf

```

import pandas as pd

import numpy as np

import seaborn as sns

from operator import attrgetter

import matplotlib.colors as mcolors

import matplotlib.pyplot as plt

import datetime as dt

```

```
from scipy.stats import skewnorm
import scipy.stats as stats
from sklearn.preprocessing import LabelEncoder
import pylab as p
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

from sklearn.model_selection import learning_curve
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
df=pd.read_excel(r'D:\OneDrive\Studies\AI - ML\Capstone Project\OnlineRetail.xlsx',sheet_name='Online Retail')
df.head()
```

Out[213]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom
2	556690	23005	printing smudges/thrown away	-9600	2011-06-14 10:37:00	0.00	NaN	United Kingdom
3	556691	23005	printing smudges/thrown away	-9600	2011-06-14 10:37:00	0.00	NaN	United Kingdom
4	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom

Project Task: Week 1 Data Cleaning:

1. Perform a preliminary data inspection and data cleaning.

8 columns are available. The item related are - Stock Code (Quantifiable), & Description Sale Realted are - Invoice number and Invoice Date & Quantity & Unit Price Customer Realted are - Customer Id & Country

The main or basic inferred data are - Spending pattern, Spending categories, Customer Spending Behaviour

- a. Check for missing data and formulate an apt strategy to treat them.

In [4]: df

Out[4]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	
0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom	
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom	
2	556690	23005	printing smudges/thrown away	-9600	2011-06-14 10:37:00	0.00	NaN	United Kingdom	
3	556691	23005	printing smudges/thrown away	-9600	2011-06-14 10:37:00	0.00	NaN	United Kingdom	
4	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom	
...	
541904	573008	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	4800	2011-10-27 12:26:00	0.21	12901.0	United Kingdom	
541905	542504	37413		NaN	5568	2011-01-28 12:03:00	0.00	NaN	United Kingdom
541906	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540	2011-11-25 15:57:00	0.00	13256.0	United Kingdom	
541907	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	2011-01-18 10:01:00	1.04	12346.0	United Kingdom	
541908	581483	23843	PAPER CRAFT , LITTLE BIRDIE	80995	2011-12-09 09:15:00	2.08	16446.0	United Kingdom	

541909 rows x 8 columns

```
In [5]: df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   InvoiceNo    541909 non-null   object 
 1   StockCode    541909 non-null   object 
 2   Description  540455 non-null   object 
 3   Quantity     541909 non-null   int64  
 4   InvoiceDate  541909 non-null   datetime64[ns]
 5   UnitPrice    541909 non-null   float64
 6   CustomerID   406829 non-null   float64
 7   Country      541909 non-null   object 
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
```

Column - Description, & CustomerID, have some Null Values in them

Description Column is No Value add so there is no problem in dropping it. Since Customer Id is the main identifying element, Unique Identifying Entity, it's absence would be difficult to fill through Unlike cost or Sale Unit, where we could use average, we cannot use any other means to treat this except deletion.

Identifying & Removing Null Values. Starting from the column that has max null values. If it clears the null values in other columns, we would not need to repeat the activity with other columns

```
In [6]: df.isnull().sum()
Out[6]: InvoiceNo      0
StockCode      0
Description    1454
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID    135080
Country        0
dtype: int64
```

```
In [7]: df.dropna(subset=['CustomerID'], inplace=True)
```

```
In [8]: df.isnull().sum()
```

```
Out[8]: InvoiceNo      0
StockCode      0
Description     0
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID    0
Country        0
dtype: int64
```

b. Identify & Remove duplicate data records.

```
In [9]: # occurrence based on all columns
duplicate = df[df.duplicated(subset=None, keep='first')]
duplicate
```

```
Out[9]:
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
44	C570556	20971	PINK BLUE FELT CRAFT TRINKET BOX	-1296	2011-10-11 11:10:00	1.06	16029.0	United Kingdom
349	C570556	22568	FELTCRAFT CUSHION OWL	-144	2011-10-11 11:10:00	3.39	16029.0	United Kingdom
350	C570556	20969	RED FLORAL FELTCRAFT SHOULDER BAG	-144	2011-10-11 11:10:00	3.39	16029.0	United Kingdom
543	C568419	51014C	FEATHER PEN,COAL BLACK	-96	2011-09-27 11:16:00	0.39	13694.0	United Kingdom
1555	C575940	23309	SET OF 60 I LOVE LONDON CAKE CASES	-24	2011-11-13 11:38:00	0.55	17838.0	United Kingdom

534720	570242	21810	CHRISTMAS HANGING STAR WITH BELL	96	2011-10-09 15:40:00	0.39	16380.0	United Kingdom
536335	564327	85099B	JUMBO BAG RED RETROSPOT	100	2011-08-24 13:33:00	1.74	16029.0	United Kingdom
536373	565475	20725	LUNCH BAG RED RETROSPOT	100	2011-09-05 10:47:00	1.45	14156.0	EIRE
541294	548910	21982	PACK OF 12 SUKI TISSUES	432	2011-04-05 08:51:00	0.20	17940.0	United Kingdom
541855	561873	84568	GIRLS ALPHABET IRON ON PATCHES	1440	2011-07-31 11:48:00	0.17	13316.0	United Kingdom

5225 rows x 8 columns

```
In [10]: df=df.drop_duplicates()
df.duplicated().sum()
```

```
Out[10]: 0
```

```
In [11]: df.describe()
```

```
Out[11]:
```

	Quantity	UnitPrice	CustomerID
count	401604.000000	401604.000000	401604.000000
mean	12.183273	3.474064	15281.160818
std	250.283037	69.764035	1714.006089
min	-80995.000000	0.000000	12346.000000
25%	2.000000	1.250000	13939.000000
50%	5.000000	1.950000	15145.000000
75%	12.000000	3.750000	16784.000000
max	80995.000000	38970.000000	18287.000000

Get names of indexes for which Column Unit Price which has value negative value. Assuming that the shop keeper does not pay customer to purchase. There is one such instance, removing this data set as incorrect

c. Perform descriptive analytics on the given data.

However, not doing the same with Quantity. The assumption here is that these might be billed in previous cycle which is not included in this database and were returned to seller in this cycle. Approx ~2% was returned

```
In [12]: indexNames = df[df['UnitPrice'] < 0 ].index
indexNames
df.drop(indexNames , inplace=True)
```

```
In [13]: #Counting % of Returns  
indexNames = df[df['Quantity'] < 0 ].index  
PercentageReturn = (indexNames.size / df['Quantity'].size)*100  
PercentageReturn
```

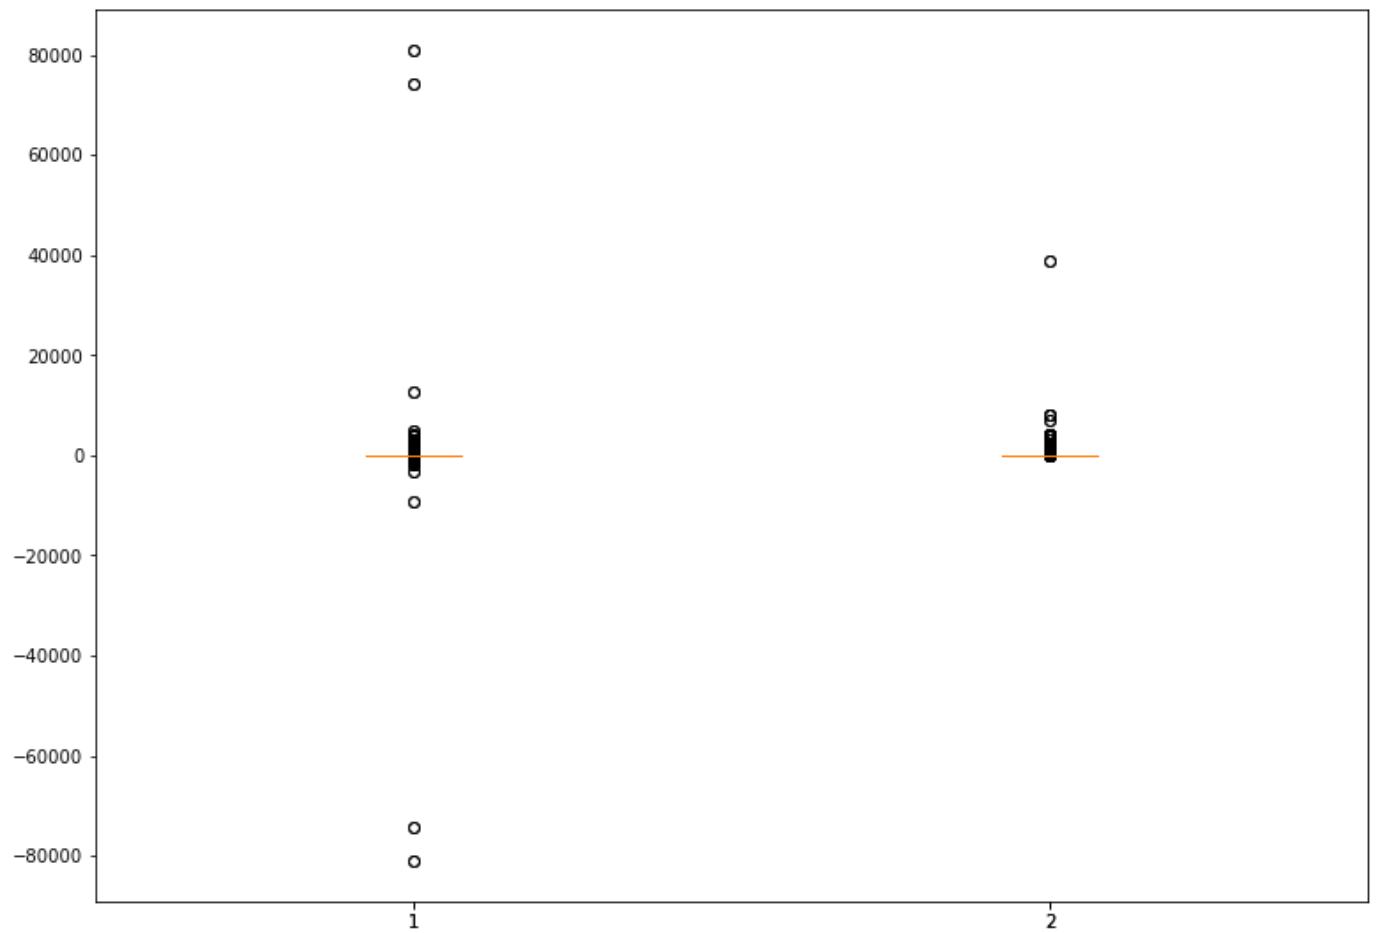
```
Out[13]: 2.209141343213713
```

```
In [14]: df.describe()
```

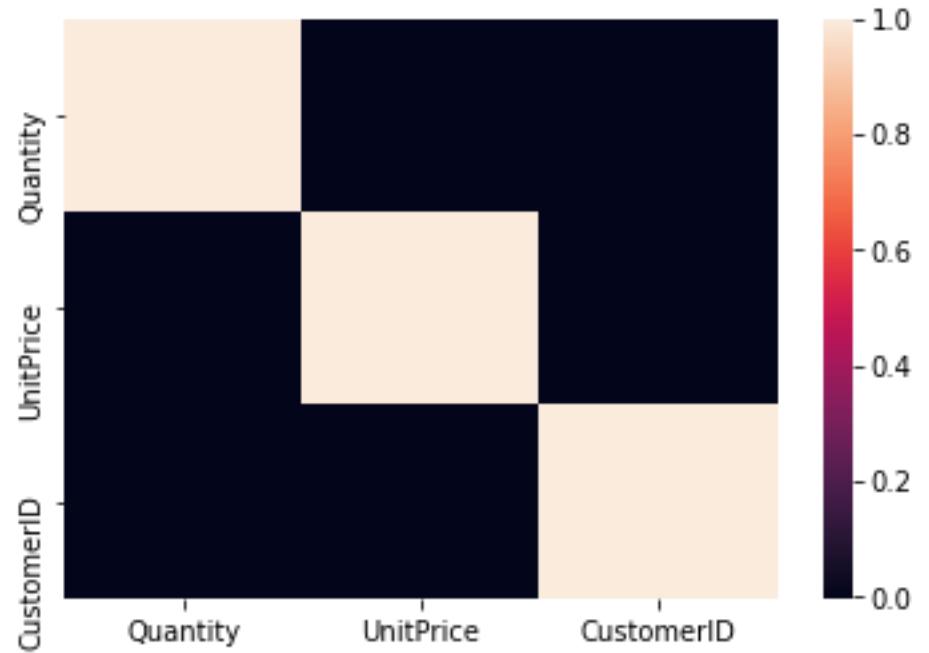
```
Out[14]:
```

	Quantity	UnitPrice	CustomerID
count	401604.000000	401604.000000	401604.000000
mean	12.183273	3.474064	15281.160818
std	250.283037	69.764035	1714.006089
min	-80995.000000	0.000000	12346.000000
25%	2.000000	1.250000	13939.000000
50%	5.000000	1.950000	15145.000000
75%	12.000000	3.750000	16784.000000
max	80995.000000	38970.000000	18287.000000

```
In [15]: fig = plt.figure(figsize =(10, 7))  
data = [df['Quantity'], df['UnitPrice']]  
fig = plt.figure(figsize =(10, 7))  
  
# Creating axes instance  
ax = fig.add_axes([0, 0, 1, 1])  
  
# Creating plot  
bp = ax.boxplot(data)  
  
# Creating plot  
plt.boxplot(data)  
  
# show plot  
plt.show()
```



```
corr = df.corr()  
print(corr)  
sns.heatmap(corr,  
            xticklabels=corr.columns,  
            yticklabels=corr.columns)
```



#Unique Countries

```
pd.DataFrame(df['Country'].unique())
```

Out[17]:

0	
0	United Kingdom
1	Japan
2	Netherlands
3	EIRE
4	Spain
5	Germany
6	France
7	Sweden
8	Switzerland
9	Australia
10	Austria
11	USA
12	Cyprus
13	Israel
14	Finland
15	Denmark
16	Czech Republic
17	Portugal
18	Norway
19	Italy
20	Belgium
21	Poland
22	Saudi Arabia
23	Malta
24	Channel Islands
25	European Community
26	Singapore

```
20      Singapore
27      Greece
28      Canada
29      Unspecified
30  United Arab Emirates
31      RSA
32      Lebanon
33      Brazil
34      Bahrain
35      Iceland
36      Lithuania
```

```
In [18]: #Unique Customers
UniqueCustomer = pd.DataFrame(df['CustomerID'].unique())
#Unique Customers are 4372 out
UniqueCustomer
```

```
Out[18]:
```

```
0
0  16446.0
1  12346.0
2  15838.0
3  15749.0
4  16938.0
...
4367 16754.0
4368 15195.0
4369 15118.0
4370 13135.0
4371 13256.0
```

```
4372 rows × 1 columns
```

```
In [19]: count = pd.DataFrame(df['CustomerID'])
x= pd.DataFrame(df['CustomerID'].value_counts())
x.rename({'CustomerID': 'Freq'}, axis='columns', inplace = True)
x
```

Out[19]:

	Freq
17841.0	7812
14911.0	5898
14096.0	5128
12748.0	4459
14606.0	2759
...	...
17948.0	1
15590.0	1
16061.0	1
18174.0	1
18068.0	1

4372 rows × 1 columns

In [20]: # column Freq has value = 1
SBuyer = x[x.Freq == 1]
SBuyer
#Only 79 buyer purchased once

Out[20]:

	Freq
16093.0	1
15524.0	1
13154.0	1
15562.0	1
16995.0	1
...	...
17948.0	1
15590.0	1
16061.0	1
18174.0	1
18068.0	1

79 rows × 1 columns

SBuyer['CustomerID']= SBuyer.index

SBuyer

	Freq	CustomerID
16093.0	1	16093.0
15524.0	1	15524.0
13154.0	1	13154.0
15562.0	1	15562.0
16995.0	1	16995.0
...
17948.0	1	17948.0
15590.0	1	15590.0
16061.0	1	16061.0
18174.0	1	18174.0
18068.0	1	18068.0

79 rows × 2 columns

```
In [22]: #Percentage of single purchaser is,  
len(SBuyer)/len(UniqueCustomer)*100
```

```
Out[22]: 1.8069533394327537
```

```
In [23]: x['CustomerID']= x.index  
# get names of indexes for which  
# column Freq has value = 1  
index_names = x[ x['Freq'] <= 1 ].index  
  
# drop these row indexes  
# from DataFrame  
x.drop(index_names, inplace = True)
```

```
In [24]: #Repeat Customers are 4293  
x
```

```
Out[24]:
```

	Freq	CustomerID
17841.0	7812	17841.0
14911.0	5898	14911.0
14096.0	5128	14096.0
12748.0	4459	12748.0
14606.0	2759	14606.0

x

Out[24]:

	Freq	CustomerID
17841.0	7812	17841.0
14911.0	5898	14911.0
14096.0	5128	14096.0
12748.0	4459	12748.0
14606.0	2759	14606.0
...
15423.0	2	15423.0
14642.0	2	14642.0
13130.0	2	13130.0
13298.0	2	13298.0
14821.0	2	14821.0

4293 rows × 2 columns

In [25]: df

Out[25]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom
4	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom
10	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	15749.0	United Kingdom
20	C550456	21175	GIN + TONIC DIET METAL SIGN	-2000	2011-04-18 13:08:00	1.85	15749.0	United Kingdom
...
541903	554868	22197	SMALL POPCORN HOLDER	4300	2011-05-27 10:52:00	0.72	13135.0	United Kingdom
541904	573008	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	4800	2011-10-27 12:26:00	0.21	12901.0	United Kingdom
541906	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540	2011-11-25 15:57:00	0.00	13256.0	United Kingdom
541907	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	2011-01-18 10:01:00	1.04	12346.0	United Kingdom
541908	581483	23843	PAPER CRAFT , LITTLE BIRDIE	80995	2011-12-09 09:15:00	2.08	16446.0	United Kingdom

401604 rows × 8 columns

```
In [27]: corr = df1.corr()
print(corr)
sns.heatmap(corr,
            xticklabels=corr.columns,
            yticklabels=corr.columns)

Quantity      Quantity   UnitPrice  CustomerID  Freq
Quantity  1.000000 -0.048550 -0.201286  NaN
UnitPrice -0.048550  1.000000  0.082662  NaN
CustomerID -0.201286  0.082662  1.000000  NaN
Freq       NaN        NaN        NaN        NaN
```

Out[27]: <AxesSubplot:>



```
In [28]: #Countries from where Buyers are from
pd.DataFrame(df['Country'].unique())
```

Out[28]:

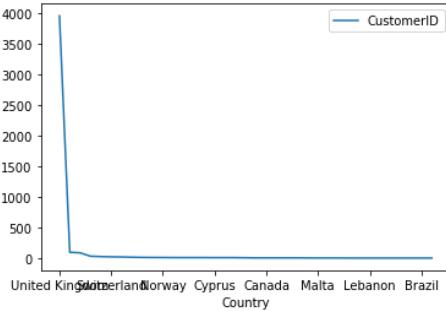
	0
0	United Kingdom
1	Japan
2	Netherlands
3	EIRE
4	Spain
5	Germany
6	France
7	Sweden
8	Switzerland
9	Australia

```
df.Country.value_counts(normalize=True).head(10).mul(100).round(1).astype(str) + '%'
```

```
Out[30]: United Kingdom    88.8%
Germany        2.4%
France         2.1%
EIRE           1.9%
Spain          0.6%
Netherlands    0.6%
Belgium        0.5%
Switzerland    0.5%
Portugal       0.4%
Australia      0.3%
Name: Country, dtype: object
```

```
In [31]: customercountrywise.plot()
#UK is the major player from where Buyers are from, which is 89%
```

```
Out[31]: <AxesSubplot:xlabel='Country'>
```



```
In [32]: #Customer who purchased more than 1 items
n_orders = df.groupby(['CustomerID'])['InvoiceNo'].nunique()
mult_orders_perc = np.sum(n_orders > 1) / df['CustomerID'].nunique()
print(f'{100 * mult_orders_perc:.2f}% of customers ordered more than one item.')
```

```
69.97% of customers ordered more than one item.
```

```
In [33]: # Creating histogram
fig, ax = plt.subplots(figsize =(10, 7))
ax.hist(n_orders, bins = [0, 25, 50, 75, 100])

# Show plot
plt.show()
ax.set(title='Distribution of number of orders per customer',
       xlabel='# of orders',
```

```
# Creating histogram
```

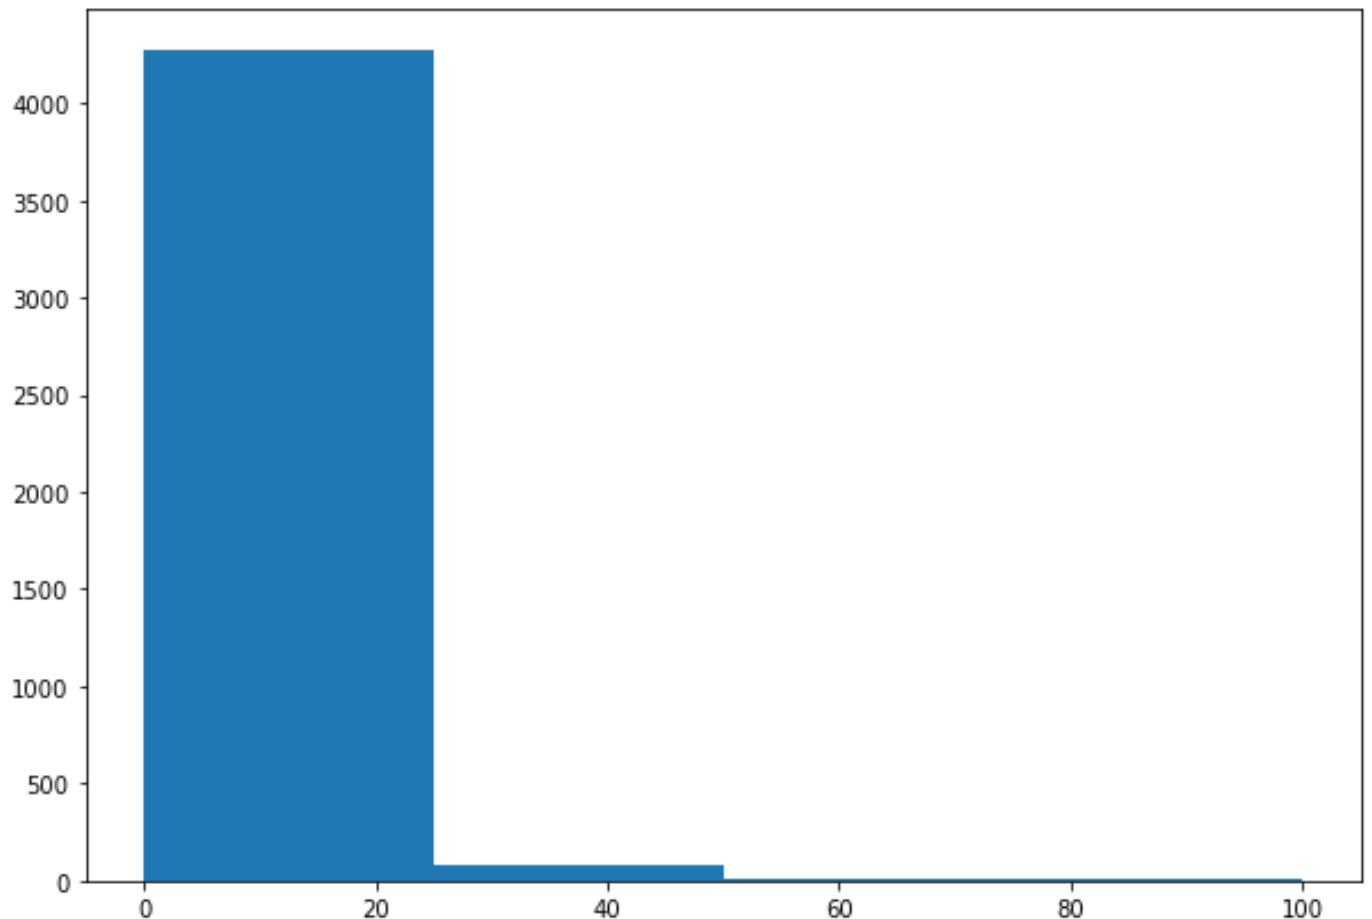
```
fig, ax = plt.subplots(figsize =(10, 7))

ax.hist(n_orders, bins = [0, 25, 50, 75, 100])
```

```
# Show plot
```

```
plt.show()

ax.set(title='Distribution of number of orders per customer',
       xlabel='# of orders',
       ylabel='# of customers');
```



```
# Check the oldest and latest date in the dataset.
```

```
print(f'Oldest date is - {df.InvoiceDate.min()}\n')
```

```
print(f'Latest date is - {df.InvoiceDate.max()}')
```

```
Oldest date is - 2010-12-01 08:26:00
```

```
Latest date is - 2011-12-09 12:50:00
```

#Monthly Sales

```
# importing DateTime module to convert extracted dates
```

```
def get_month(x):
```

```
    return dt.datetime(x.year, x.month, 1)
```

```
df['InvoiceMonth'] = df['InvoiceDate'].apply(get_month)
```

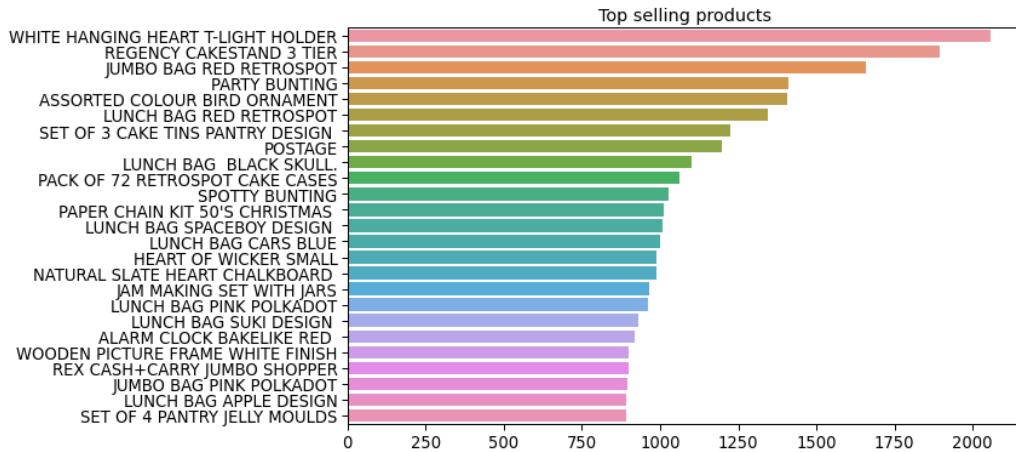
```
df.head()
```

Out[35]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth
0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom	2011-12-01
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom	2011-01-01
4	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom	2010-12-01
10	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	15749.0	United Kingdom	2011-04-01
20	C550456	21175	GIN + TONIC DIET METAL SIGN	-2000	2011-04-18 13:08:00	1.85	15749.0	United Kingdom	2011-04-01

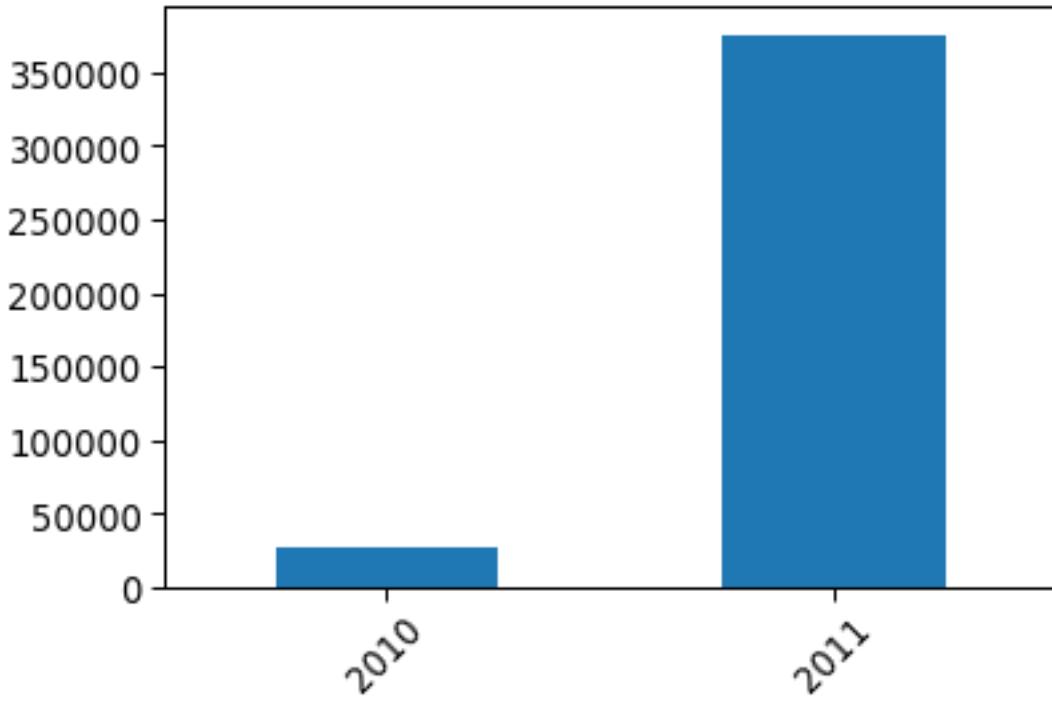
In [36]: # Top selling products

```
top_products = df['Description'].value_counts()[:25]
plt.figure(figsize=(10,6))
sns.set_context("paper", font_scale=1.5)
sns.barplot(y = top_products.index,
            x = top_products.values)
plt.title("Top selling products")
plt.show();
```



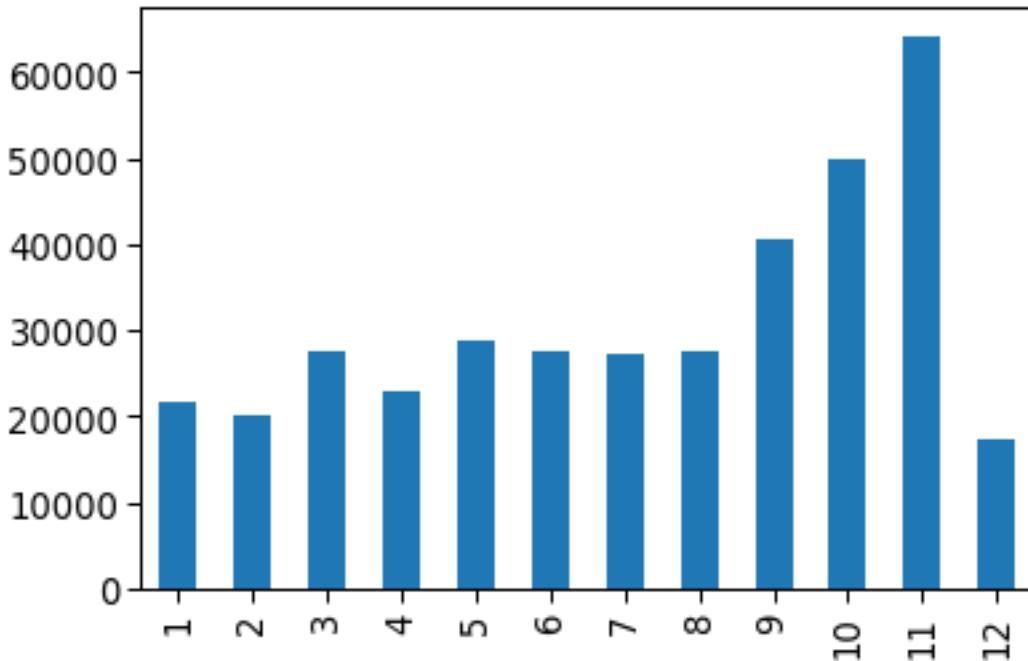
Count of transactions in different years

```
df.InvoiceDate.dt.year.value_counts(sort=False).plot(kind='bar', rot=45);
```



Most of the records belong to 2011. Now doing monthly break up. And we see that max transaction is in Nov & Oct. Could be Black Friday or Halloween. Dec is lesser does not indicating advance buying, because in this data, sales till 09th is considered.

```
df[df.InvoiceDate.dt.year==2011].InvoiceDate.dt.month.value_counts(sort=False).plot(kind='bar')
```



```
df['Total_cost'] = df['UnitPrice']*df['Quantity']
```

df

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	Total_cost
0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom	2011-12-01	-168469.6
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom	2011-01-01	-77183.6
4	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom	2010-12-01	-280.8
10	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	15749.0	United Kingdom	2011-04-01	-6539.4
20	C550456	21175	GIN + TONIC DIET METAL SIGN	-2000	2011-04-18 13:08:00	1.85	15749.0	United Kingdom	2011-04-01	-3700.0
...
541903	554868	22197	SMALL POPCORN HOLDER	4300	2011-05-27 10:52:00	0.72	13135.0	United Kingdom	2011-05-01	3096.0
541904	573008	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	4800	2011-10-27 12:26:00	0.21	12901.0	United Kingdom	2011-10-01	1008.0
541906	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540	2011-11-25 15:57:00	0.00	13256.0	United Kingdom	2011-11-01	0.0
541907	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	2011-01-18 10:01:00	1.04	12346.0	United Kingdom	2011-01-01	77183.6
541908	581483	23843	PAPER CRAFT , LITTLE BIRDIE	80995	2011-12-09 09:15:00	2.08	16446.0	United Kingdom	2011-12-01	168469.6

401604 rows × 10 columns

```
In [40]: monthly_gross = df[df.InvoiceDate.dt.year==2011].groupby(df.InvoiceDate.dt.month).Total_cost.sum()
plt.figure(figsize=(10,5))
sns.lineplot(y=monthly_gross.values,x=monthly_gross.index, marker='o');
plt.xticks(range(1,13))
plt.show();
```

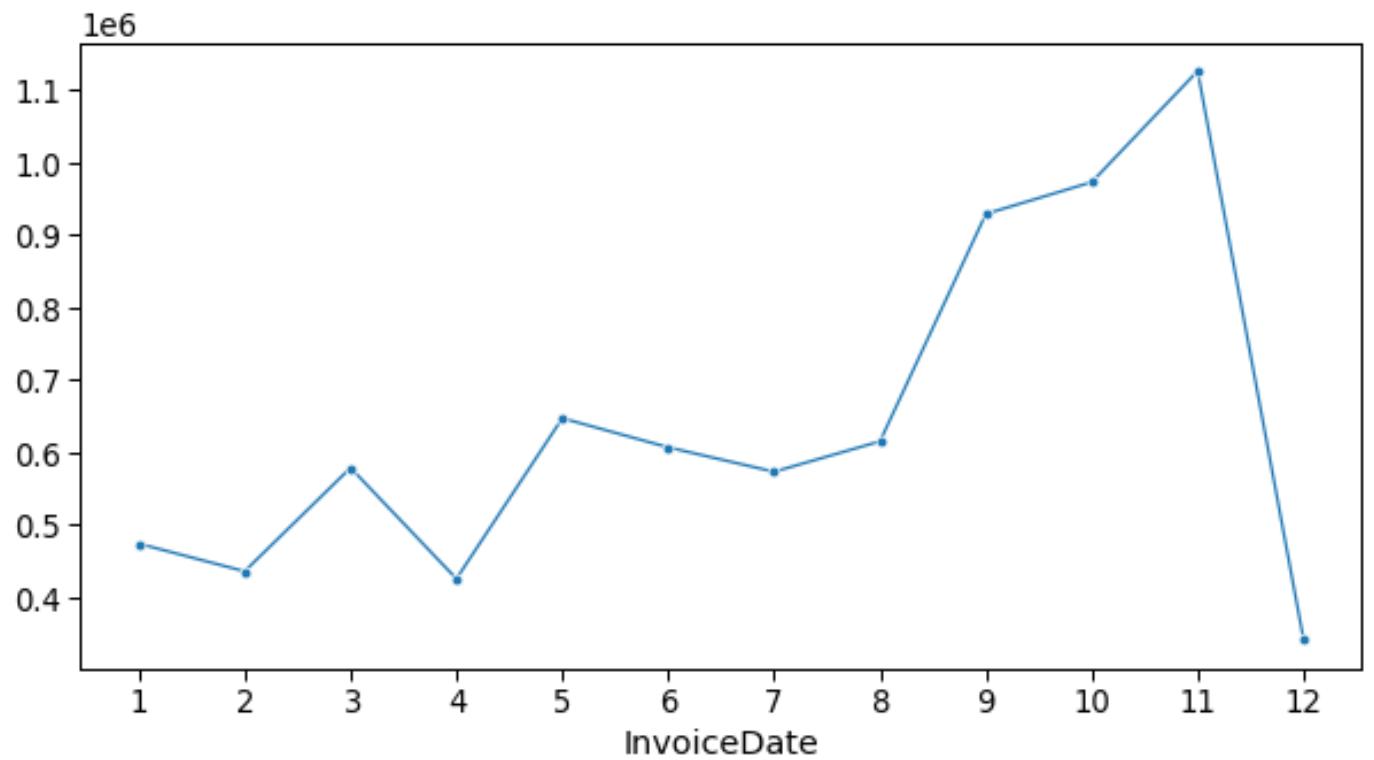
```
monthly_gross = df[df.InvoiceDate.dt.year==2011].groupby(df.InvoiceDate.dt.month).Total_cost.sum()

plt.figure(figsize=(10,5))

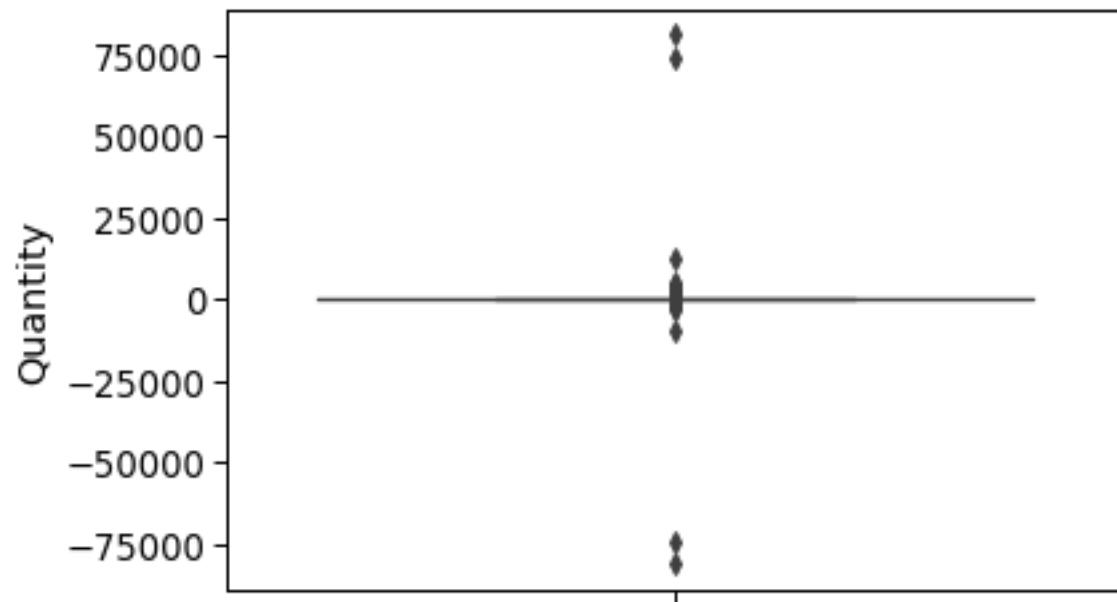
sns.lineplot(y=monthly_gross.values,x=monthly_gross.index, marker='o');

plt.xticks(range(1,13))

plt.show();
```



```
# Boxplot to for Quantity distribution  
sns.boxplot(y='Quantity', data=df, orient='h');
```



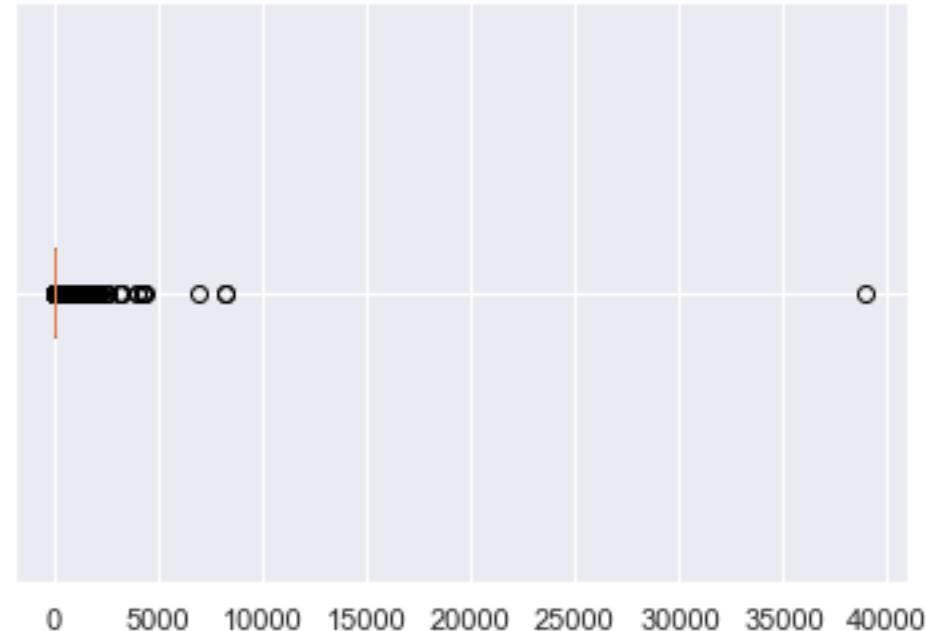
```
#Unit price distribution  
sns.set(style="darkgrid")
```

```
plt.boxplot(df['UnitPrice'], vert = 0)
```

```
#plt.tight_layout()
```

```
plt.show()
```

```
#sns.boxplot(y='UnitPrice', data=df)
```



Perform cohort analysis (a cohort is a group of subjects that share a defining characteristic). Observe how a cohort behaves across time and compare it to other cohorts. a. Create month cohorts and analyze active customers for each cohort. b. Analyze the retention rate of customers.

Time cohorts Time cohorts are customers who signed up for a product or service during a particular time frame. Analysing these cohorts shows the customers' behaviour depending on the time they started using the company's products or services. The time may be monthly or quarterly, even daily.

```
#Assigning Cohort to each group
```

```
group = df.groupby('CustomerID')['InvoiceMonth']
```

```
group.head()
```

```
0      2011-12-01  
1      2011-01-01  
4      2010-12-01  
10     2011-04-01  
20     2011-04-01  
..
```

```

541900 2011-07-01
541903 2011-05-01
541906 2011-11-01
541907 2011-01-01
541908 2011-12-01

```

Name: InvoiceMonth, Length: 21206, dtype: datetime64[ns]

```
df['Month'] = df.groupby('CustomerID')['InvoiceMonth'].transform('min')
```

df

Invoic eNo	StockC ode	Description	Quantit y	Invoice Date	UnitPr ice	Custom erID	Coun try	InvoiceM onth	Total_c ost	Mont h
0	C58148 4	23843 PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	1644 6.0	United Kingdom	2011-12-01	16846 9.6	2011-05-01
1	C54143 3	23166 MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	1234 6.0	United Kingdom	2011-01-01	77183 .6	2011-01-01
4	C53675 7	84347 ROTATING SILVER ANGEL ST LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	1583 8.0	United Kingdom	2010-12-01	-280.8	2010-12-01
10	C55045 6	21108 FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	1574 9.0	United Kingdom	2011-04-01	6539. 4	2011-01-01
20	C55045 6	21175 GIN + TONIC DIET METAL SIGN	-2000	2011-04-18 13:08:00	1.85	1574 9.0	United Kingdom	2011-04-01	3700. 0	2011-01-01

Invoic eNo	StockC ode	Description	Quantit y	Invoice Date	UnitPr ice	Custom erID	Coun try	InvoiceM onth	Total_c ost	Mont h
541903	554868	22197	SMALL POPCO RN HOLDE R	4300	2011- 05-27 10:52: 00	0.72	1313 5.0	United Kingdom	2011- 05-01	3096. 0
541904	573008	84077	WORLD WAR 2 GLIDER S ASSTD DESIGN S	4800	2011- 10-27 12:26: 00	0.21	1290 1.0	United Kingdom	2011- 10-01	1008. 0
541906	578841	84826	ASSTD DESIGN 3D PAPER STICKE RS	12540	2011- 11-25 15:57: 00	0.00	1325 6.0	United Kingdom	2011- 11-01	0.0
541907	541431	23166	MEDIU M CERAM IC TOP STORA GE JAR	74215	2011- 01-18 10:01: 00	1.04	1234 6.0	United Kingdom	2011- 01-01	77183 .6
541908	581483	23843	PAPER CRAFT , LITTLE BIRDIE	80995	2011- 12-09 09:15: 00	2.08	1644 6.0	United Kingdom	2011- 12-01	16846 9.6

401604 rows × 11 columns

#monthly cohorts based on the month each customer has made their first transaction.

```
def get_month(x):
```

```
    return dt.datetime(x.year,x.month,1)
```

```
# Create InvoiceMonth column
df['InvoiceMonth'] = df['InvoiceDate'].apply(get_month)

# Group by CustomerID and select the InvoiceMonth value
grouping = df.groupby('CustomerID')['InvoiceMonth']

# Assign a minimum InvoiceMonth value to the dataset
df['Month'] = grouping.transform('min')
```

Calculate time offset in months Calculating time offset for each transaction allows you to report the metrics for each cohort in a comparable fashion.

First, we will create some variables that capture the integer value of years and months for Invoice and Cohort Date

```
def get_date_int(df, column):
    year = df[column].dt.year
    month = df[column].dt.month
    return year, month
```

```
# Get the integers for date parts from the `InvoiceMonth` column
invoice_year, invoice_month = get_date_int(df,'InvoiceMonth')
```

```
# Get the integers for date parts from the `CohortMonth` column
cohort_year, cohort_month = get_date_int(df,'Month')
```

```
# Calculate difference in years
years_diff = invoice_year - cohort_year
```

```
# Calculate difference in months
months_diff = invoice_month - cohort_month
```

```

# Extract the difference in months from all previous values
df['CohortIndex'] = years_diff * 12 + months_diff + 1

#Sanity Check to see if the Cohort Index is of different number
df['CohortIndex']

0          8
1          1
4          1
10         4
20         4
..
541903     1
541904     8
541906     1
541907     1
541908     8
Name: CohortIndex, Length: 401604, dtype: int64

```

```
df.head()
```

	Invoic eNo	StockC ode	Descrip tion	Quant ity	Invoice Date	UnitP rice	Custome rID	Count ry	InvoiceM onth	Total_ cost	Mon th	CohortI ndex
0	C5814 84	23843	PAPER CRAFT , LITTLE BIRDIE	- 80995	2011- 12-09 09:27:00	2.08	16446.0	Unite d Kingd om	2011-12- 01	- 168469 .6	2011 -05- 01	8
1	C5414 33	23166	MEDIU M CERAM IC TOP STORA GE JAR	- 74215	2011- 01-18 10:17:00	1.04	12346.0	Unite d Kingd om	2011-01- 01	- 77183. 6	2011 -01- 01	1
4	C5367 57	84347	ROTA TI NG SILVER ANGEL S T- LIGHT HLDR	-9360	2010- 12-02 14:23:00	0.03	15838.0	Unite d Kingd om	2010-12- 01	-280.8	2010 -12- 01	1

	Invoic eNo	StockC ode	Description	Quant ity	Invoice Date	UnitP rice	Custome rID	Count ry	InvoiceM onth	Total_ cost	Mon th	CohortI ndex
1 0	C5504 56	21108	FAIRY CAKE FLANN EL ASSOR TED COLOU R	-3114	2011- 04-18 13:08:00	2.10	15749.0	Unite d Kingd om	2011-04- 01	-6539.4	2011 -01- 01	4
2 0	C5504 56	21175	GIN + TONIC DIET METAL SIGN	-2000	2011- 04-18 13:08:00	1.85	15749.0	Unite d Kingd om	2011-04- 01	-3700.0	2011 -01- 01	4

#retention rate or Active Customers

```
grouping = df.groupby(['Month', 'CohortIndex'])
```

Count the number of unique values per customer ID

```
cohort_data = grouping['CustomerID'].apply(pd.Series.nunique).reset_index()
```

Create a pivot

```
cohort_counts = cohort_data.pivot(index='Month', columns='CohortIndex', values='CustomerID')
```

```
cohort_counts
```

CohortIndex	1	2	3	4	5	6	7	8	9	10	11	12	13
Month													
2010-12-01	948.0	362.0	317.0	367.0	341.0	376.0	360.0	336.0	336.0	374.0	354.0	474.0	260.0
2011-01-01	421.0	101.0	119.0	102.0	138.0	126.0	110.0	108.0	131.0	146.0	155.0	63.0	NaN
2011-02-01	380.0	94.0	73.0	106.0	102.0	94.0	97.0	107.0	98.0	119.0	35.0	NaN	NaN
2011-03-01	440.0	84.0	112.0	96.0	102.0	78.0	116.0	105.0	127.0	39.0	NaN	NaN	NaN


```
In [53]: # Select the first column and store it to cohort_sizes
cohort_sizes = cohort_counts.iloc[:,0]
cohort_sizes
```

```
Out[53]: Month
2010-12-01    948.0
2011-01-01    421.0
2011-02-01    388.0
2011-03-01    440.0
2011-04-01    299.0
2011-05-01    279.0
2011-06-01    235.0
2011-07-01    191.0
2011-08-01    167.0
2011-09-01    298.0
2011-10-01    352.0
2011-11-01    321.0
2011-12-01     41.0
Name: 1, dtype: float64
```

```
In [54]: # Divide the cohort count by cohort sizes along the rows
retention = cohort_counts.divide(cohort_sizes, axis=0)*100
```

```
In [55]: retention
```

```
Out[55]: CohortIndex   1      2      3      4      5      6      7      8      9      10     11     12     13
Month
2010-12-01  100.0  38.185654  33.438819  38.713080  35.970464  39.662447  37.974684  35.443038  35.443038  39.451477  37.341772  50.000000  27.42616
2011-01-01  100.0  23.990499  28.266033  24.228029  32.779097  29.928741  26.128266  25.653207  31.116390  34.679335  36.817102  14.964371  NaN
2011-02-01  100.0  24.736842  19.210526  27.894737  26.842105  24.736842  25.526316  28.157895  25.789474  31.315789  9.210526  NaN  NaN
2011-03-01  100.0  19.090909  25.454545  21.818182  23.181818  17.727273  26.363636  23.863636  28.863636  8.863636  NaN  NaN  NaN
2011-04-01  100.0  22.742475  22.073579  21.070234  20.735786  23.745819  23.076923  26.086957  8.361204  NaN  NaN  NaN  NaN
2011-05-01  100.0  23.655914  17.204301  17.204301  21.505376  24.372760  26.523297  10.394265  NaN  NaN  NaN  NaN  NaN
2011-06-01  100.0  20.851064  18.723404  27.234043  24.680851  33.617021  10.212766  NaN  NaN  NaN  NaN  NaN  NaN
2011-07-01  100.0  20.942408  20.418848  23.036649  27.225131  11.518325  NaN  NaN  NaN  NaN  NaN  NaN  NaN
2011-08-01  100.0  25.149701  25.149701  25.149701  13.772455  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN
2011-09-01  100.0  29.865772  32.550336  12.080537  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN
2011-10-01  100.0  26.420455  13.068182  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN
2011-11-01  100.0  13.395639  NaN  NaN
2011-12-01  100.0  NaN  NaN
```

retention.min(), retention.max()

```
(CohortIndex
1      100.000000
2      13.395639
3      13.068182
4      12.080537
5      13.772455
6      11.518325
7      10.212766
8      10.394265
9      8.361204
10     8.863636
11     9.210526
12     14.964371
13     27.426160
dtype: float64,
CohortIndex
1      100.000000
```

```
2      38.185654
3      33.438819
4      38.713080
5      35.970464
6      39.662447
7      37.974684
8      35.443038
9      35.443038
10     39.451477
11     37.341772
12     50.000000
13     27.426160
dtype: float64)

month_list = ["Dec '10", "Jan '11", "Feb '11", "Mar '11", "Apr '11", \
    "May '11", "Jun '11", "Jul '11", "Aug '11", "Sep '11", \
    "Oct '11", "Nov '11", "Dec '11"]

retention = retention/100

# Initialize inches plot figure
plt.figure(figsize=(15,7))

# Add a title
plt.title('Retention by Monthly Cohorts')

# Create the heatmap
sns.heatmap(data=retention,
            annot = True,
            fmt= '.0%',
            cmap = "GnBu",
            vmin = 0.0,
            vmax = list(retention.max().sort_values(ascending = False))[1]+3,
            fmt = '.1%',
            linewidth = 0.3,
            yticklabels=month_list)
```

```
plt.show();
```



```
#average price per cohort
```

```
# Create a groupby object and pass the monthly cohort and cohort index as a list
```

```
grouping = df.groupby(['Month', 'CohortIndex'])
```

```
# Calculate the average of the unit price column
```

```
cohort_data = grouping['UnitPrice'].mean()
```

```
# Reset the index of cohort_data
```

```
cohort_data = cohort_data.reset_index()
```

```
# Create a pivot
```

```
average_price = cohort_data.pivot(index='Month', columns='CohortIndex', values='UnitPrice')
```

```
average_price.round(1)
```

```
average_price.index = average_price.index.date
```

average_price

Cohort Index	1	2	3	4	5	6	7	8	9	10	11	12	13
2011-12-01	2.288 479	NaN											

```

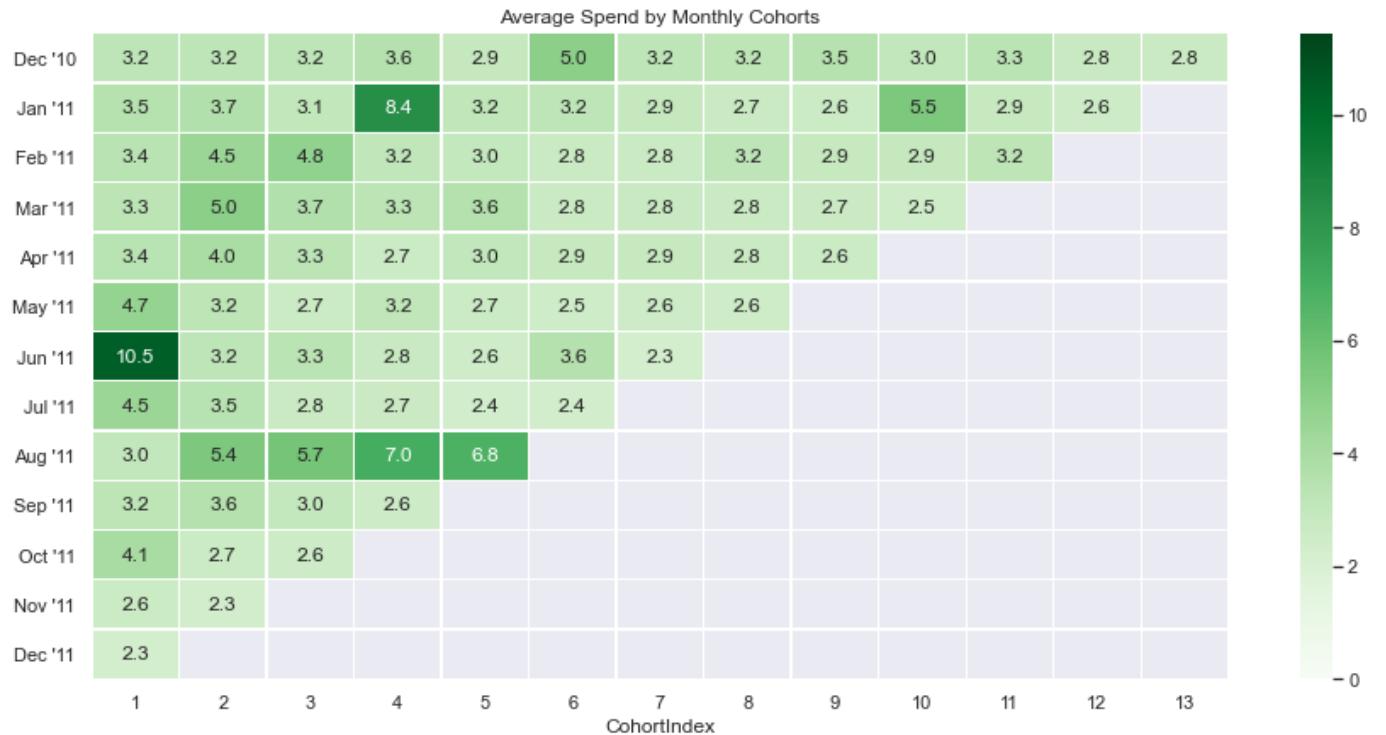
plt.figure(figsize=(15, 7))

# Add a title
plt.title('Average Spend by Monthly Cohorts')

# Create the heatmap
sns.heatmap(data = average_price,
             annot=True,
             vmin = 0.0,
#             vmax =20,
             cmap='Greens',
             vmax = list(average_price.max().sort_values(ascending = False))[1]+3,
             fmt = '.1f',
             linewidth = 0.3,
             yticklabels=month_list)

plt.show();

```



```
#average quantity per cohort
```

```
# Create a groupby object and pass the monthly cohort and cohort index as a list
```

```
grouping = df.groupby(['Month', 'CohortIndex'])
```

```
# Calculate the average of the Quantity column
```

```
cohort_data = grouping['Quantity'].mean()
```

```
# Reset the index of cohort_data
```

```
cohort_data = cohort_data.reset_index()
```

```
# Create a pivot
```

```
average_quantity = cohort_data.pivot(index='Month', columns='CohortIndex', values='Quantity')
```

```
average_quantity
```

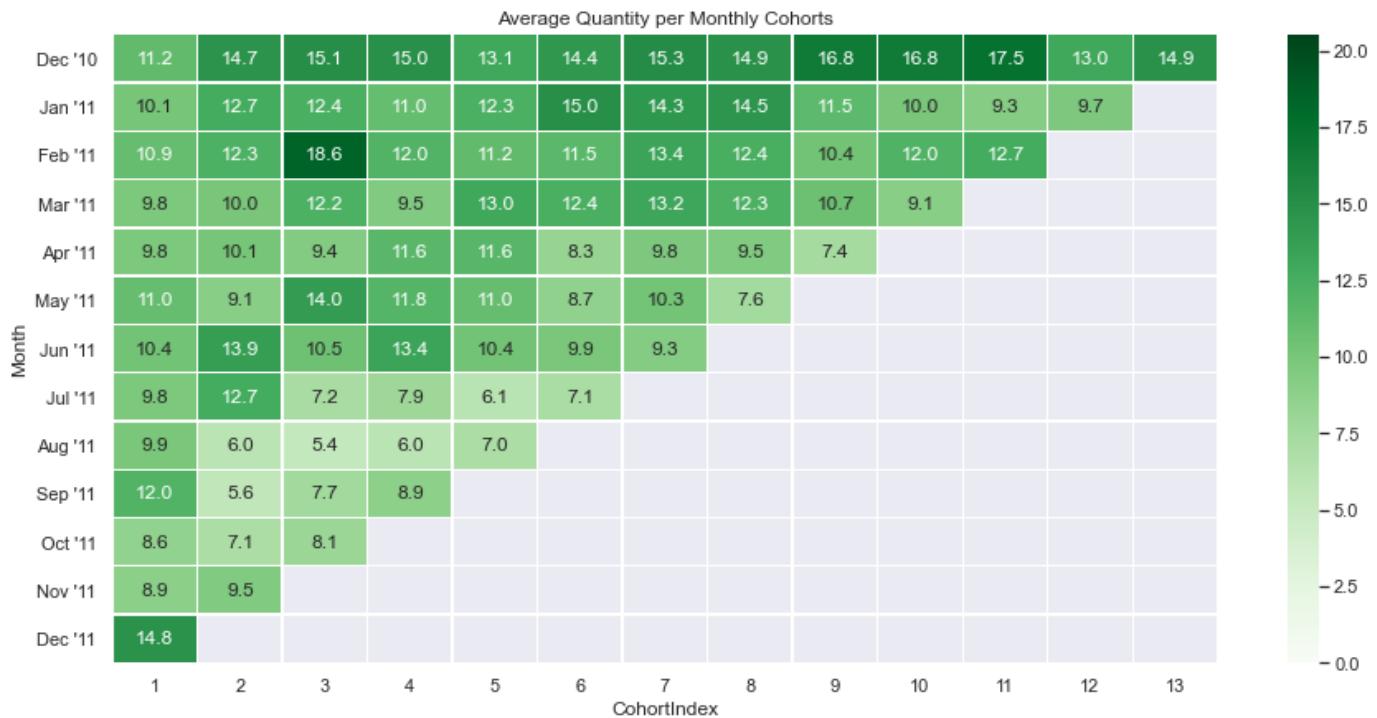
Out[61]:

CohortIndex	1	2	3	4	5	6	7	8	9	10	11	12	13
Month													
2010-12-01	11.200463	14.691852	15.108447	14.954097	13.054649	14.416287	15.306910	14.879447	16.764934	16.809158	17.528956	13.019471	14.901201
2011-01-01	10.127231	12.704190	12.429557	11.032382	12.288608	15.006101	14.302480	14.519414	11.451025	9.982762	9.256968	9.737305	NaN
2011-02-01	10.924450	12.251366	18.563808	12.018144	11.167271	11.476727	13.378526	12.448602	10.381961	12.043074	12.702765	NaN	NaN
2011-03-01	9.818050	9.972109	12.249296	9.483094	13.037510	12.369617	13.221102	12.263293	10.662973	9.091004	NaN	NaN	NaN
2011-04-01	9.803935	10.130252	9.432453	11.622102	11.645560	8.315994	9.777895	9.480778	7.403071	NaN	NaN	NaN	NaN
2011-05-01	10.977360	9.138087	14.023864	11.805435	10.973613	8.740725	10.275862	7.576774	NaN	NaN	NaN	NaN	NaN
2011-06-01	10.411028	13.859783	10.509642	13.384102	10.360800	9.901184	9.348609	NaN	NaN	NaN	NaN	NaN	NaN
2011-07-01	9.804225	12.700952	7.229385	7.929151	6.101961	7.111538	NaN						
2011-08-01	9.941459	5.983114	5.371409	5.972992	6.980110	NaN							
2011-09-01	12.003023	5.551129	7.657590	8.873418	NaN								
2011-10-01	8.553545	7.056196	8.079686	NaN									
2011-11-01	8.901297	9.508021	NaN										
2011-12-01	14.795478	NaN											

In [62]:

```
plt.figure(figsize=(15, 7))
plt.title('Average Quantity per Monthly Cohorts')

# Create the heatmap
sns.heatmap(data = average_quantity,
            annot=True,
            vmin = 0.0,
            cmap='Greens',
            vmax = list(average_quantity.max().sort_values(ascending = False))[1]+3,
            fmt = '.1f',
            linewidth = 0.3,
            ticklabels=month_list)
plt.show();
```



#Creating a copy of df as safe copy. Will be using df1 for changes

```
df1 = df
```

```
df1
```

#For recency, need to get the date difference since the last purchase.

#For this using the last purchase date on the database as today's date

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	Total_cost	Month	CohortIndex
0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom	2011-12-01	-168469.6	2011-05-01	8
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom	2011-01-01	-77183.6	2011-01-01	1
4	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom	2010-12-01	-280.8	2010-12-01	1
10	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	15749.0	United Kingdom	2011-04-01	-6539.4	2011-01-01	4
20	C550456	21175	GIN + TONIC DIET METAL SIGN	-2000	2011-04-18 13:08:00	1.85	15749.0	United Kingdom	2011-04-01	-3700.0	2011-01-01	4
...
541903	554868	22197	SMALL POPCORN HOLDER	4300	2011-05-27 10:52:00	0.72	13135.0	United Kingdom	2011-05-01	3096.0	2011-05-01	1
541904	573008	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	4800	2011-10-27 12:26:00	0.21	12901.0	United Kingdom	2011-10-01	1008.0	2011-03-01	8
541906	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540	2011-11-25 15:57:00	0.00	13256.0	United Kingdom	2011-11-01	0.0	2011-11-01	1
541907	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	2011-01-18 10:01:00	1.04	12346.0	United Kingdom	2011-01-01	77183.6	2011-01-01	1
541908	581483	23843	PAPER CRAFT , LITTLE BIRDIE	80995	2011-12-09 09:15:00	2.08	16446.0	United Kingdom	2011-12-01	168469.6	2011-05-01	8

401604 rows × 12 columns

```
In [64]: #last date available in our dataset  
df1['InvoiceDate'].max()
```

```
Out[64]: Timestamp('2011-12-09 12:50:00')
```

```
In [65]: current_date = df1['InvoiceDate'].max()  
current_date = pd.to_datetime(current_date).date()  
current_date
```

```
Out[65]: datetime.date(2011, 12, 9)
```

```
In [66]: # Lets create a date column for date values only  
df1['Purchase_Date'] = df1.InvoiceDate.dt.date
```

```
In [67]: recency = df1.groupby('CustomerID')['Purchase_Date'].max().reset_index()
# Creating a separate column for this date.
recency['Current_Date'] = current_date
# Compute the number of days since last purchase
recency['Recency'] = recency.Purchase_Date.apply(lambda x: (current_date - x).days)
recency
```

Out[67]:

	CustomerID	Purchase_Date	Current_Date	Recency
0	12346.0	2011-01-18	2011-12-09	325
1	12347.0	2011-12-07	2011-12-09	2
2	12348.0	2011-09-25	2011-12-09	75
3	12349.0	2011-11-21	2011-12-09	18
4	12350.0	2011-02-02	2011-12-09	310
...
4367	18280.0	2011-03-07	2011-12-09	277
4368	18281.0	2011-06-12	2011-12-09	180
4369	18282.0	2011-12-02	2011-12-09	7
4370	18283.0	2011-12-06	2011-12-09	3
4371	18287.0	2011-10-28	2011-12-09	42

4372 rows × 4 columns

```
In [68]: # Drop Current tdate that we took for calculation. That is no more required
recency.drop(['Purchase_Date', 'Current_Date'], axis=1, inplace=True)

#Now finding out Frequency - how often or how many a customer used the product of a company.

frequency = df1.groupby('CustomerID').InvoiceNo.nunique().reset_index().rename(columns={'InvoiceNo':'Frequency'})
frequency
```

Out[68]:

	CustomerID	Frequency
0	12346.0	2
1	12347.0	7
2	12348.0	4
3	12349.0	1
4	12350.0	1
...

4367	18280.0	1
4368	18281.0	1
4369	18282.0	3
4370	18283.0	16
4371	18287.0	3

4372 rows × 2 columns

Doing the same on the Spending or Monetary Monetary is the total amount of money a customer spent in that given period. Therefore big spenders will be differentiated with other customers such as MVP or VIP.

In [69]: *#We had already calculated Total Cost Earlier, using that*

```
monetary = df1.groupby('CustomerID').Total_cost.sum().reset_index().rename(columns={'Total_cost':'Monetary'})  
monetary
```

Out[69]:

	CustomerID	Monetary
0	12346.0	0.00
1	12347.0	4310.00
2	12348.0	1797.24
3	12349.0	1757.55
4	12350.0	334.40
...
4367	18280.0	180.60
4368	18281.0	80.82
4369	18282.0	176.60
4370	18283.0	2045.53
4371	18287.0	1837.28

4372 rows × 2 columns

In [70]: *#Also creating a separate table for Customers for RFM. First merging frequency with recency and then that table with monetary*

```
temp_ = recency.merge(frequency, on='CustomerID')  
rfm_table = temp_.merge(monetary, on='CustomerID')
```

```
In [71]: rfm_table
```

```
Out[71]:
```

	CustomerID	Recency	Frequency	Monetary
0	12346.0	325	2	0.00
1	12347.0	2	7	4310.00
2	12348.0	75	4	1797.24
3	12349.0	18	1	1757.55
4	12350.0	310	1	334.40
...
4367	18280.0	277	1	180.60
4368	18281.0	180	1	80.82
4369	18282.0	7	3	176.60
4370	18283.0	3	16	2045.53
4371	18287.0	42	3	1837.28

4372 rows × 4 columns

```
In [72]: #RFM Table integrity Check
```

```
# Fetch the records corresponding to the first customer id in above table
df1.groupby('CustomerID').Total_cost.sum()
#Data matches
```

```
Out[72]: CustomerID
```

```
12346.0      0.00
12347.0    4310.00
12348.0    1797.24
12349.0    1757.55
12350.0     334.40
...
18280.0     180.60
18281.0     80.82
18282.0     176.60
18283.0   2045.53
18287.0   1837.28
Name: Total_cost, Length: 4372, dtype: float64
```

```
In [73]: temp = df1.groupby('CustomerID').InvoiceDate.max().dt.date
```

```
temp = current_date - temp
temp
#Data Matches
```

```
Out[73]: CustomerID
12346.0    325 days
12347.0      2 days
12348.0     75 days
12349.0    18 days
12350.0    310 days
...
18280.0    277 days
18281.0    180 days
18282.0      7 days
18283.0      3 days
18287.0    42 days
Name: InvoiceDate, Length: 4372, dtype: timedelta64[ns]
```

```
In [74]: # RFM Quantiles
quantiles = rfm_table.quantile(q=[0.25,0.5,0.75, 1])
quantiles
```

```
Out[74]:
CustomerID  Recency  Frequency  Monetary
0.25        13812.75   16.0       1.0    291.795
0.50        15300.50   50.0       3.0    644.070
0.75        16778.25   143.0      5.0    1608.335
1.00        18287.00   373.0      248.0   279489.020
```

```
In [75]: #convert quartile information into a dictionary so that cutoffs can be picked up. Like a Lookup table
quantiles=quantiles.to_dict()
quantiles
```

```
Out[75]: {'CustomerID': {0.25: 13812.75, 0.5: 15300.5, 0.75: 16778.25, 1.0: 18287.0},
'Recency': {0.25: 16.0, 0.5: 50.0, 0.75: 143.0, 1.0: 373.0},
'Frequency': {0.25: 1.0, 0.5: 3.0, 0.75: 5.0, 1.0: 248.0},
'Monetary': {0.25: 291.795,
0.5: 644.070000000002,
0.75: 1608.335,
1.0: 279489.0199999991}}
```

#RFM Segments

```
# Arguments (x = value, p = recency, monetary_value, frequency, d = quantiles dict)
```

```
def RScore(x,p,d):
```

```
    if x <= d[p][0.25]:
```

```
        return 4
```

```
    elif x <= d[p][0.50]:
```

```
        return 3
```

```
    elif x <= d[p][0.75]:
```

```
        return 2
```

```
    else:
```

```
        return 1
```

```
# Arguments (x = value, p = recency, monetary_value, frequency, k = quantiles dict)
```

```
def FMScore(x,p,d):
```

```
    if x <= d[p][0.25]:
```

```

    return 1

elif x <= d[p][0.50]:
    return 2

elif x <= d[p][0.75]:
    return 3

else:
    return 4

#rfm_table['segment'] = rfm_table.copy()

rfm_table['R_Quartile'] = rfm_table['Recency'].apply(RScore, args=('Recency',quantiles,))
rfm_table['F_Quartile'] = rfm_table['Frequency'].apply(FMScore, args=('Frequency',quantiles,))
rfm_table['M_Quartile'] = rfm_table['Monetary'].apply(FMScore, args=('Monetary',quantiles,))

rfm_table

```

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile
0	12346.0	325	2	0.00	1	2	1
1	12347.0	2	7	4310.00	4	4	4
2	12348.0	75	4	1797.24	2	3	4
3	12349.0	18	1	1757.55	3	1	4
4	12350.0	310	1	334.40	1	1	2
...
4367	18280.0	277	1	180.60	1	1	1
4368	18281.0	180	1	80.82	1	1	1
4369	18282.0	7	3	176.60	4	2	1

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile
4370	18283.0	3	16	2045.53	4	4	4
4371	18287.0	42	3	1837.28	3	2	4

4372 rows × 7 columns

```
rfm_table['RFMScore'] = rfm_table.R_Quartile.map(str) \
    + rfm_table.F_Quartile.map(str) \
    + rfm_table.M_Quartile.map(str)
```

rfm_table

CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore
0	12346.0	325	2	0.00	1	2	1 121
1	12347.0	2	7	4310.00	4	4	4 444
2	12348.0	75	4	1797.24	2	3	4 234
3	12349.0	18	1	1757.55	3	1	4 314
4	12350.0	310	1	334.40	1	1	2 112
...
4367	18280.0	277	1	180.60	1	1	1 111
4368	18281.0	180	1	80.82	1	1	1 111
4369	18282.0	7	3	176.60	4	2	1 421

CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score
4370	18283.0	3	16	2045.53	4	4	4	444
4371	18287.0	42	3	1837.28	3	2	4	324

4372 rows × 8 columns

```
In [79]: #Integrity Check - Passed
rfm_table.iloc[0,7] < rfm_table.iloc[1,7]
```

```
Out[79]: True
```

```
In [80]: rfm_table['RFM_Score'] = rfm_table[['R_Quartile','F_Quartile','M_Quartile']].sum(axis=1)
#Assigning Score to the RFM before categorization. Will help later in plotting
rfm_table
```

```
Out[80]:
```

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score
0	12346.0	325	2	0.00	1	2	1	121	4
1	12347.0	2	7	4310.00	4	4	4	444	12
2	12348.0	75	4	1797.24	2	3	4	234	9
3	12349.0	18	1	1757.55	3	1	4	314	8
4	12350.0	310	1	334.40	1	1	2	112	4
...
4367	18280.0	277	1	180.60	1	1	1	111	3
4368	18281.0	180	1	80.82	1	1	1	111	3
4369	18282.0	7	3	176.60	4	2	1	421	7
4370	18283.0	3	16	2045.53	4	4	4	444	12
4371	18287.0	42	3	1837.28	3	2	4	324	9

4372 rows × 9 columns

```
In [81]: # Create a dictionary for each segment to map them against each customer
segment_dict = {
    'Best Customers':'444',      # Highest frequency as well as monetary value with Least recency
    'Loyal Customers':'344',      # High frequency as well as monetary value with good recency
    'Big Spenders':'334',        # High monetary value but good recency and frequency values
    'Almost Lost':'244',         # Customer's shopping less often now who used to shop a lot
    'Lost Customers':'144',       # Customer's shopped Long ago who used to shop a lot.
    'Recent Customers':'443',     # Customer's who recently started shopping a lot but with less monetary value
    'Lost Cheap Customers':'122', # Customer's shopped long ago but with less frequency and monetary value
    'No Harm to Lose Cheap Customers':'211' # Customer's shopped sometime back ago but with less frequency and monetary value
}
```

```
In [82]: # Swap the key and value of dictionary. So that Lookup is from value to Customer type and not vice-versa
dict_segment = dict(zip(segment_dict.values(),segment_dict.keys()))
rfm_table['Segment'] = rfm_table.RFMScore.map(lambda x: dict_segment.get(x))
```

```
In [83]: rfm_table
```

```
Out[83]:
```

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score	Segment
0	12346.0	325	2	0.00	1	2	1	121	4	None
1	12347.0	2	7	4310.00	4	4	4	444	12	Best Customers
2	12348.0	75	4	1797.24	2	3	4	234	9	None
3	12349.0	18	1	1757.55	3	1	4	314	8	None
4	12350.0	310	1	334.40	1	1	2	112	4	None
...
4367	18280.0	277	1	180.60	1	1	1	111	3	None
4368	18281.0	180	1	80.82	1	1	1	111	3	None
4369	18282.0	7	3	176.60	4	2	1	421	7	None
4370	18283.0	3	16	2045.53	4	4	4	444	12	Best Customers
4371	18287.0	42	3	1837.28	3	2	4	324	9	None

4372 rows × 10 columns

```
In [84]: rfm_table.groupby('Segment').CustomerID.count()
```

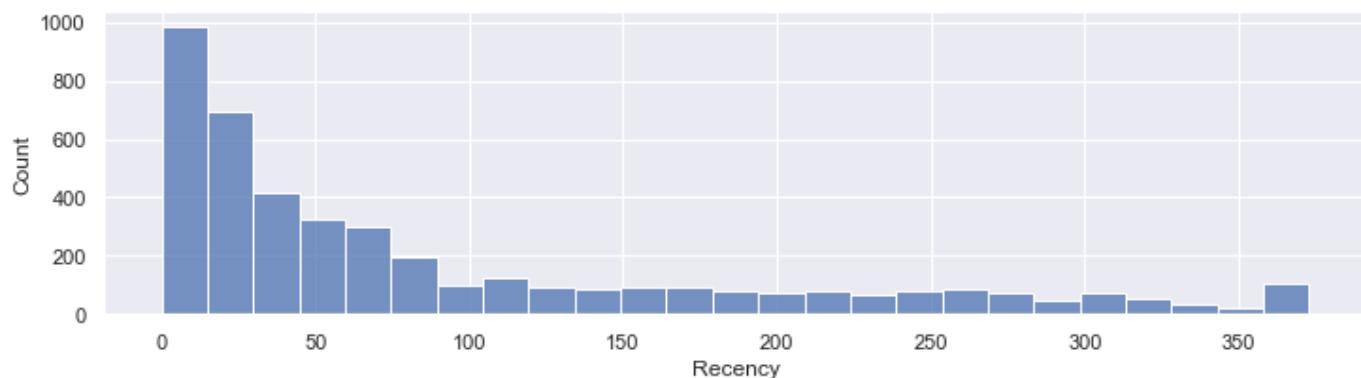
```
Out[84]: Segment
```

```
Almost Lost          90
Best Customers      482
Big Spenders         55
Lost Cheap Customers 151
Lost Customers        13
Loyal Customers       225
No Harm to Lose Cheap Customers 177
Recent Customers      99
Name: CustomerID, dtype: int64
```

Almost Lost 90 Best Customers 482 - These need Promotion materials and other engagement Big Spenders 55 These need exclusive product - high end newsletters Lost Cheap Customers 151 They may come back but focus spending is not recommended Lost Customers 13 No effort to win them back Loyal Customers 225 They need focused product list and AI built-in recommendation No Harm to Lose Cheap Customers 177 Won't spend of this category Recent Customers 99 Need to keep their interest alive - promote newer things on their spending using product recommendation

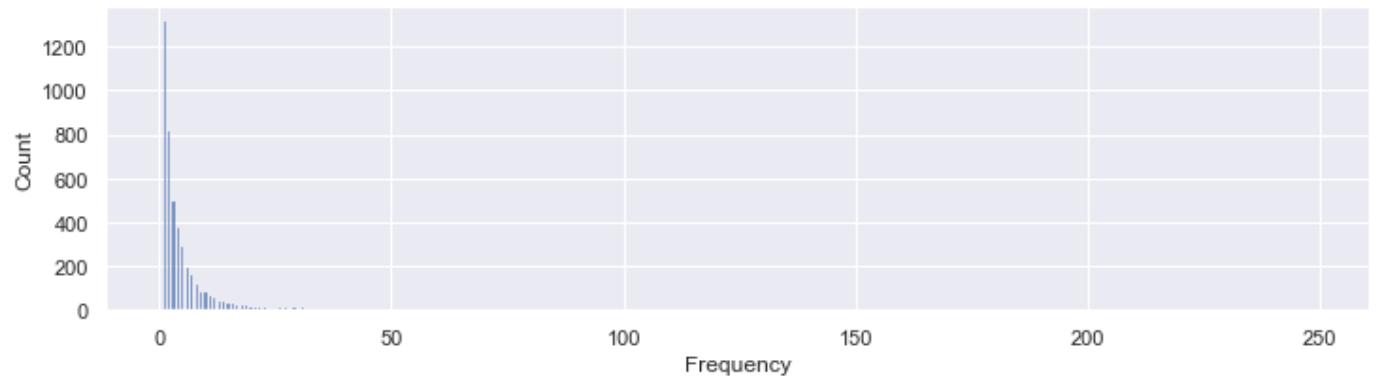
```
In [85]: # Checking the distribution of variables.
```

```
plt.figure(figsize=(12,10))
# Plot distribution of Recency
plt.subplot(3, 1, 1); sns.histplot(rfm_table['Recency'])
```



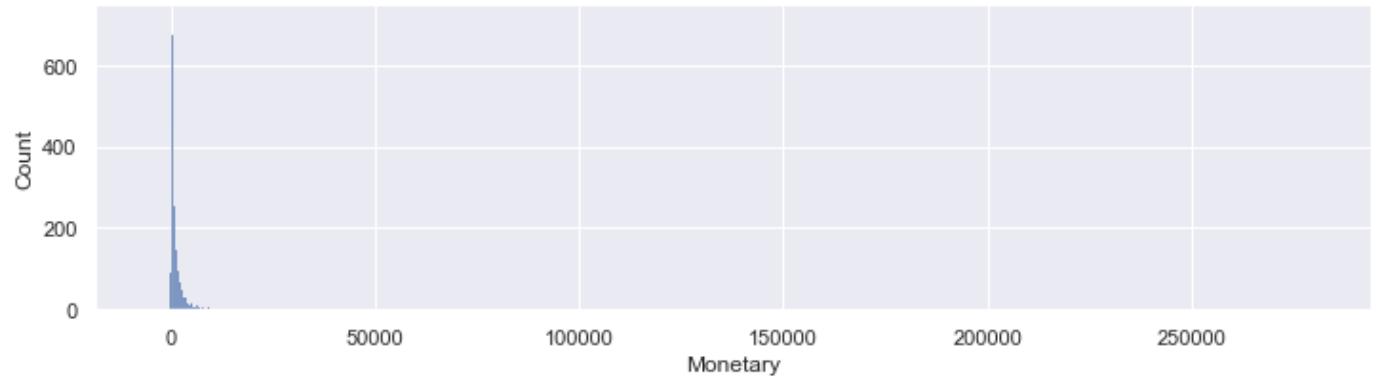
```
# Plot distribution of Frequency
```

```
plt.figure(figsize=(12,10))
plt.subplot(3, 1, 2); sns.histplot(rfm_table['Frequency'])
```



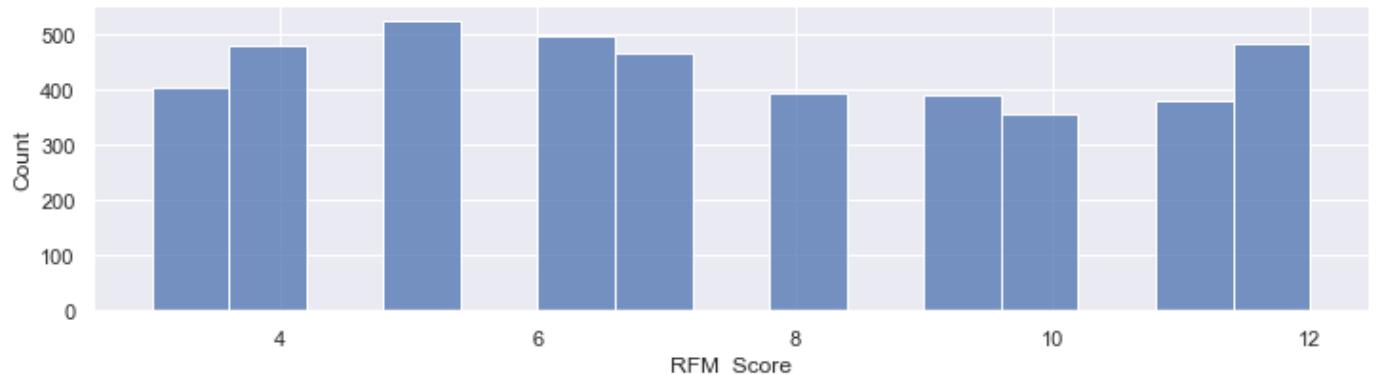
```
# Checking the distribution of variables.
```

```
plt.figure(figsize=(12,10))  
  
# Plot distribution of Monetary  
  
plt.subplot(3, 1, 3); sns.histplot(rfm_table['Monetary'])
```



```
# Checking the distribution of variables.
```

```
plt.figure(figsize=(12,10))  
  
# Plot distribution of RFM_Score Segment  
  
plt.subplot(3, 1, 3); sns.histplot(rfm_table['RFM_Score'])
```



```
rfm_table.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4372 entries, 0 to 4371
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   CustomerID  4372 non-null   float64
 1   Recency     4372 non-null   int64  
 2   Frequency   4372 non-null   int64  
 3   Monetary    4372 non-null   float64
 4   R_Quartile  4372 non-null   int64  
 5   F_Quartile  4372 non-null   int64  
 6   M_Quartile  4372 non-null   int64  
 7   RFMScore    4372 non-null   object 
 8   RFM_Score   4372 non-null   int64  
 9   Segment     1292 non-null   object 
dtypes: float64(2), int64(6), object(2)
memory usage: 375.7+ KB
```

```
rfm_table.describe()
```

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFM_Score
count	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000
mean	15299.677722	91.581199	5.075480	1893.531433	2.510979	2.349039	2.500000	7.360018
std	1722.390705	100.772139	9.338754	8218.696204	1.117084	1.151264	1.118162	2.872703

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFM_Score
min	12346.000000	0.000000	1.000000	-4287.630000	1.000000	1.000000	1.000000	3.000000
25%	13812.750000	16.000000	1.000000	291.795000	2.000000	1.000000	1.750000	5.000000
50%	15300.500000	50.000000	3.000000	644.070000	3.000000	2.000000	2.500000	7.000000
75%	16778.250000	143.000000	5.000000	1608.335000	4.000000	3.000000	3.250000	10.000000
max	18287.000000	373.000000	248.000000	279489.020000	4.000000	4.000000	4.000000	12.000000

```
# Create a copy of rfm table for scaled calculation
```

```
rfm_s = rfm_table.copy()
```

```
# Shift all values in the column by adding absolute of minimum value to each value, thereby making each value positive.
```

```
rfm_s.Monetary = rfm_s.Monetary + abs(rfm_s.Monetary.min()) + 1
```

```
rfm_s.Recency = rfm_s.Recency + abs(rfm_s.Recency.min()) + 1
```

```
# Check the summary of new values
```

```
rfm_s.describe()
```

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFM_Score
count	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000	4372.000000
mean	15299.677722	92.581199	5.075480	6182.161433	2.510979	2.349039	2.500000	7.360018

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFM_Score
std	1722.390705	100.772139	9.338754	8218.696204	1.117084	1.151264	1.118162	2.872703
min	12346.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	3.000000
25%	13812.750000	17.000000	1.000000	4580.425000	2.000000	1.000000	1.750000	5.000000
50%	15300.500000	51.000000	3.000000	4932.700000	3.000000	2.000000	2.500000	7.000000
75%	16778.250000	144.000000	5.000000	5896.965000	4.000000	3.000000	3.250000	10.000000
max	18287.000000	374.000000	248.000000	283777.650000	4.000000	4.000000	4.000000	12.000000

rfm_s.head()

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score	Segment
0	12346.0	326	2	4288.63	1	2	1	121	4	None
1	12347.0	3	7	8598.63	4	4	4	444	12	Best Customers
2	12348.0	76	4	6085.87	2	3	4	234	9	None
3	12349.0	19	1	6046.18	3	1	4	314	8	None
4	12350.0	311	1	4623.03	1	1	2	112	4	None

raw_data = rfm_s[['Recency','Frequency','Monetary']]

```
data_log = np.log(raw_data)

# Initialize a standard scaler and fit it
scaler = StandardScaler()
scaler.fit(data_log)

# Scale and center the data
data_normalized = scaler.transform(data_log)

# Create a pandas DataFrame
data_norm = pd.DataFrame(data=data_log, index=raw_data.index, columns=raw_data.columns)

data_norm.head()
```

Out[93]:

	Recency	Frequency	Monetary
0	5.786897	0.693147	8.363723
1	1.098612	1.945910	9.059358
2	4.330733	1.386294	8.713725
3	2.944439	0.000000	8.707182
4	5.739793	0.000000	8.438806

In [94]: *#Plotting the figures again to see if it is normalized*

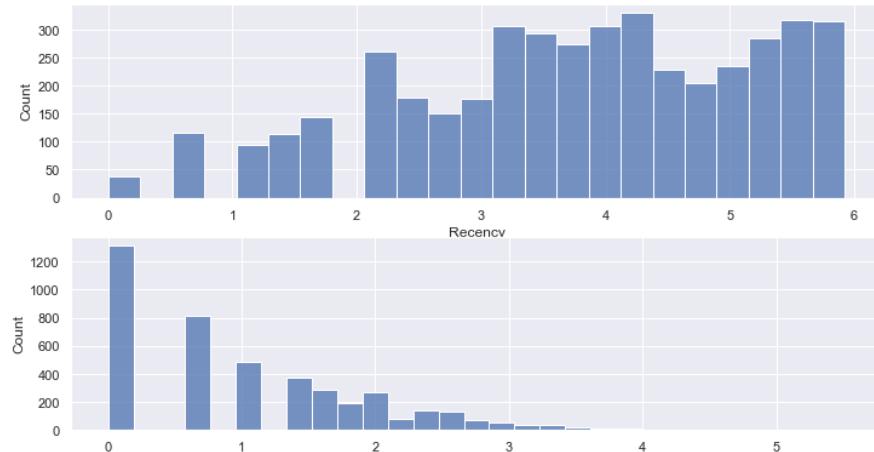
```
plt.figure(figsize=(12,10))

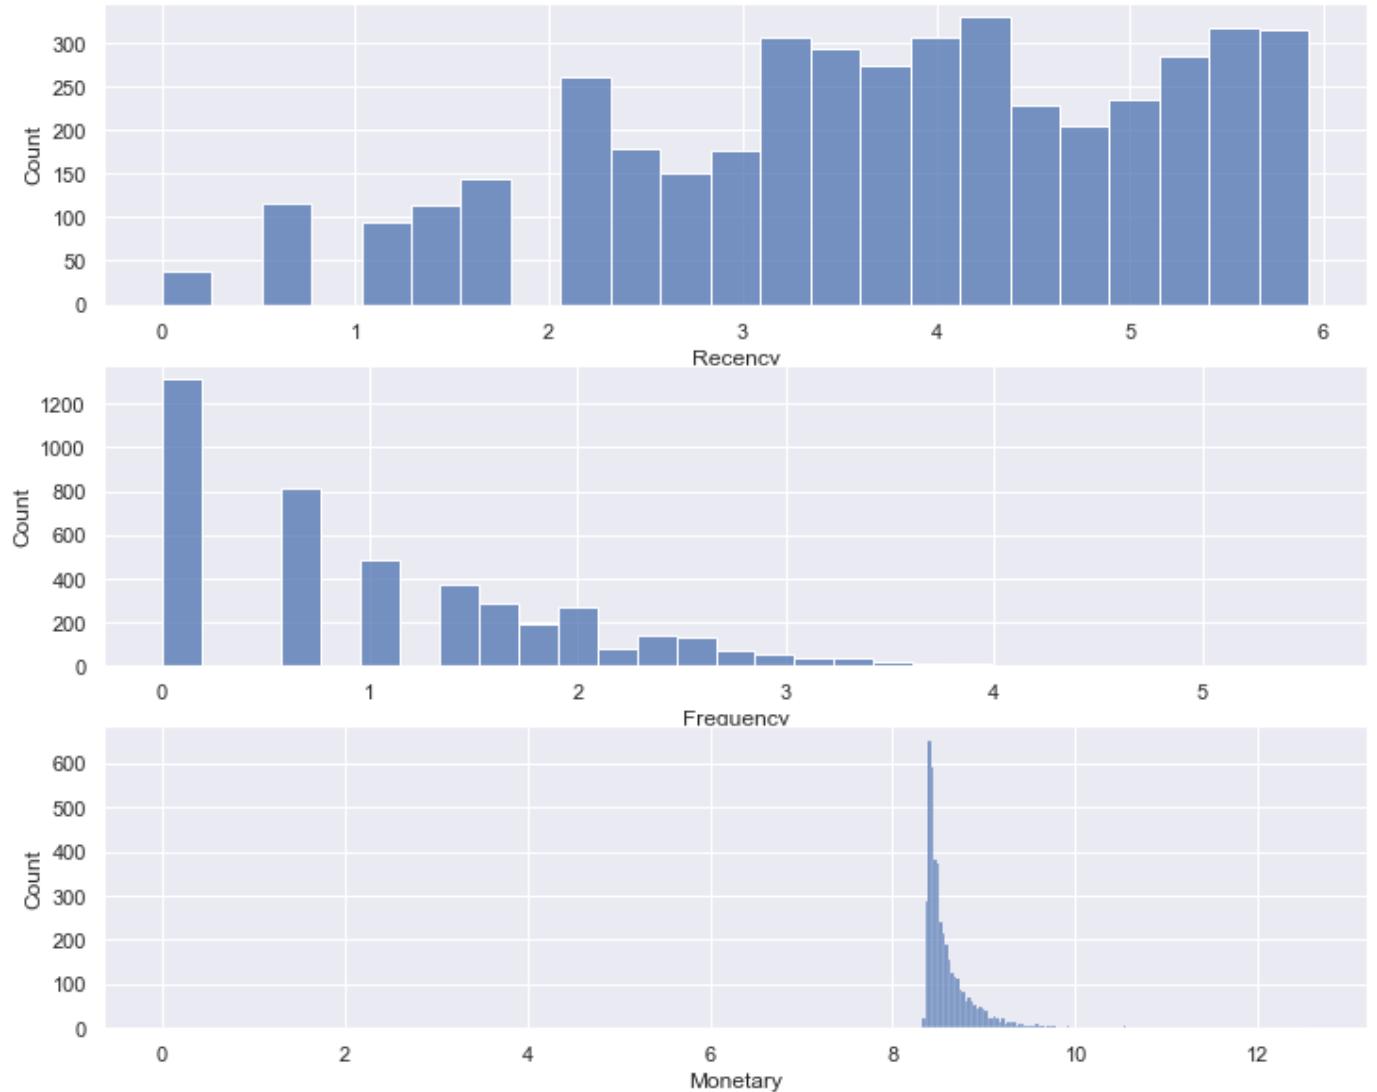
# Plot recency distribution
plt.subplot(3, 1, 1); sns.histplot(data_norm['Recency'])

# Plot frequency distribution
plt.subplot(3, 1, 2); sns.histplot(data_norm['Frequency'])

# Plot monetary value distribution
plt.subplot(3, 1, 3); sns.histplot(data_norm['Monetary'])

# Show the plot
plt.show()
```





```
sse = []
```

```
# Fit KMeans and calculate SSE for each k
```

```
for k in range(1, 21):
```

```
# Initialize KMeans with k clusters
```

```
kmeans = KMeans(n_clusters=k, random_state=1)
```

```
# Fit KMeans on the normalized dataset
```

```

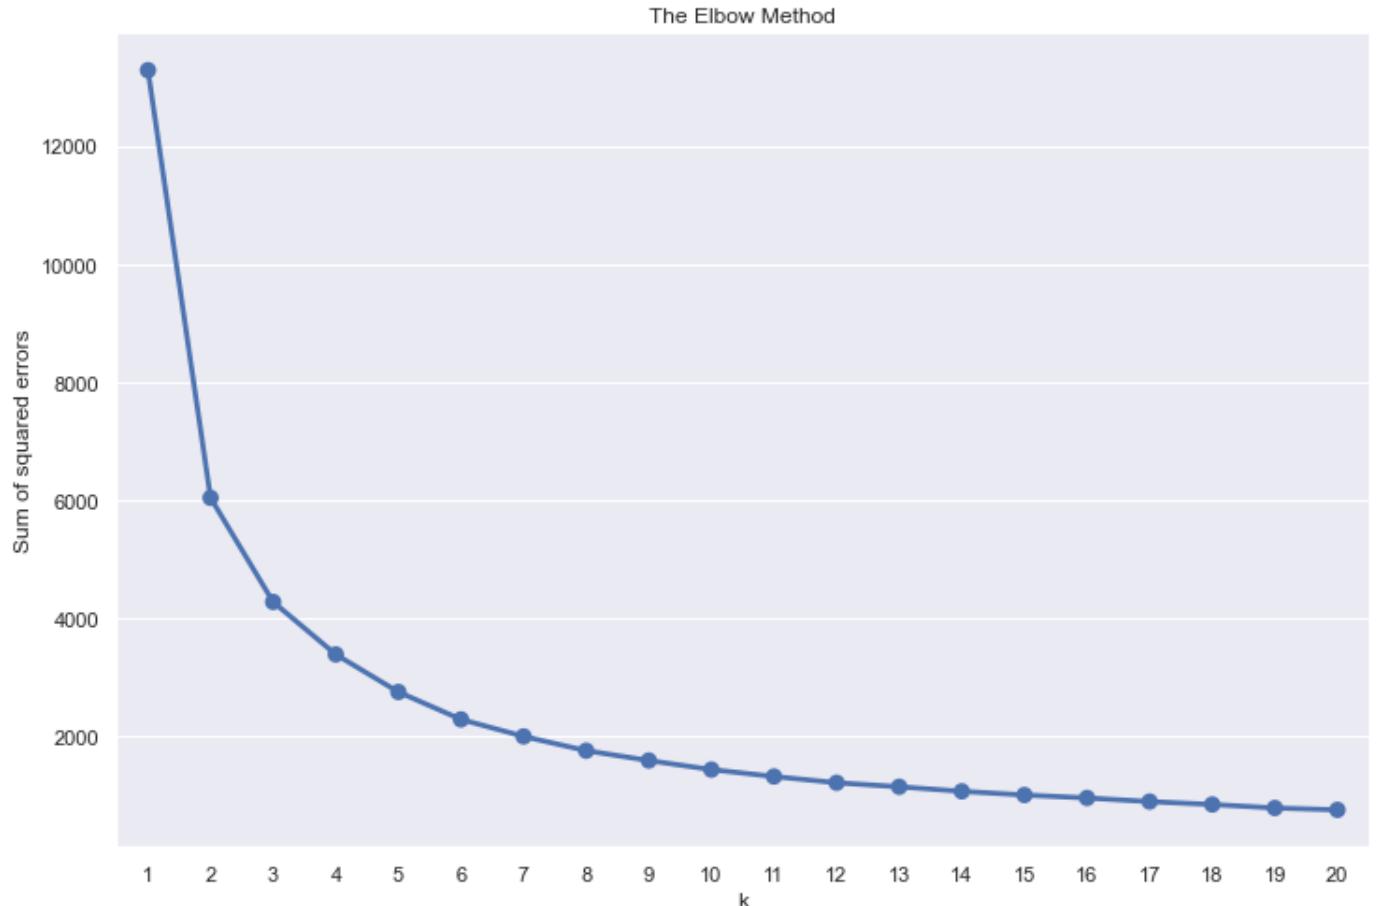
kmeans.fit(data_norm)

# Assign sum of squared distances to k element of dictionary
sse[k] = kmeans.inertia_

plt.figure(figsize=(12,8))

plt.title('The Elbow Method')
plt.xlabel('k');
plt.ylabel('Sum of squared errors')
sns.pointplot(x=list(sse.keys()), y=list(sse.values()))
plt.show()

```



#Confirming as the same

```

from sklearn.metrics import silhouette_score
wcss_silhouette = []
for i in range(2,12):
    km = KMeans(n_clusters=i, random_state=0, init='k-means++').fit(data_norm)
    preds = km.predict(data_norm)
    silhouette = silhouette_score(data_norm, preds)
    wcss_silhouette.append(silhouette)
    print("Silhouette score for number of cluster(s) {}: {}".format(i, silhouette))

```

5-

```

plt.figure(figsize=(10,5))
plt.title("The silhouette coefficient method \nfor determining number of clusters\n", fontsize=16)
plt.scatter(x=[i for i in range(2,12)], y=wcss_silhouette, s=150, edgecolor='k')
plt.grid(True)
plt.xlabel("Number of clusters", fontsize=14)
plt.ylabel("Silhouette score", fontsize=15)
plt.xticks([i for i in range(2,12)], fontsize=14)
plt.yticks(fontsize=15)
plt.show()

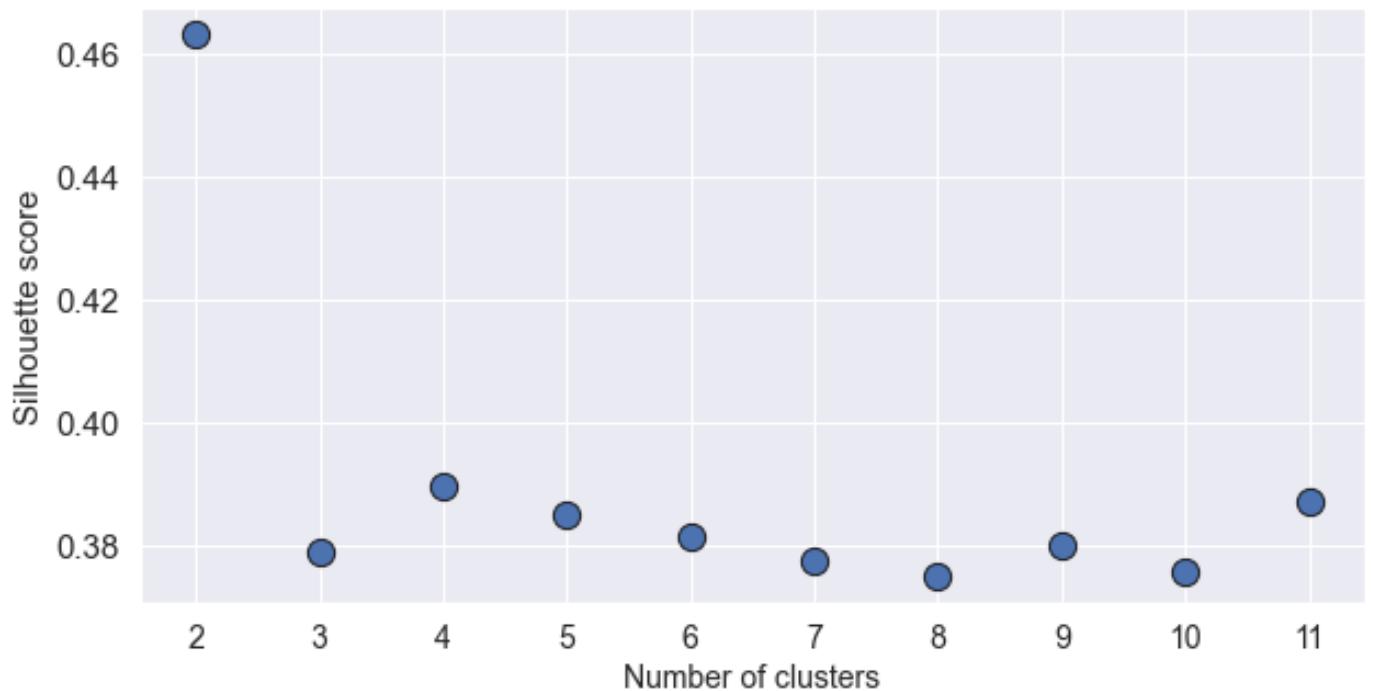
```

```

Silhouette score for number of cluster(s) 2: 0.46320292333881263
Silhouette score for number of cluster(s) 3: 0.37882385562792953
Silhouette score for number of cluster(s) 4: 0.38942882326747985
Silhouette score for number of cluster(s) 5: 0.38480602348390036
Silhouette score for number of cluster(s) 6: 0.3814620438903689
Silhouette score for number of cluster(s) 7: 0.37740088768033014
Silhouette score for number of cluster(s) 8: 0.3749214782348163
Silhouette score for number of cluster(s) 9: 0.379746242901208
Silhouette score for number of cluster(s) 10: 0.37572690058058267
Silhouette score for number of cluster(s) 11: 0.3869596058631919

```

The silhouette coefficient method for determining number of clusters



```
#Implementation of K-Means Clustering
```

```
plt.figure(figsize=(15,5))

model = KMeans(n_clusters = 5)
model.fit(data_norm)

#Extract cluster labels from labels_ attribute
cluster_labels = model.labels_
centers = np.array(model.cluster_centers_)

plt.plot()
```

```

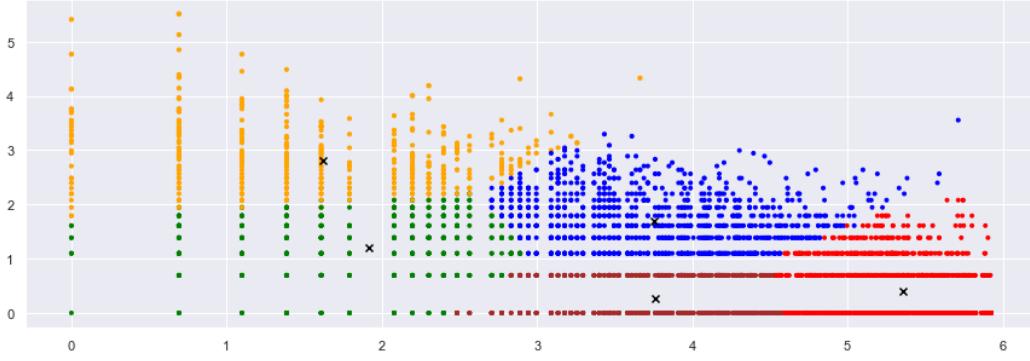
colormap = np.array(['Red', 'Blue', 'Green', 'Orange', 'Brown', 'Black'])

z = plt.scatter(data_norm.Recency, data_norm.Frequency, data_norm.Monetary, c = colormap[cluster_labels])

plt.scatter(centers[:,0], centers[:,1], marker="x", color='Black')

```

Out[141]: <matplotlib.collections.PathCollection at 0x2cae3bece50>



```

In [100]: # Create a cluster label column in the original DataFrame
data_norm_k5 = data_norm.assign(Cluster = model.labels_)
data_k5 = raw_data.assign(Cluster = model.labels_)

# Calculate average RFM values and size for each cluster
summary_k5 = data_k5.groupby(['Cluster']).agg({'Recency': 'mean',
                                                'Frequency': 'mean',
                                                'Monetary': ['mean', 'count'],}).round(0)

summary_k5

```

Out[100]:

Cluster	Recency	Frequency	Monetary	
	mean	mean	mean	count
0	229.0	2.0	4743.0	1279
1	9.0	4.0	5404.0	646
2	51.0	1.0	4751.0	873
3	53.0	6.0	6197.0	1120
4	7.0	21.0	14061.0	454

```

In [101]: data_norm_k5.index = rfm_s['CustomerID'].astype(int)
data_norm_k5

```

CustomerID	Recency	Frequency	Monetary	Cluster
12346	5.786897	0.693147	8.363723	0
12347	1.098612	1.945910	9.059358	4
12348	4.330733	1.386294	8.713725	3
12349	2.944439	0.000000	8.707182	2
12350	5.739793	0.000000	8.438806	0
...
18280	5.627621	0.000000	8.404971	0
18281	5.198497	0.000000	8.382392	0
18282	2.079442	1.098612	8.404076	1
18283	1.386294	2.772589	8.753712	4
18287	3.761200	1.098612	8.720283	3

4372 rows × 4 columns

```
In [102]: # Assign the clusters as column to each customer
Cluster_table = rfm_s.assign(Cluster = cluster_labels)
Cluster_table
```

Out[102]:

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score	Segment	Cluster
0	12346.0	326	2	4288.63	1	2	1	121	4	None	0
1	12347.0	3	7	8598.63	4	4	4	444	12	Best Customers	4
2	12348.0	76	4	6085.87	2	3	4	234	9	None	3
3	12349.0	19	1	6046.18	3	1	4	314	8	None	2
4	12350.0	311	1	4623.03	1	1	2	112	4	None	0
...
4367	18280.0	278	1	4469.23	1	1	1	111	3	None	0
4368	18281.0	181	1	4369.45	1	1	1	111	3	None	0
4369	18282.0	8	3	4465.23	4	2	1	421	7	None	1
4370	18283.0	4	16	6334.16	4	4	4	444	12	Best Customers	4
4371	18287.0	43	3	6125.91	3	2	4	324	9	None	3

4372 rows × 11 columns

```
In [103]: # Check counts of records assigned to different clusters  
Cluster_table.Cluster.value_counts()
```

```
Out[103]: 0    1279  
3    1120  
2    873  
1    646  
4    454  
Name: Cluster, dtype: int64
```

```
In [104]: #Cluster_table  
Inference = Cluster_table.groupby(['Cluster']).agg({'RFM_Score': 'mean'}).round(0)  
Inference  
#summary_k5 = data_k5.groupby(['Cluster'])
```

```
Out[104]:  
RFM_Score  
Cluster  
0    4.0  
1    9.0  
2    6.0  
3    9.0  
4    12.0
```

```
In [105]: Cluster_table[Cluster_table.Cluster == 0].sample(10)
```

```
Out[105]:  
CustomerID  Recency  Frequency  Monetary  R_Quartile  F_Quartile  M_Quartile  RFMScore  RFM_Score  Segment  Cluster  
4081        17890.0   323       2   4875.52      1          2          2       122         5  Lost Cheap Customers     0  
2016        15083.0   257       1   4376.83      1          1          1       111         3           None     0  
3223        16714.0   219       4   5167.69      1          3          3       133         7           None     0  
1826        14816.0   198       1   4560.48      1          1          1       111         3           None     0  
2550        15789.0   359       1   4639.93      1          1          2       112         4           None     0  
2765        16093.0   107       1   4305.63      2          1          1       211         4  No Harm to Lose Cheap Customers     0  
242         12641.0   116       1   4474.53      2          1          1       211         4  No Harm to Lose Cheap Customers     0  
758         13343.0   173       2   4592.56      1          2          2       122         5  Lost Cheap Customers     0  
841         13466.0   101       2   4586.58      2          2          2       222         6           None     0  
3798        17508.0   281       1   4675.94      1          1          2       112         4           None     0
```

```
In [106]: Cluster_table[Cluster_table.Cluster == 1].sample(10)
```

```
Out[106]:
```

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score	Segment	Cluster
520	13017.0	8	1	4492.63	4	1	1	411	6	None	1
3890	17631.0	2	3	4816.96	4	2	2	422	8	None	1
3885	17624.0	15	2	5132.11	4	2	3	423	9	None	1
232	12627.0	11	7	8767.16	4	4	4	444	12	Best Customers	1
4280	18167.0	4	6	5644.14	4	4	3	443	11	Recent Customers	1
566	13083.0	5	2	4592.38	4	2	2	422	8	None	1
3931	17682.0	11	7	6747.96	4	4	4	444	12	Best Customers	1
1471	14335.0	17	2	4743.49	4	2	2	422	8	None	1
1346	14157.0	20	3	4681.07	3	2	2	322	7	None	1
247	12646.0	5	2	5635.60	4	2	3	423	9	None	1

```
In [107]: Cluster_table[Cluster_table.Cluster == 2].sample(10)
```

```
Out[107]:
```

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score	Segment	Cluster
570	13091.0	22	2	4608.45	3	2	2	322	7	None	2
4188	18040.0	20	2	4645.83	3	2	2	322	7	None	2
1602	14508.0	23	2	4556.68	3	2	1	321	6	None	2
1202	13960.0	22	2	4518.85	3	2	1	321	6	None	2
1249	14029.0	64	2	4756.29	2	2	2	222	6	None	2
1404	14236.0	81	2	4779.49	2	2	2	222	6	None	2
2381	15565.0	51	2	4461.79	3	2	1	321	6	None	2
58	12420.0	64	1	4889.02	2	1	2	212	5	None	2
3446	17011.0	31	1	4559.53	3	1	1	311	5	None	2
3349	16878.0	25	2	4288.63	3	2	1	321	6	None	2

```
In [108]: Cluster_table[Cluster_table.Cluster == 3].sample(10)
```

Out[108]:

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score	Segment	Cluster
3556	17166.0	39	3	4478.12	3	2	1	321	6	None	3
1430	14273.0	52	4	4848.60	2	3	2	232	7	None	3
3147	16609.0	16	9	9739.59	4	4	4	444	12	Best Customers	3
2310	15468.0	36	5	4727.09	3	3	2	332	8	None	3
1477	14342.0	22	6	4896.34	3	4	2	342	9	None	3
1338	14147.0	50	3	4757.03	3	2	2	322	7	None	3
3246	16743.0	30	8	6455.51	3	4	4	344	11	Loyal Customers	3
3769	17462.0	52	3	4951.57	2	2	3	223	7	None	3
4208	18069.0	27	7	6283.02	3	4	4	344	11	Loyal Customers	3
1336	14145.0	47	5	7436.83	3	3	4	334	10	Big Spenders	3

In [109]: Cluster_table[Cluster_table.Cluster == 4].sample(10)

Out[109]:

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score	Segment	Cluster
2590	15838.0	12	21	37639.39	4	4	4	444	12	Best Customers	4
3215	16705.0	1	29	18234.76	4	4	4	444	12	Best Customers	4
3923	17673.0	2	7	5876.70	4	4	3	443	11	Recent Customers	4
3985	17754.0	1	6	5920.94	4	4	4	444	12	Best Customers	4
4340	18241.0	10	18	6346.72	4	4	4	444	12	Best Customers	4
3537	17139.0	16	15	14856.50	4	4	4	444	12	Best Customers	4
3014	16422.0	18	75	38094.32	3	4	4	344	11	Loyal Customers	4
115	12490.0	6	10	9706.56	4	4	4	444	12	Best Customers	4
277	12682.0	4	31	16568.45	4	4	4	444	12	Best Customers	4
4192	18044.0	5	11	6374.28	4	4	4	444	12	Best Customers	4

CLuster 0 & 4 does not matter for us. Their RFM Score avg is low, and as we see they do not fall under any specialized marketing plan. They are the ones, who have very low RF&M. Though their number is high. They may be chaced customer who happen to drop in by some add etc.

1 & 3 has few categories that we had defined. Still a lot of effort is not to be directed on this cluster. Normal exposure to brand is good enough.

2 is the category we should be foucsing out attention to. Their RFM avg is 12

#Melt the data into along format so RFM values and metric names are stored in 1 column each

data_melt = pd.melt(data_norm_k5.reset_index(),

```
    id_vars=['CustomerID', 'Cluster'],
    value_vars=['Recency', 'Frequency', 'Monetary'],
    var_name='Attribute',
    value_name='Value')
```

data_melt

CustomerID	Cluster	Attribute	Value
------------	---------	-----------	-------

0	12346	0	Recency	5.786897
---	-------	---	---------	----------

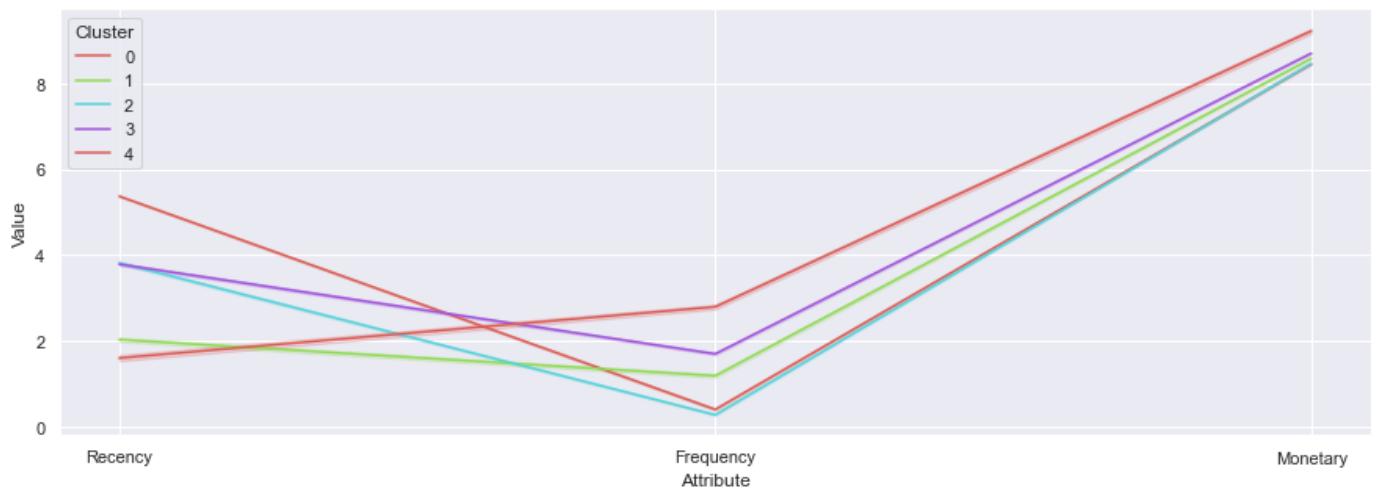
1	12347	4	Recency	1.098612
---	-------	---	---------	----------

	CustomerID	Cluster	Attribute	Value
2	12348	3	Recency	4.330733
3	12349	2	Recency	2.944439
4	12350	0	Recency	5.739793
...
13111	18280	0	Monetary	8.404971
13112	18281	0	Monetary	8.382392
13113	18282	1	Monetary	8.404076
13114	18283	4	Monetary	8.753712
13115	18287	3	Monetary	8.720283

13116 rows × 4 columns

```
plt.figure(figsize=(15,5))

sns.lineplot(x="Attribute", y="Value", hue='Cluster', palette = 'hls', data=data_melt)
```



```
In [112]: cluster_avg = data_k5.groupby(['Cluster']).mean()
population_avg = raw_data.head().mean()
```

```
In [113]: population_avg
```

```
Out[113]: Recency      147.000
Frequency     3.000
Monetary    5928.468
dtype: float64
```

```
In [114]: cluster_avg
```

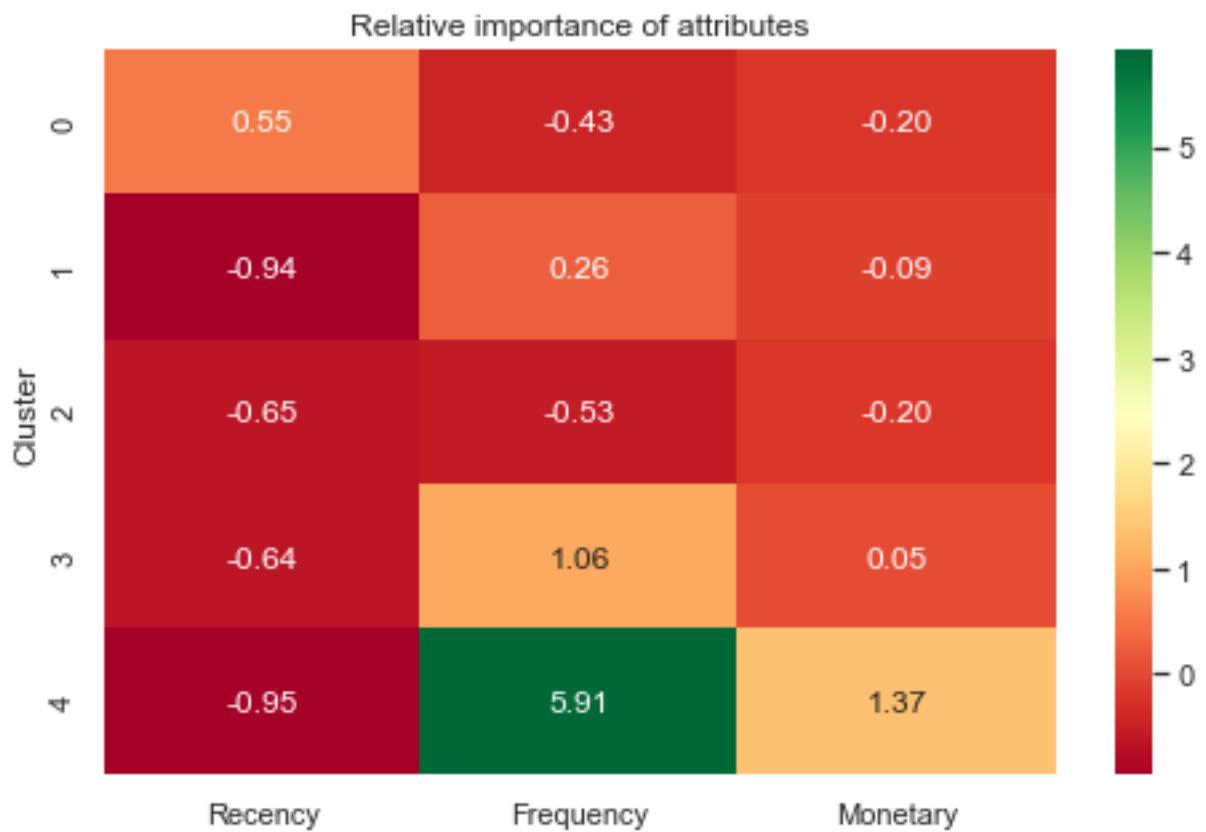
```
Out[114]:
   Recency  Frequency  Monetary
Cluster
  0  228.553557  1.718530  4742.588234
  1   9.105263  3.773994  5404.087043
  2  50.983963  1.395189  4750.991639
  3  52.691071  6.180357  6196.646055
  4   6.696035 20.735683 14061.101145
```

```
In [115]: relative_imp = cluster_avg / population_avg - 1
relative_imp.round(2)
```

```
Out[115]:
   Recency  Frequency  Monetary
```

```
Cluster
  0    0.55    -0.43   -0.20
  1   -0.94     0.26   -0.09
  2   -0.65    -0.53   -0.20
  3   -0.64     1.06    0.05
  4   -0.95     5.91    1.37
```

```
In [119]: # Plot heatmap
plt.figure(figsize=(8, 5))
plt.title('Relative importance of attributes')
sns.heatmap(data=relative_imp, annot=True, fmt=".2f", cmap='RdYlGn')
plt.show()
```



In [116]: Cluster_table

Out[116]:

	CustomerID	Recency	Frequency	Monetary	R_Quartile	F_Quartile	M_Quartile	RFMScore	RFM_Score	Segment	Cluster
0	12346.0	326	2	4288.63	1	2	1	121	4	None	0
1	12347.0	3	7	8598.63	4	4	4	444	12	Best Customers	4
2	12348.0	76	4	6085.87	2	3	4	234	9	None	3
3	12349.0	19	1	6046.18	3	1	4	314	8	None	2
4	12350.0	311	1	4623.03	1	1	2	112	4	None	0
...
4367	18280.0	278	1	4469.23	1	1	1	111	3	None	0
4368	18281.0	181	1	4369.45	1	1	1	111	3	None	0
4369	18282.0	8	3	4465.23	4	2	1	421	7	None	1
4370	18283.0	4	16	6334.16	4	4	4	444	12	Best Customers	4
4371	18287.0	43	3	6125.91	3	2	4	324	9	None	3

4372 rows × 11 columns

In [120]: df1

Out[120]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	Total_cost	Month	CohortIndex	Purchase_Da
0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom	2011-12-01	-168469.6	2011-05-01	8	2011-12-
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom	2011-01-01	-77183.6	2011-01-01	1	2011-01-
4	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom	2010-12-01	-280.8	2010-12-01	1	2010-12-
10	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	15749.0	United Kingdom	2011-04-01	-6539.4	2011-01-01	4	2011-04-

```
In [121]: FinalFile = pd.merge(df1, Cluster_table, how = 'left', on='CustomerID')
```

```
In [122]: FinalFile
```

```
Out[122]:
```

		InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	Total_cost	...	Recency	Frequency	Monetary
	0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom	2011-12-01	-168469.6	...	1	3	4291.5
	1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom	2011-01-01	-77183.6	...	326	2	4288.6
	2	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom	2010-12-01	-280.8	...	12	21	37639.3
	3	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	15749.0	United Kingdom	2011-04-01	-6539.4	...	236	4	25824.5
	4	C550456	21175	GIN + TONIC DIET METAL SIGN	-2000	2011-04-18 13:08:00	1.85	15749.0	United Kingdom	2011-04-01	-3700.0	...	236	4	25824.5

	401599	554868	22197	SMALL POPCORN HOLDER	4300	2011-05-27 10:52:00	0.72	13135.0	United Kingdom	2011-05-01	3096.0	...	197	1	7384.6
	401600	573008	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	4800	2011-10-27 12:26:00	0.21	12901.0	United Kingdom	2011-10-01	1008.0	...	9	36	20581.7
	401601	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540	2011-11-25 15:57:00	0.00	13256.0	United Kingdom	2011-11-01	0.0	...	15	1	4288.6
	401602	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	2011-01-18 10:01:00	1.04	12346.0	United Kingdom	2011-01-01	77183.6	...	326	2	4288.6

```
In [128]: relative_imp.round(2)
```

```
Out[128]:
```

Cluster	Recency	Frequency	Monetary
0	0.55	-0.43	-0.20
1	-0.94	0.26	-0.09
2	-0.65	-0.53	-0.20
3	-0.64	1.06	0.05
4	-0.95	5.91	1.37

```
In [129]: FF = pd.merge(FinalFile, relative_imp, how = 'left', on='Cluster')
```

```
In [130]: FF
```

jupyter Capstone Project 3 - Retail Last Checkpoint: an hour ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 Logout

InvoiceNo StockCode Description Quantity InvoiceDate UnitPrice CustomerID Country InvoiceMonth Total_cost ... R_Quartile F_Quartile M_Qua

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	Total_cost	R_Quartile	F_Quartile	M_Qua
0	C581484	23843	PAPER CRAFT LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom	2011-12-01	-168499.6	4	2	
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom	2011-01-01	-77183.6	1	2	
2	C536757	84347	ROTATING SILVER ANGEL LIGHT HOLDER	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom	2010-12-01	-280.8	4	4	
3	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	15749.0	United Kingdom	2011-04-01	-6539.4	1	3	
4	C550456	21175	GIN + TONIC DIET METAL SIGN	-2000	2011-04-18 13:08:00	1.85	15749.0	United Kingdom	2011-04-01	-3700.0	1	3	
...	
401599	554868	22197	SMALL POPCORN HOLDER	4300	2011-05-27 10:52:00	0.72	13135.0	United Kingdom	2011-05-01	3096.0	1	1	
401600	573008	84077	WORLD VIEWS GLIDERS ASS'D DESIGNS	4800	2011-10-27 12:26:00	0.21	12901.0	United Kingdom	2011-10-01	1008.0	4	4	
401601	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540	2011-11-25 15:57:00	0.00	13256.0	United Kingdom	2011-11-01	0.0	4	1	
401602	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	2011-01-18 10:01:00	1.04	12346.0	United Kingdom	2011-01-01	-77183.6	1	2	
401603	581483	23843	PAPER CRAFT LITTLE BIRDIE	80995	2011-12-09 09:15:00	2.08	16446.0	United Kingdom	2011-12-01	-168499.6	4	2	

401604 rows x 26 columns

```
In [131]: FF.rename(columns = {'Recency_y':'Recency_Img'}, inplace = True)
FF.rename(columns = {'Frequency_y':'Frequency_Img'}, inplace = True)
FF.rename(columns = {'Monetary_y':'Monetary_Img'}, inplace = True)
FF
```

Out[131]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	InvoiceMonth	Total_cost	...	R_Quartile	F_Quartile	M_Qua
0	C581484	23843	PAPER CRAFT , LITTLE BIRDIE	-80995	2011-12-09 09:27:00	2.08	16446.0	United Kingdom	2011-12-01	-168469.6	...	4	2	
1	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	2011-01-18 10:17:00	1.04	12346.0	United Kingdom	2011-01-01	-77183.6	...	1	2	
2	C536757	84347	ROTATING SILVER ANGELS T-LIGHT HLDR	-9360	2010-12-02 14:23:00	0.03	15838.0	United Kingdom	2010-12-01	-280.8	...	4	4	
3	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-3114	2011-04-18 13:08:00	2.10	15749.0	United Kingdom	2011-04-01	-6539.4	...	1	3	
4	C550456	21175	GIN + TONIC DIET METAL SIGN	-2000	2011-04-18 13:08:00	1.85	15749.0	United Kingdom	2011-04-01	-3700.0	...	1	3	
...	
401599	554868	22197	SMALL POPCORN HOLDER	4300	2011-05-27 10:52:00	0.72	13135.0	United Kingdom	2011-05-01	3096.0	...	1	1	
401600	573008	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	4800	2011-10-27 12:26:00	0.21	12901.0	United Kingdom	2011-10-01	1008.0	...	4	4	
401601	578841	84826	ASSTD DESIGN 3D PAPER STICKERS	12540	2011-11-25 15:57:00	0.00	13256.0	United Kingdom	2011-11-01	0.0	...	4	1	
401602	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	2011-01-18 10:01:00	1.04	12346.0	United Kingdom	2011-01-01	77183.6	...	1	2	

Assessment | Query Update - Project Related - | Home Page - Select or create a | Capstone Project 3 - Retail - Jupyter Notebook | Untitled | Untitled | Logout

jupyter Capstone Project 3 - Retail Last Checkpoint: an hour ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

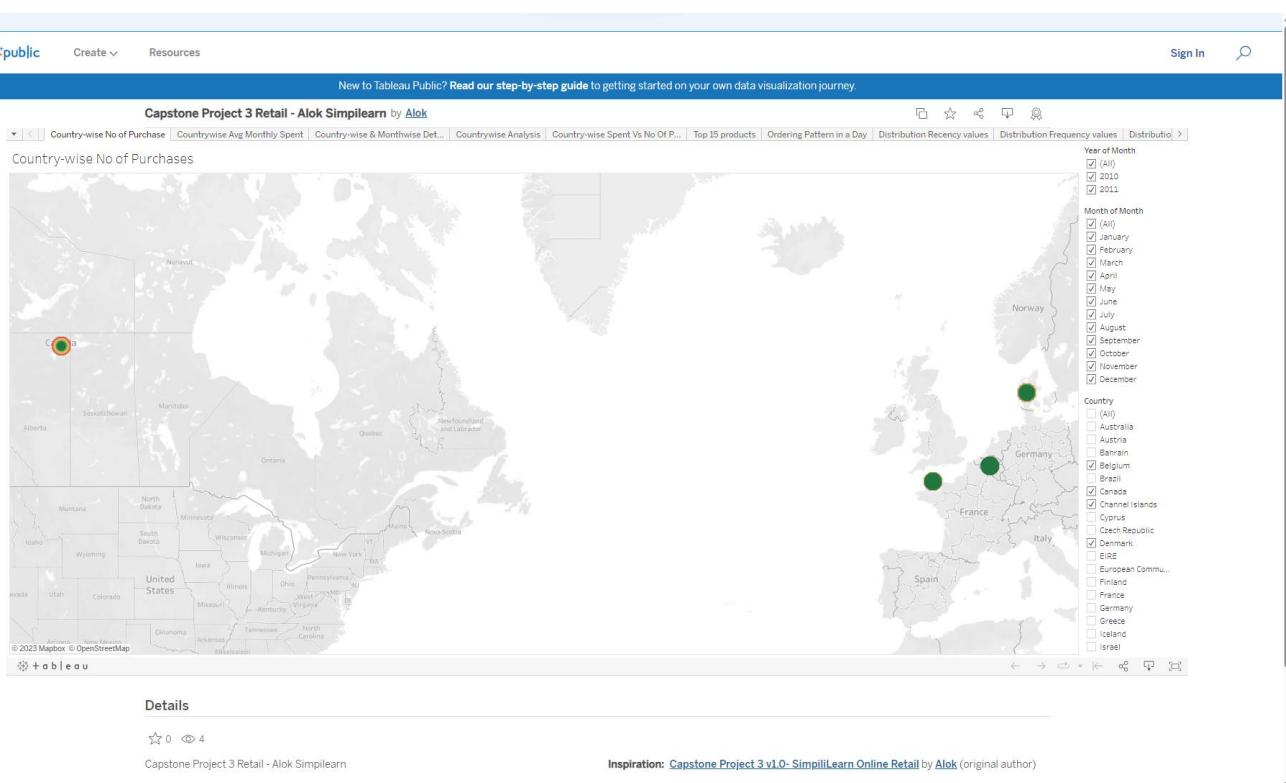
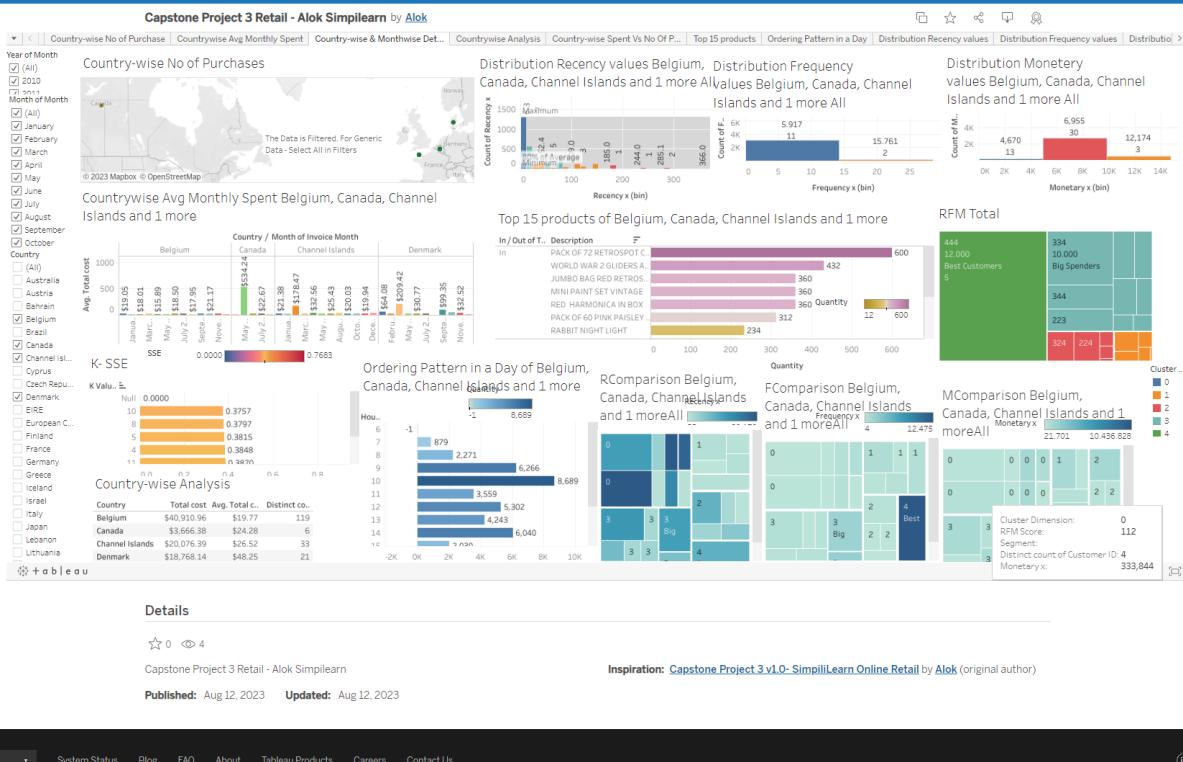
In [222]: `i= 0
wcss_silhouette
sse = pd.DataFrame(wcss_silhouette)
sse`

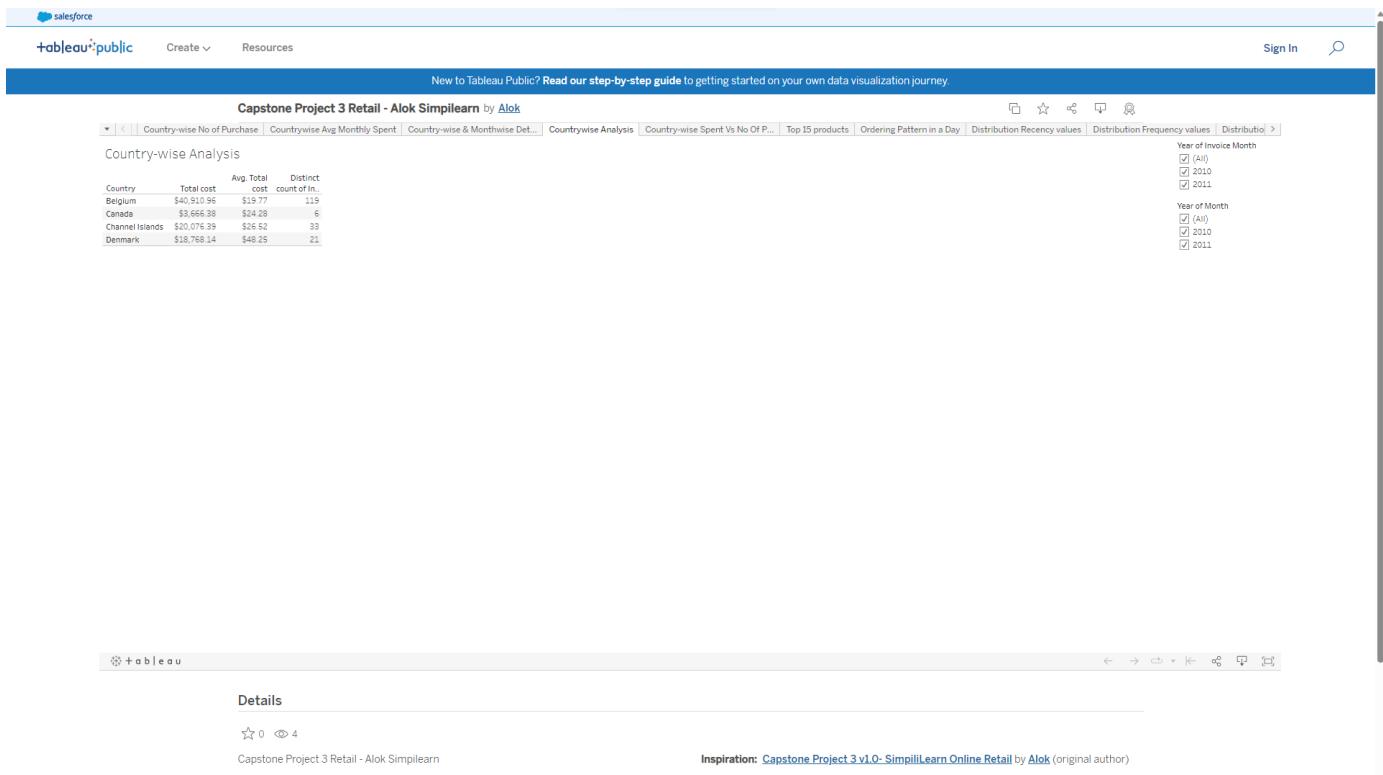
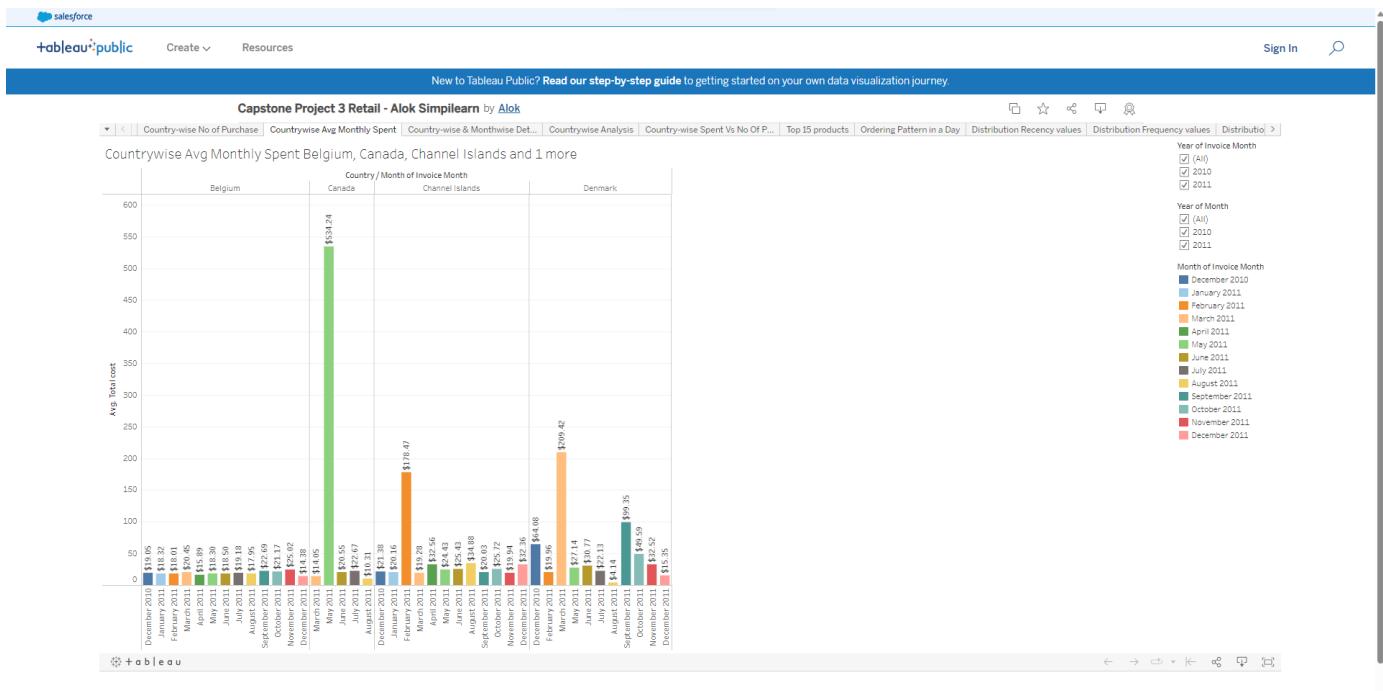
Out[222]: `0
0 0.463203
1 0.378824
2 0.389429
3 0.384806
4 0.381462
5 0.377401
6 0.374921
7 0.379746
8 0.375727
9 0.388960`

In [223]: `sse.to_excel(r'D:\OneDrive\Studies\AI - ML\Capstone Project\SSE.xlsx')`

In [132]: `FF.to_excel('Cluster_file.xlsx')`

SSE file is to create the SSE in Tableau. These two Tables are picked added as source in Tableau as independent tables. The last part is in Tableau Public. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following a. Country-wise analysis to demonstrate average spend. - Use a bar chart to show the monthly figures b. Bar graph of top 15 products which are mostly ordered by the users to show the number of products sold c. Bar graph to show the count of orders vs. hours throughout the day d. Plot the distribution of RFM values using histogram and frequency charts e. Plot error (cost) vs. number of clusters selected f. Visualize to compare the RFM values of the clusters using heatmap







Details

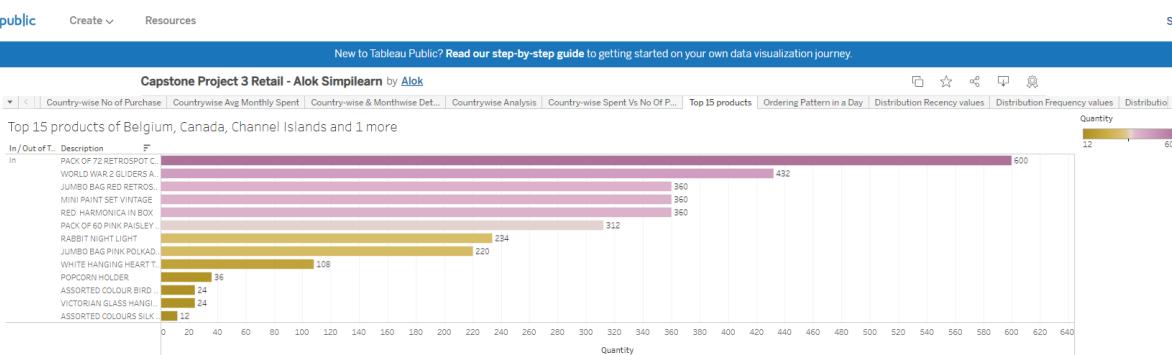
☆ 0 4

Capstone Project 3 Retail - Alok Simplearn

Inspiration: [Capstone Project 3 v1.0- SimpliiLearn Online Retail](#) by Alok (original author)

Published: Aug 12, 2023 Updated: Aug 12, 2023

Sign in



Details

☆ 0 4

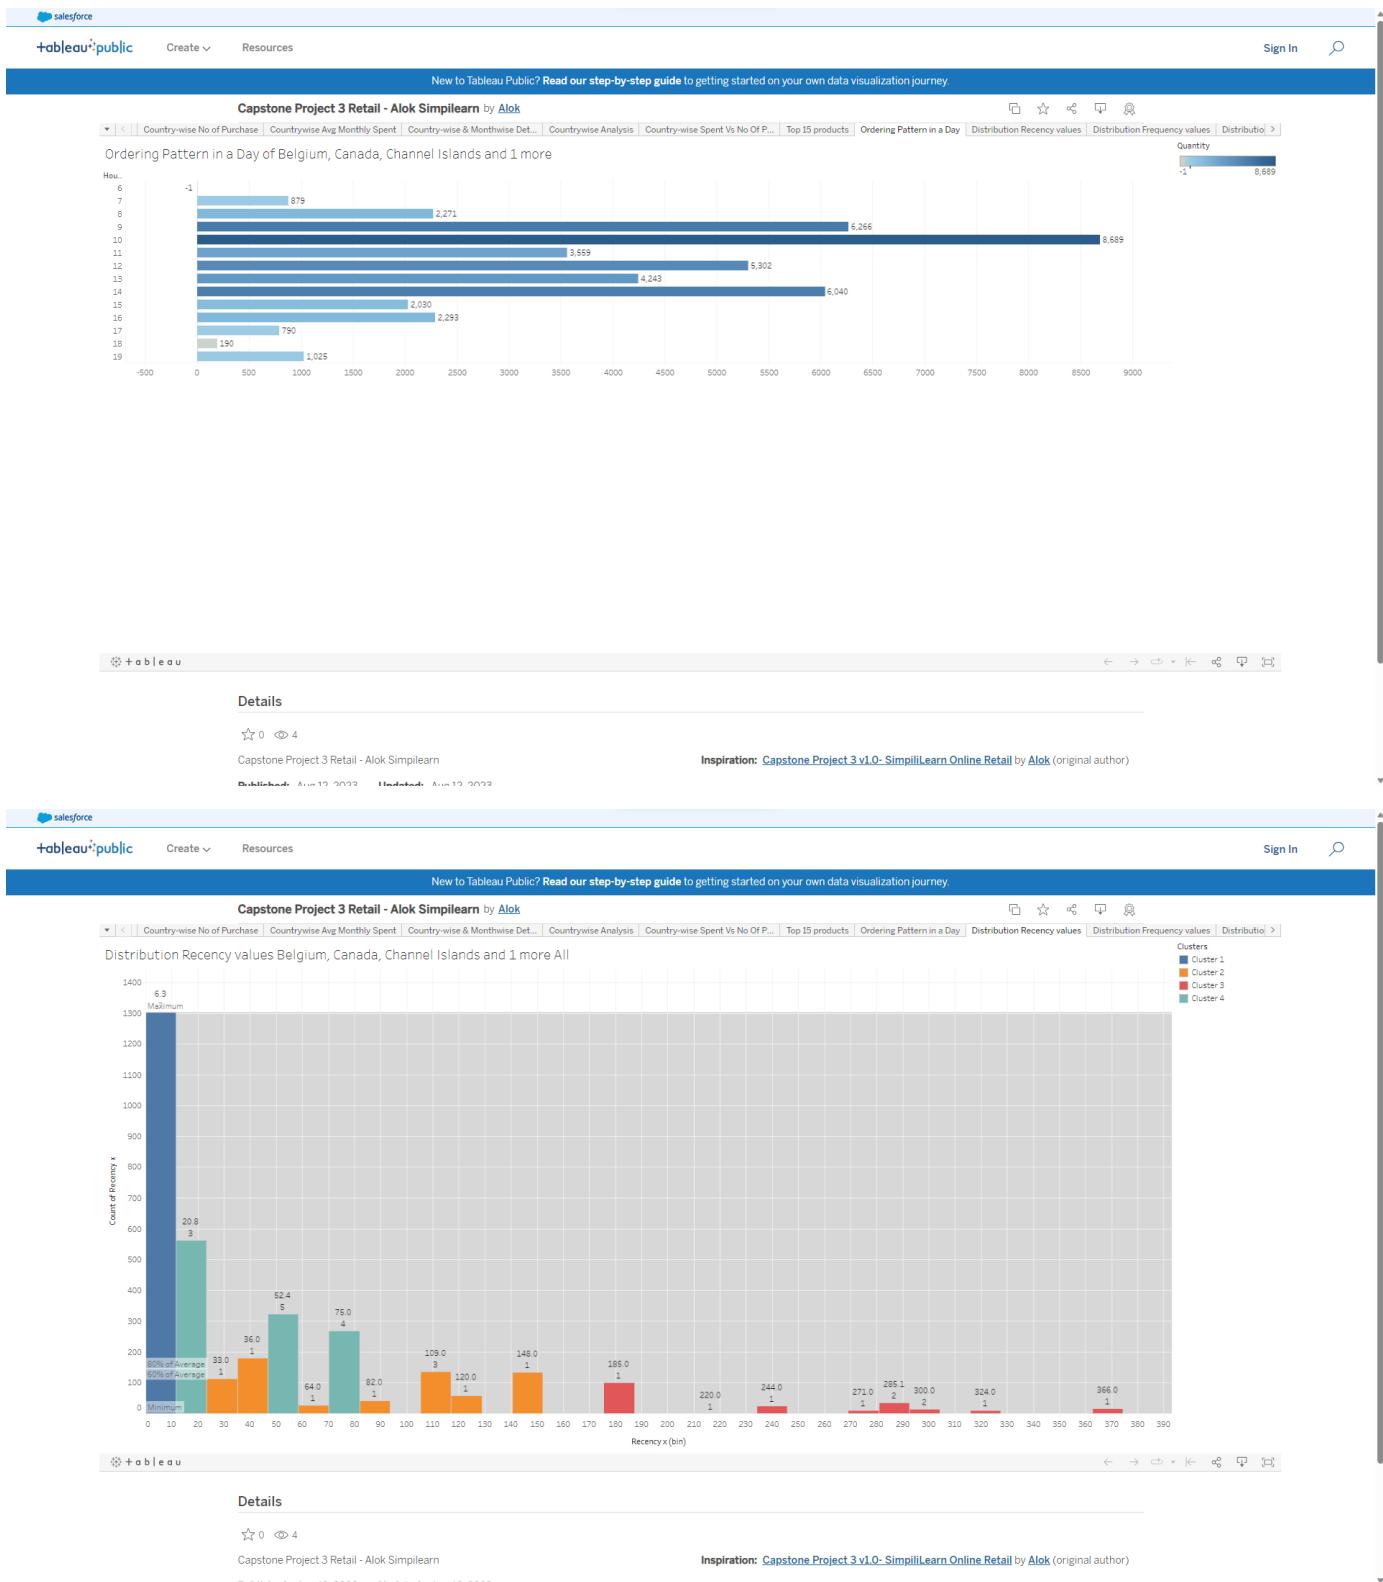
Capstone Project 3 Retail - Alok Simplearn

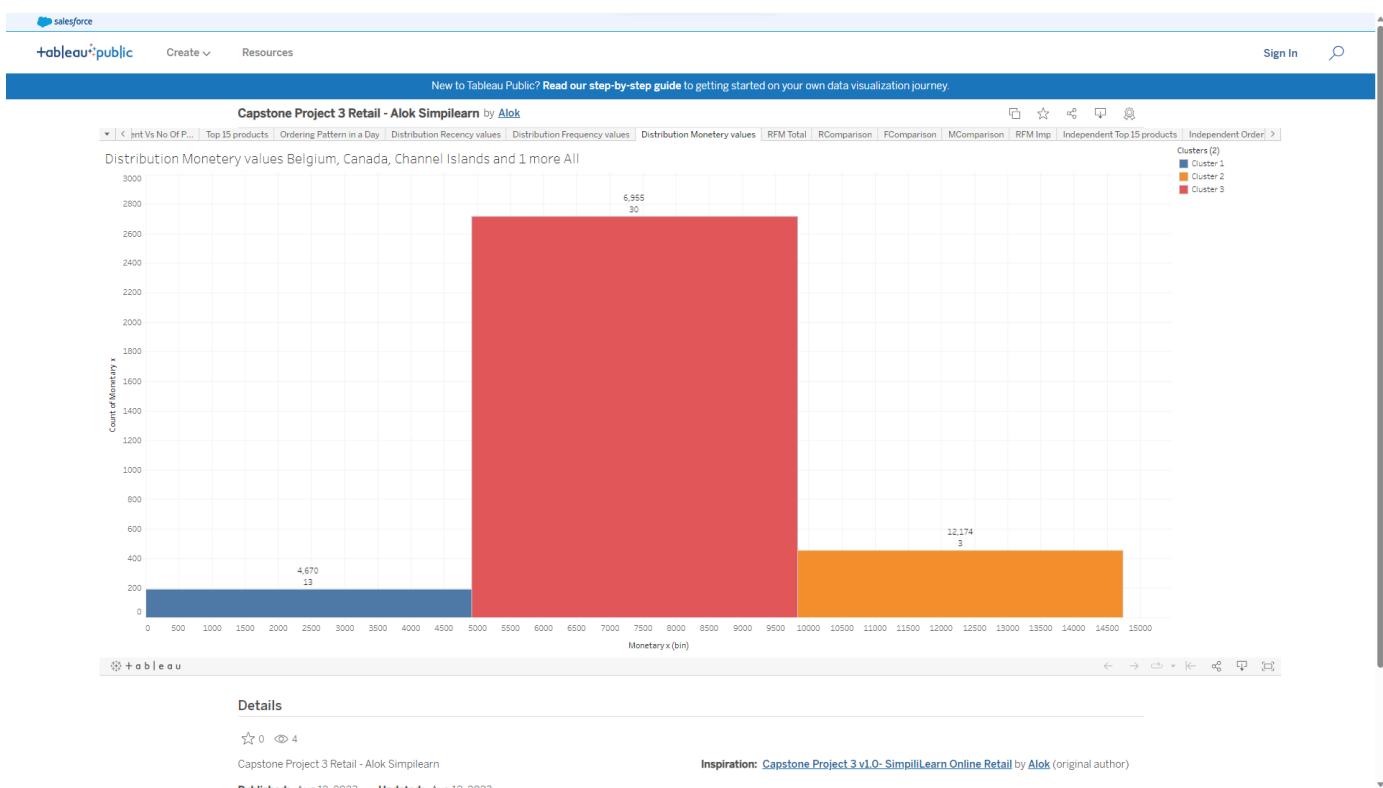
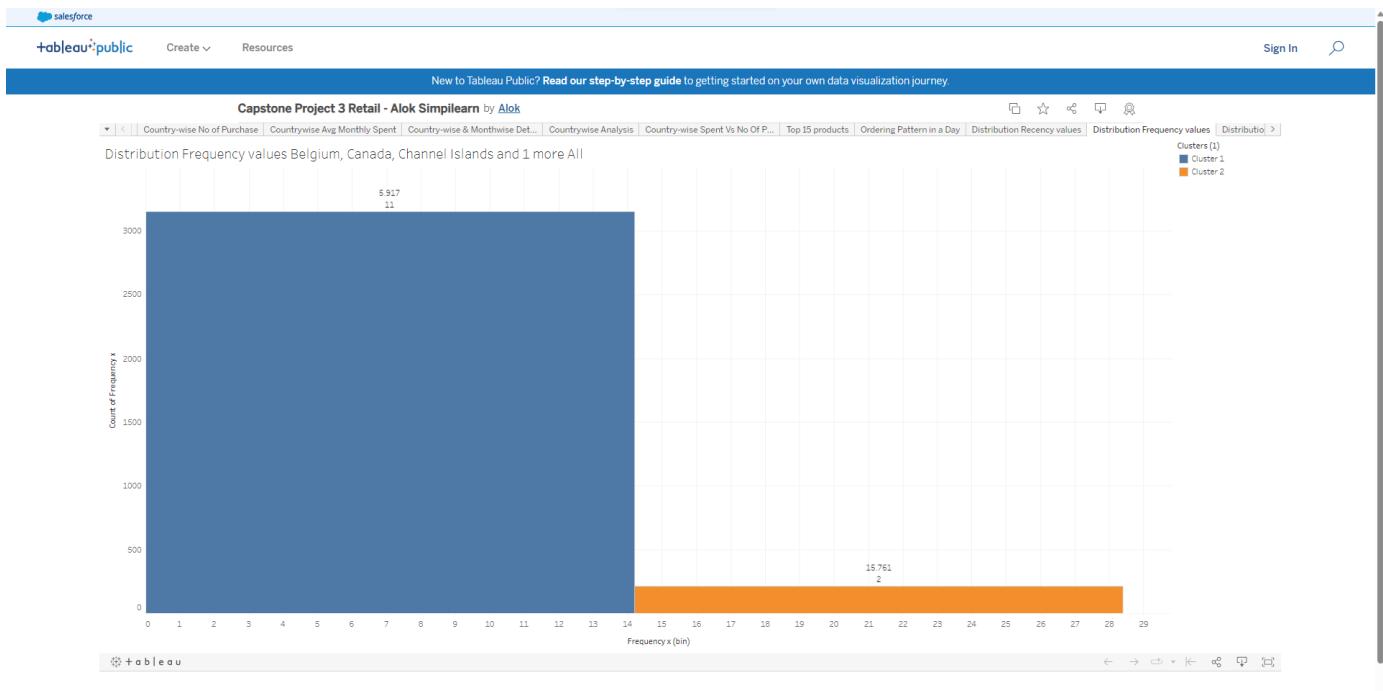
Inspiration: [Capstone Project 3 v1.0- SimpliiLearn Online Retail](#) by Alok (original author)

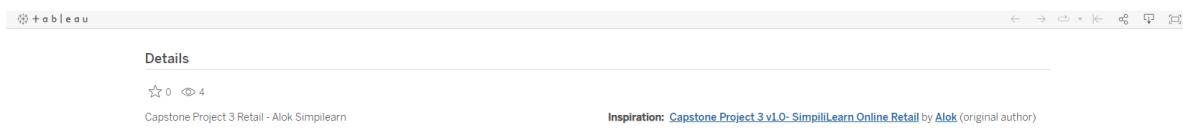
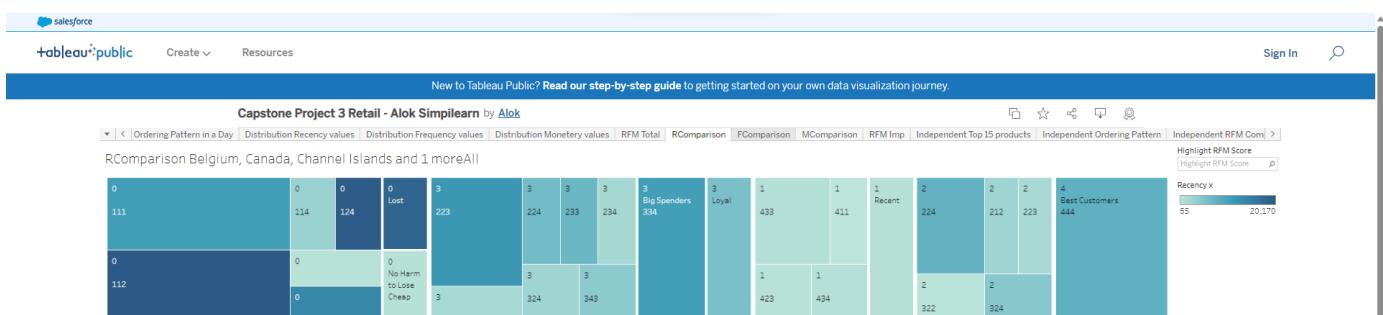
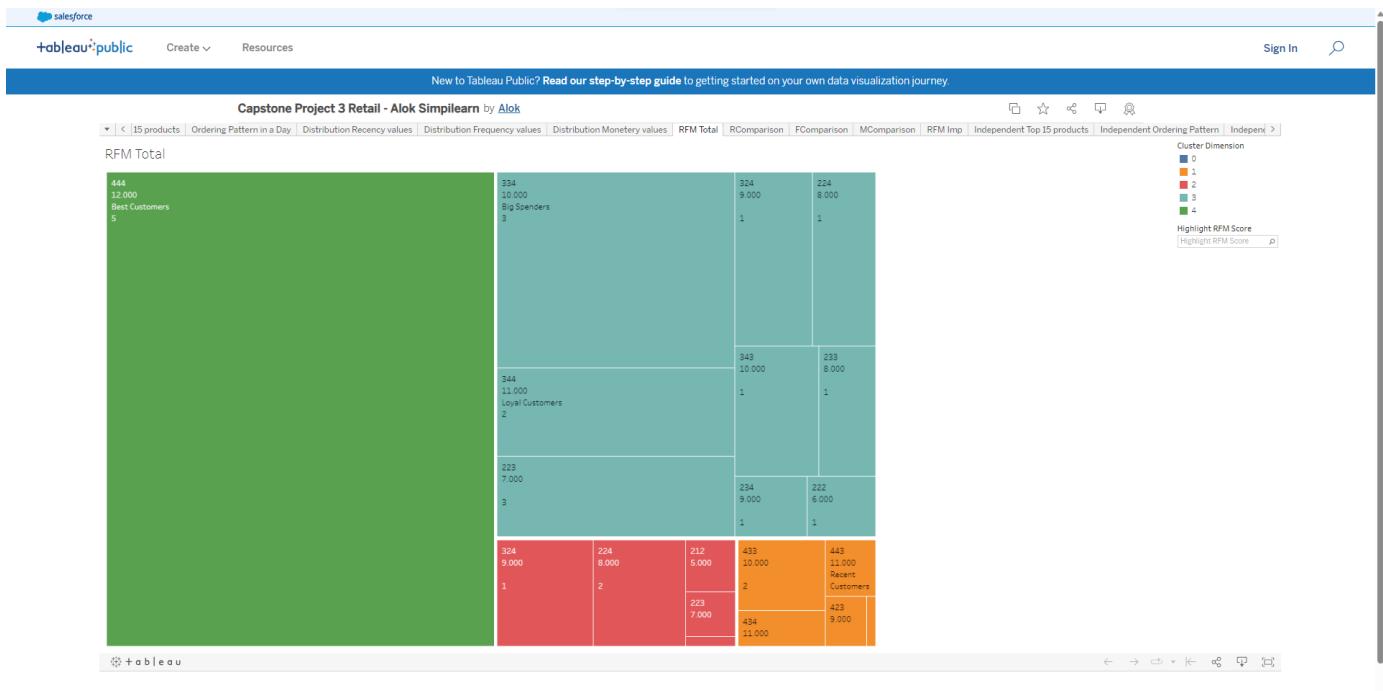
Published: Aug 12, 2023 Updated: Aug 12, 2023

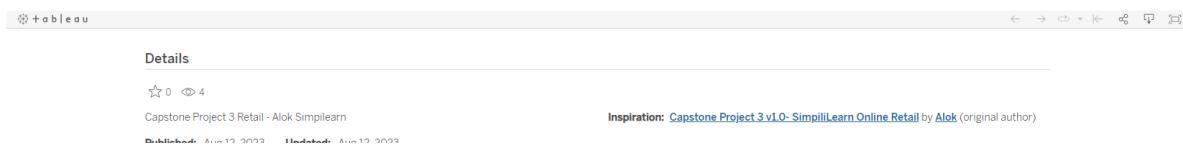
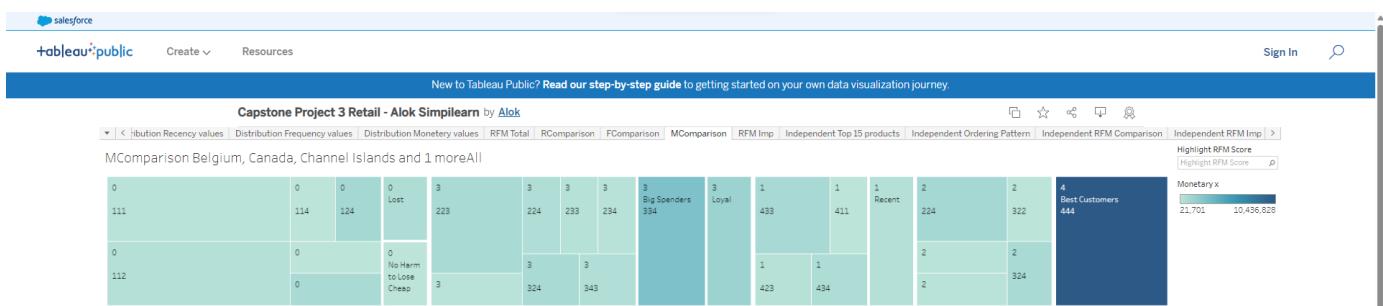
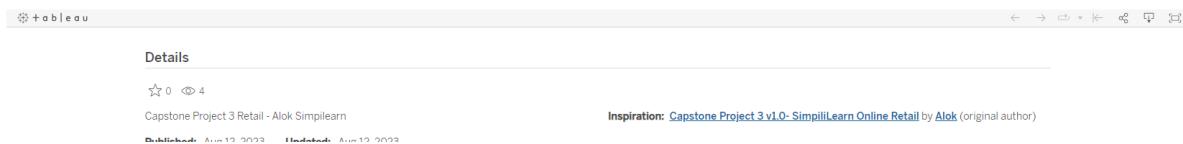
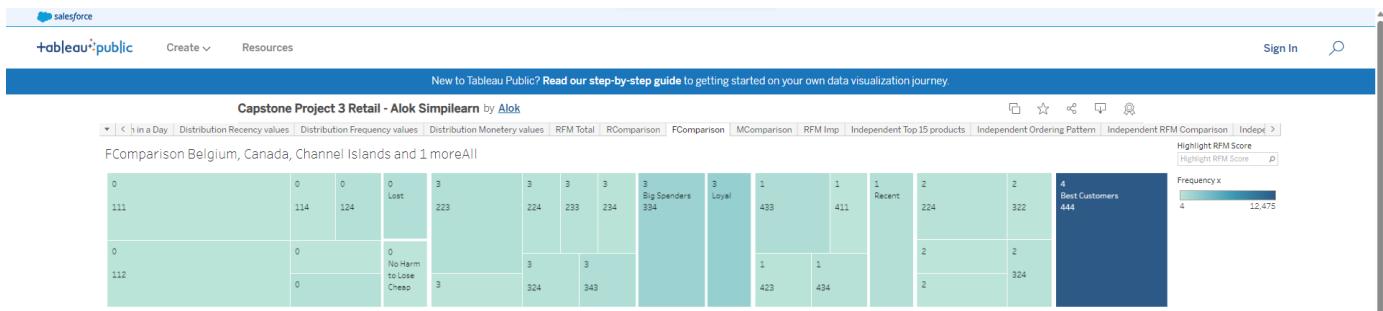
Sign in

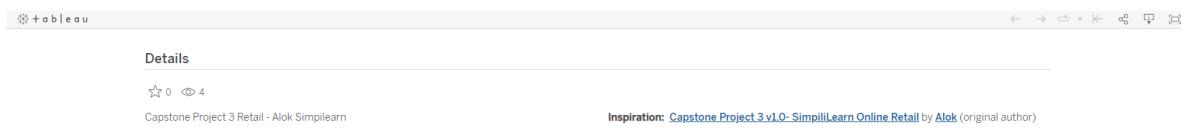
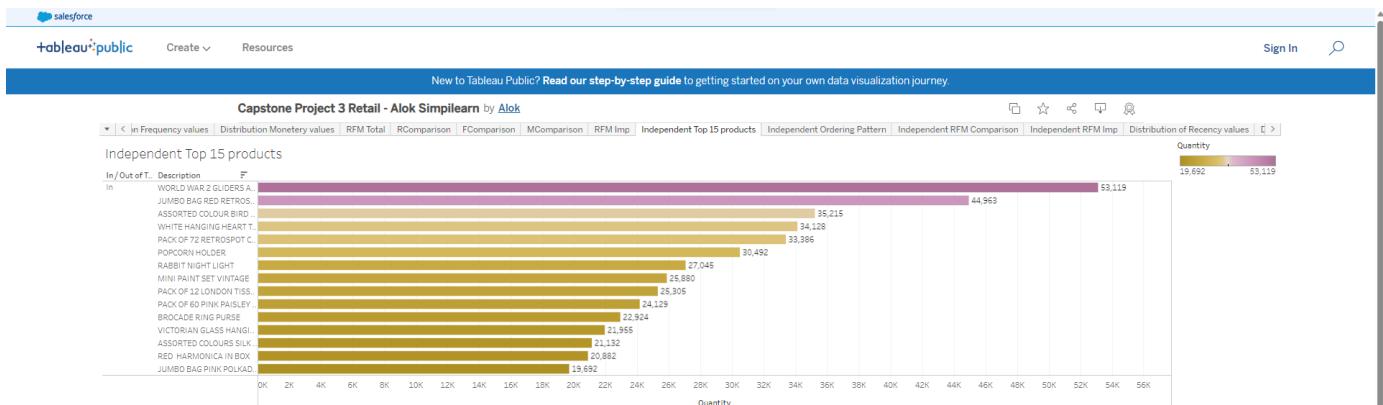
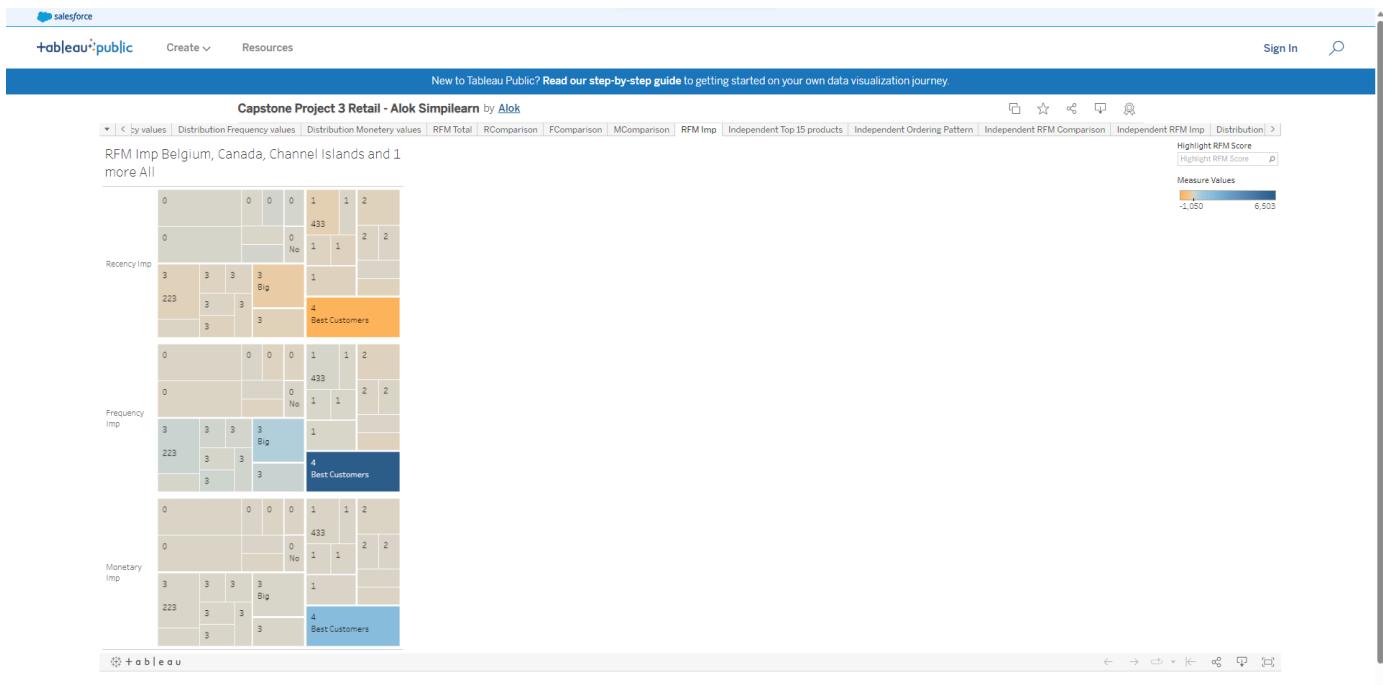


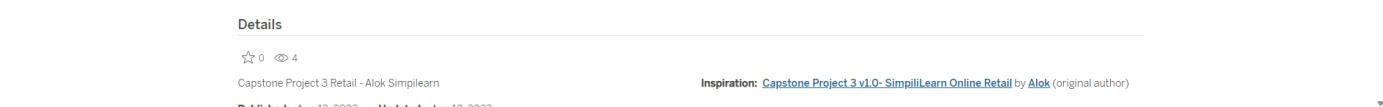
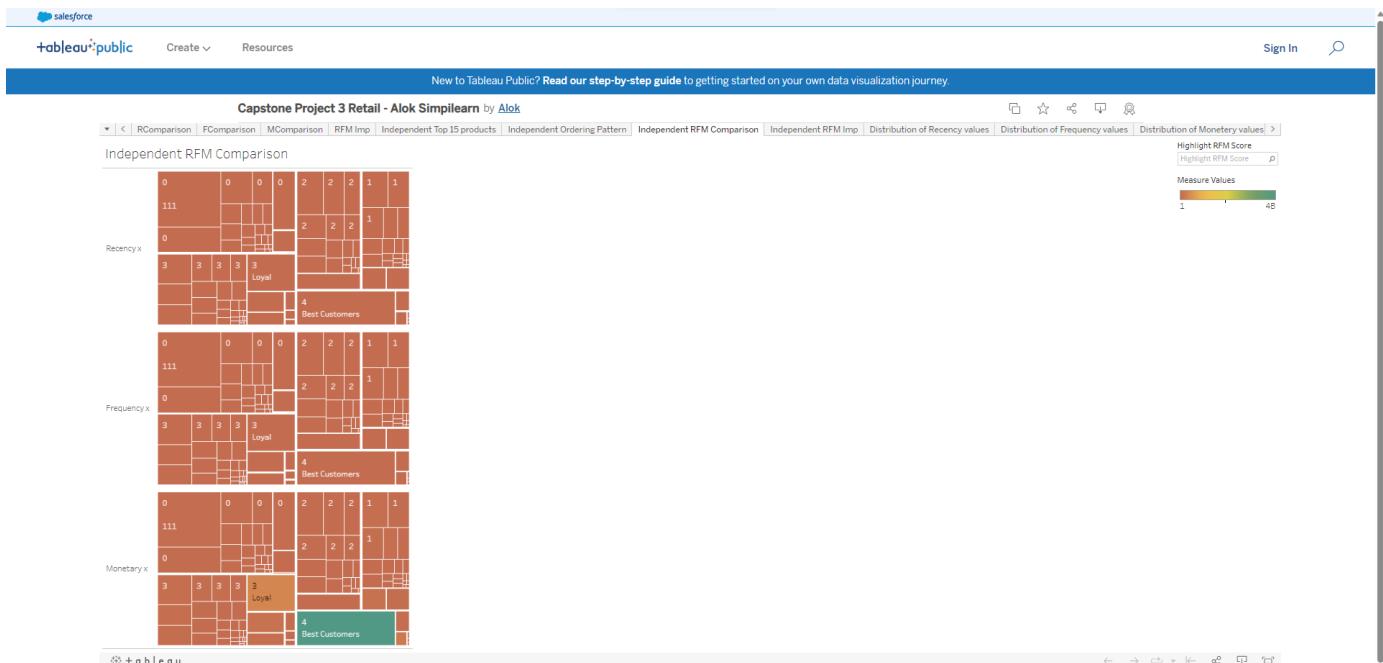
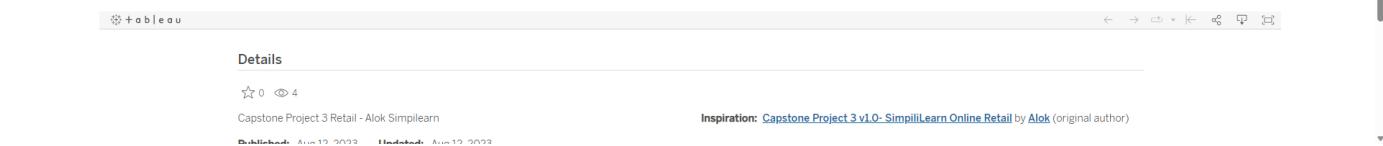
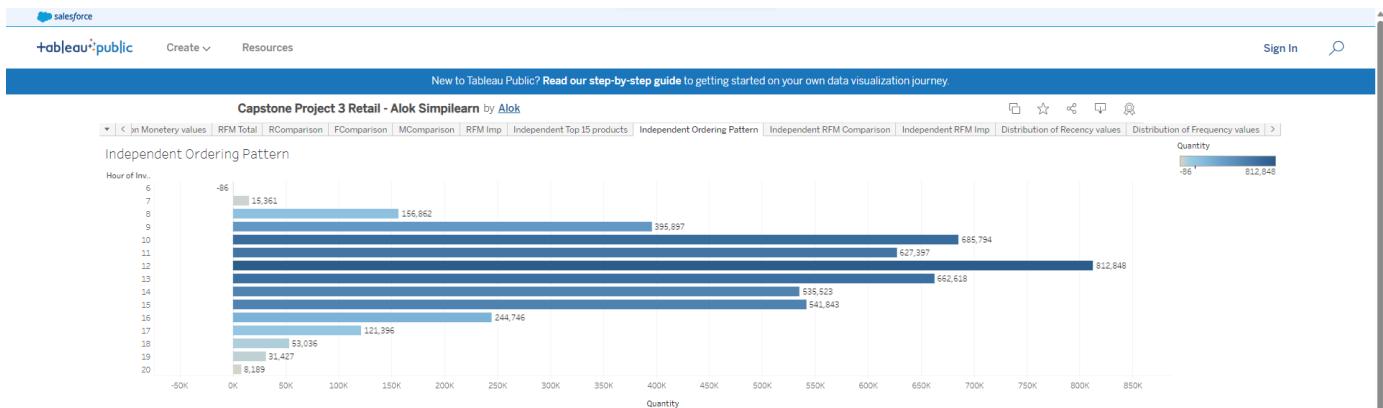


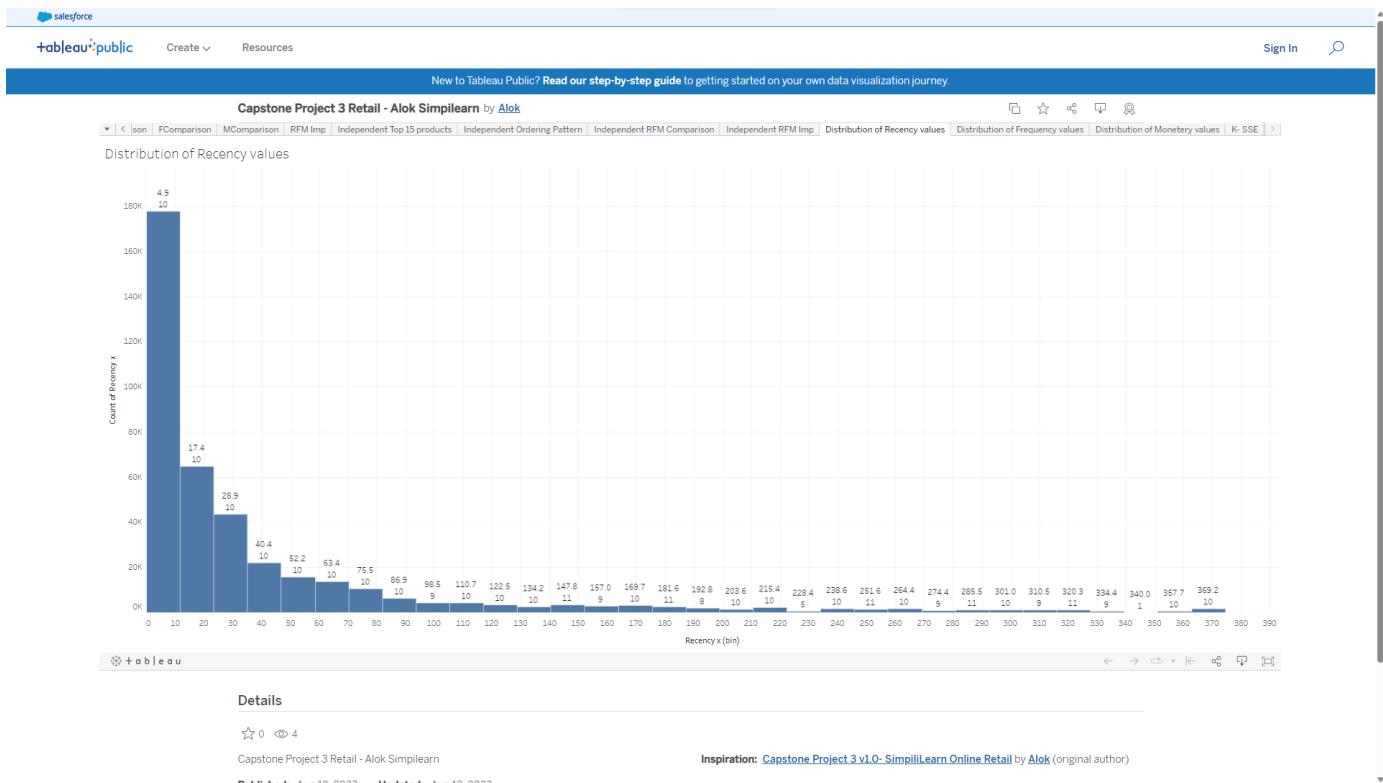
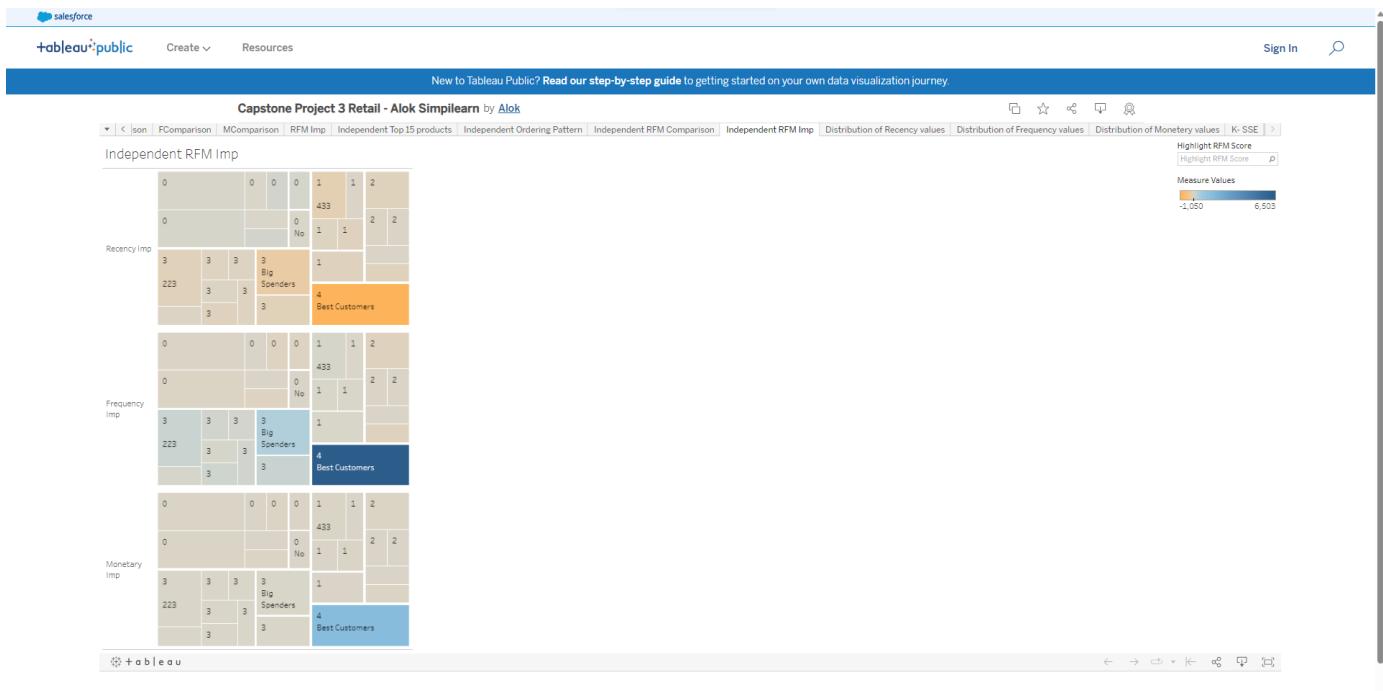


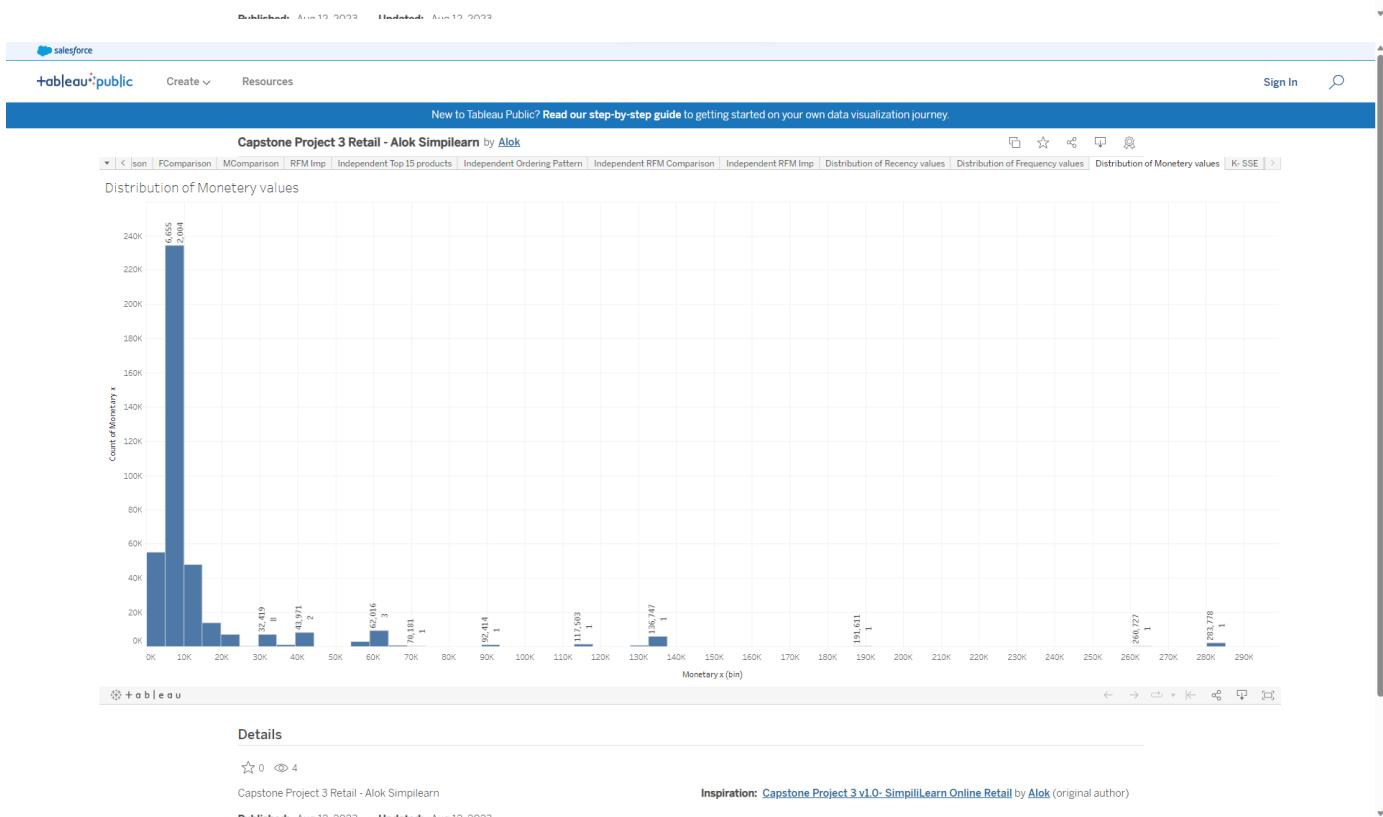
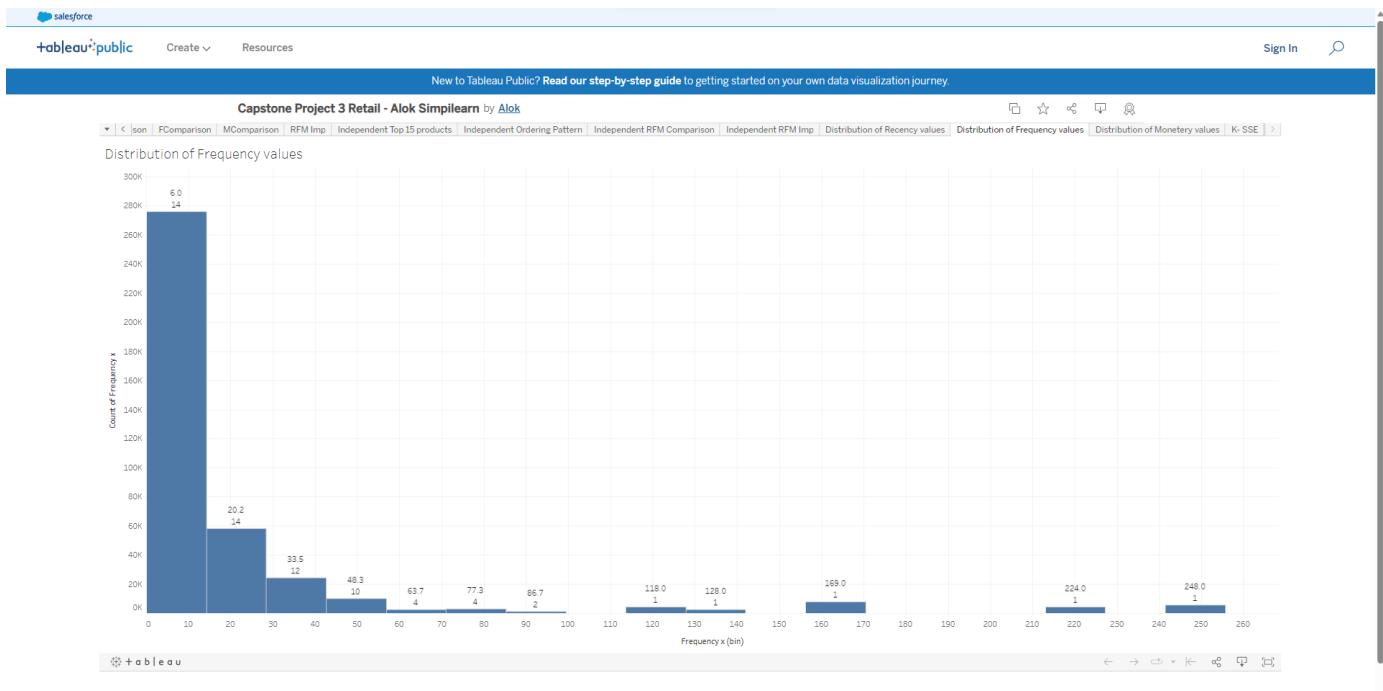


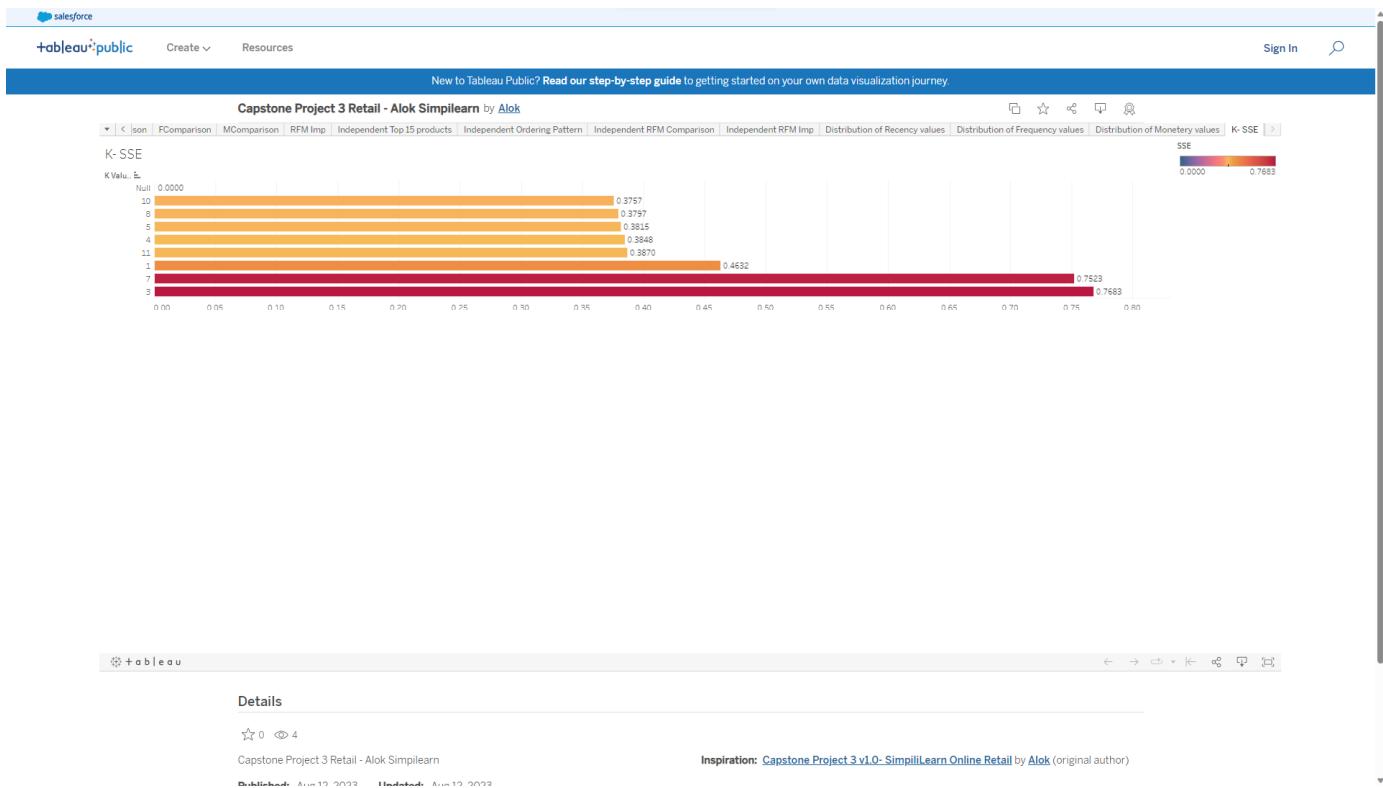












Links

GitHub: https://github.com/ks-alokranjan/Simplilearn_Capstone_Retail

Tableau Link: [Capstone Project 3 Retail - Alok Simplelearn | Tableau Public](#)